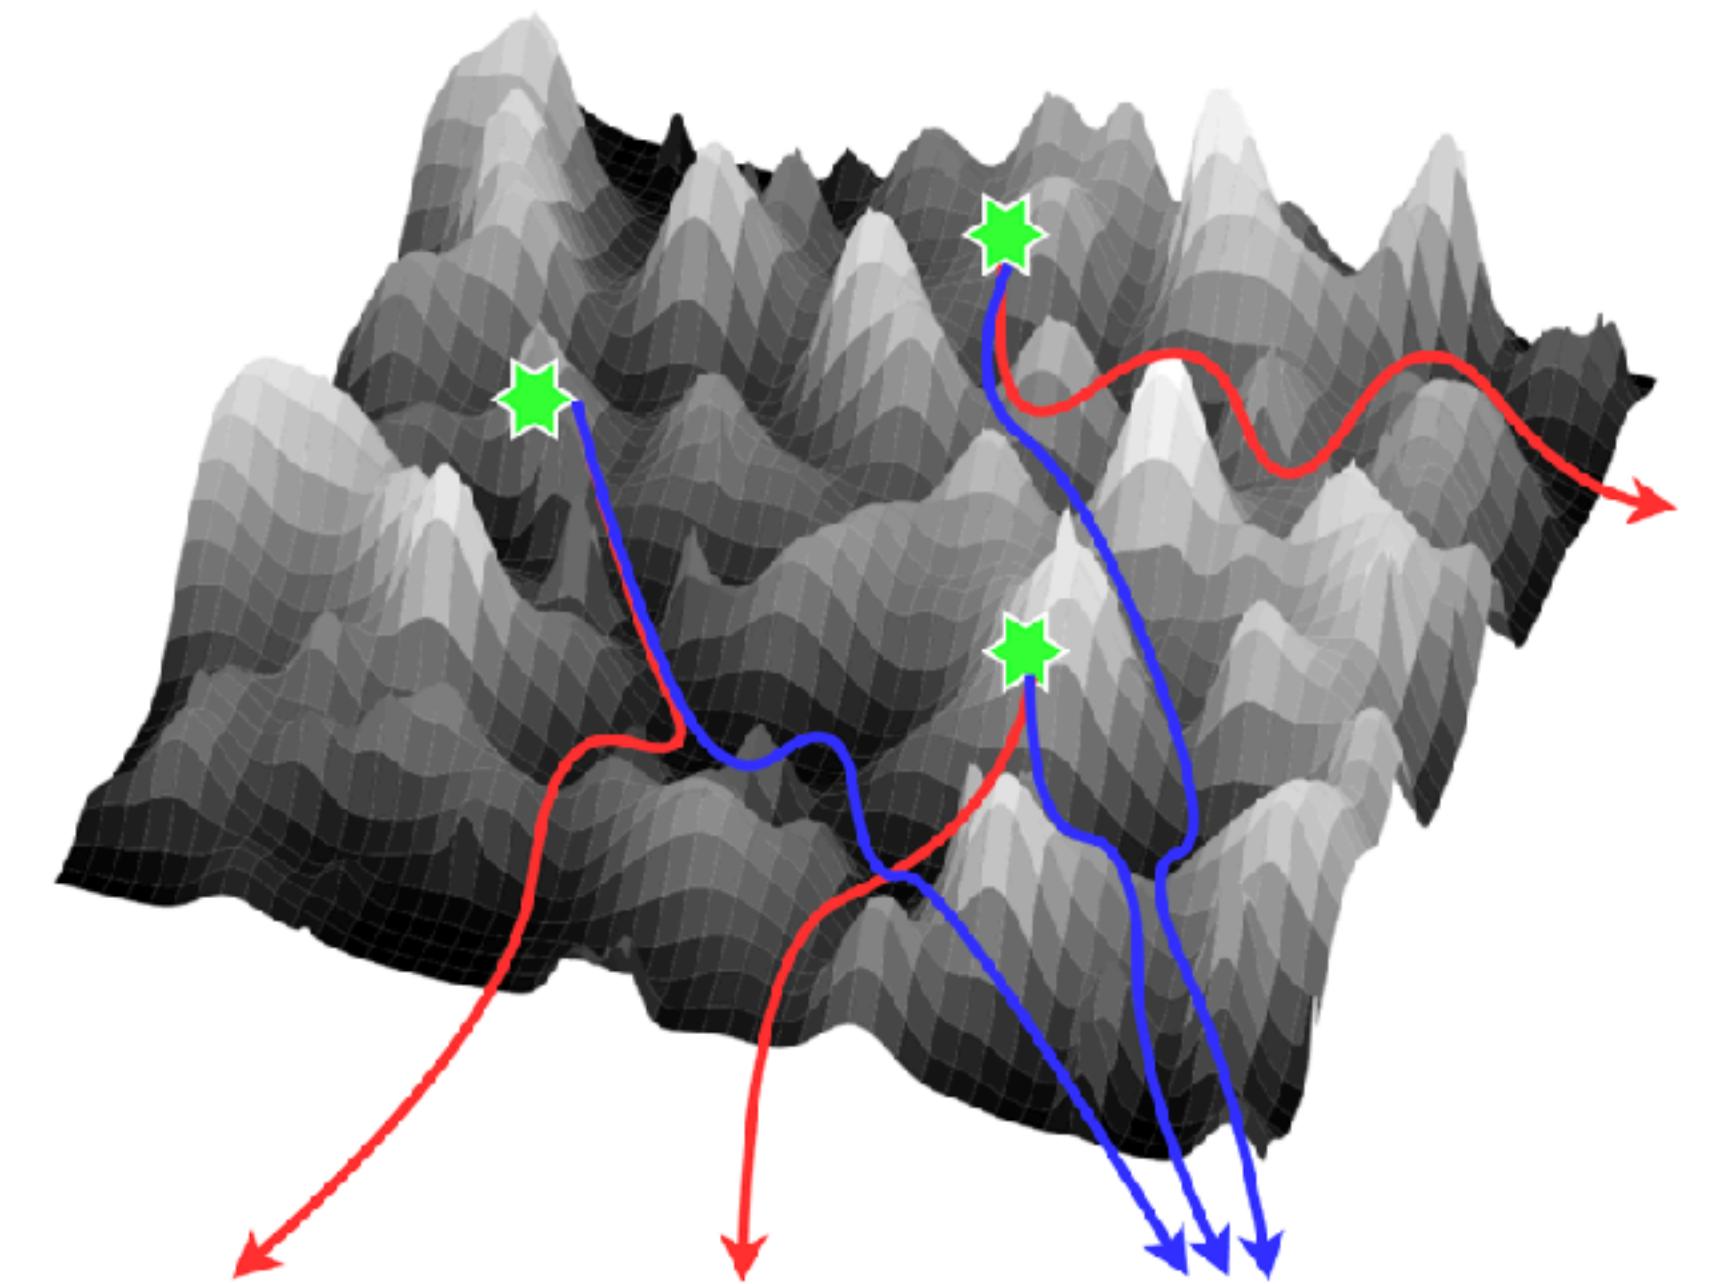


Feature learning from non-Gaussian inputs

Sebastian Goldt (SISSA, Trieste)

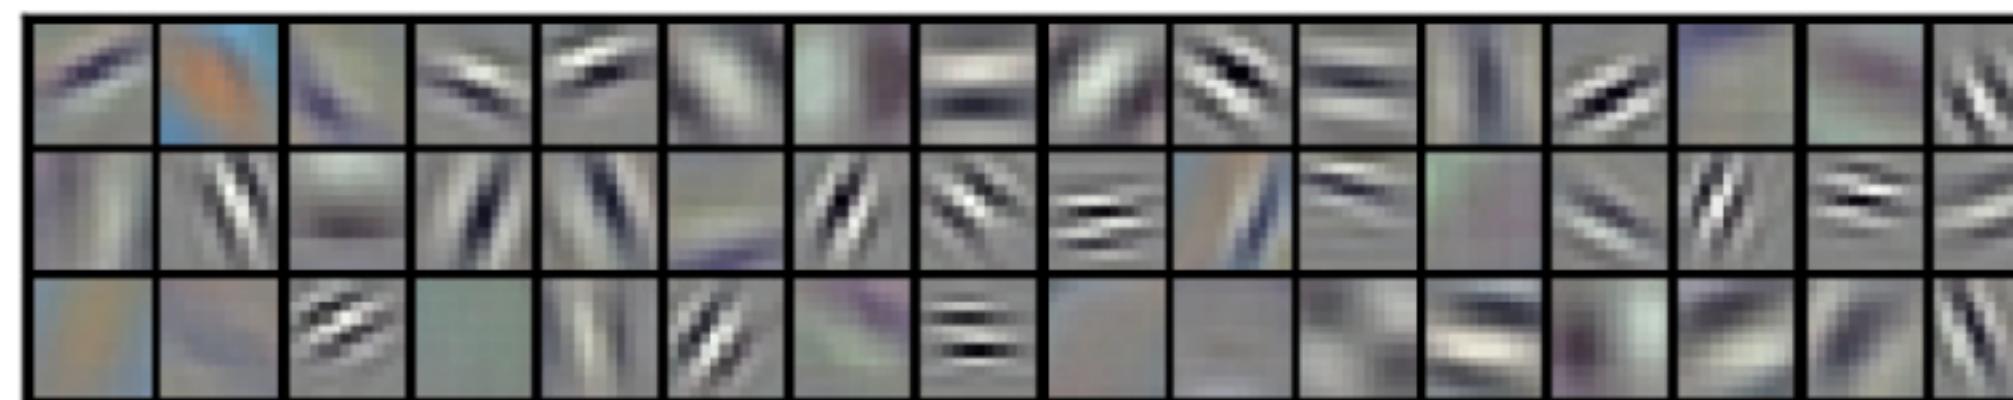
joint work w/ Lorenzo Bardone and Fabiola Ricci



**What do neural networks learn
from their inputs?**

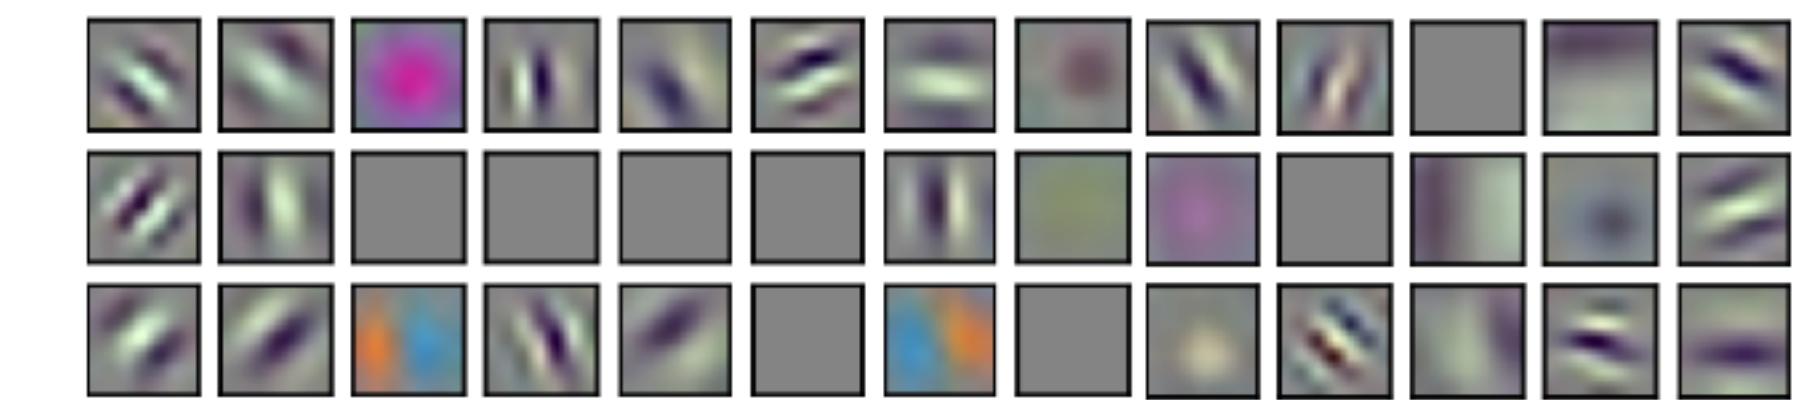
Neural networks learn stereotypical features

First-layer filters learnt from ImageNet resemble Gabor filters across architectures

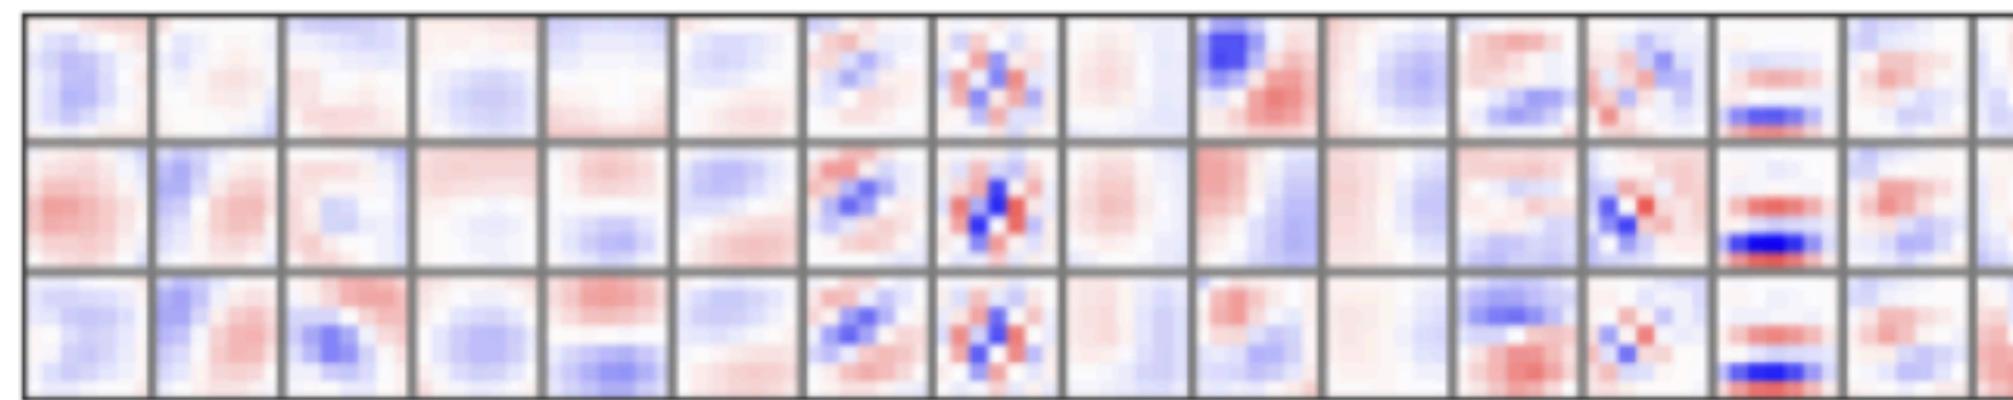


AlexNet

Krizhevsky, Sutskever, Hinton (2012)

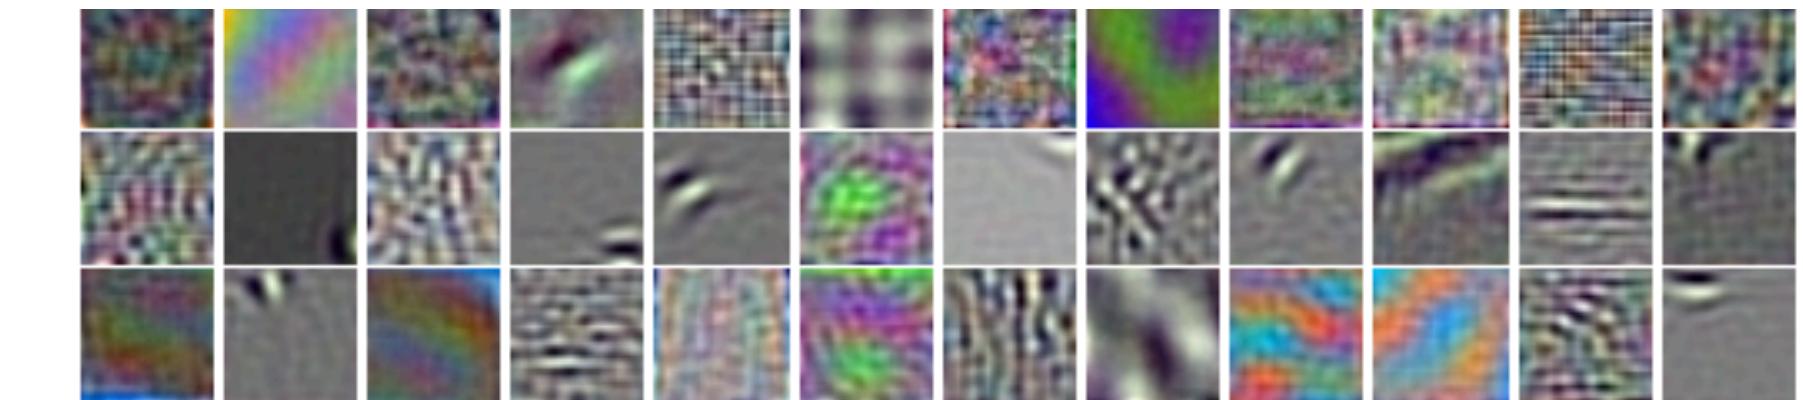


DenseNet121



VGG-11

Guth & Ménard (2024)



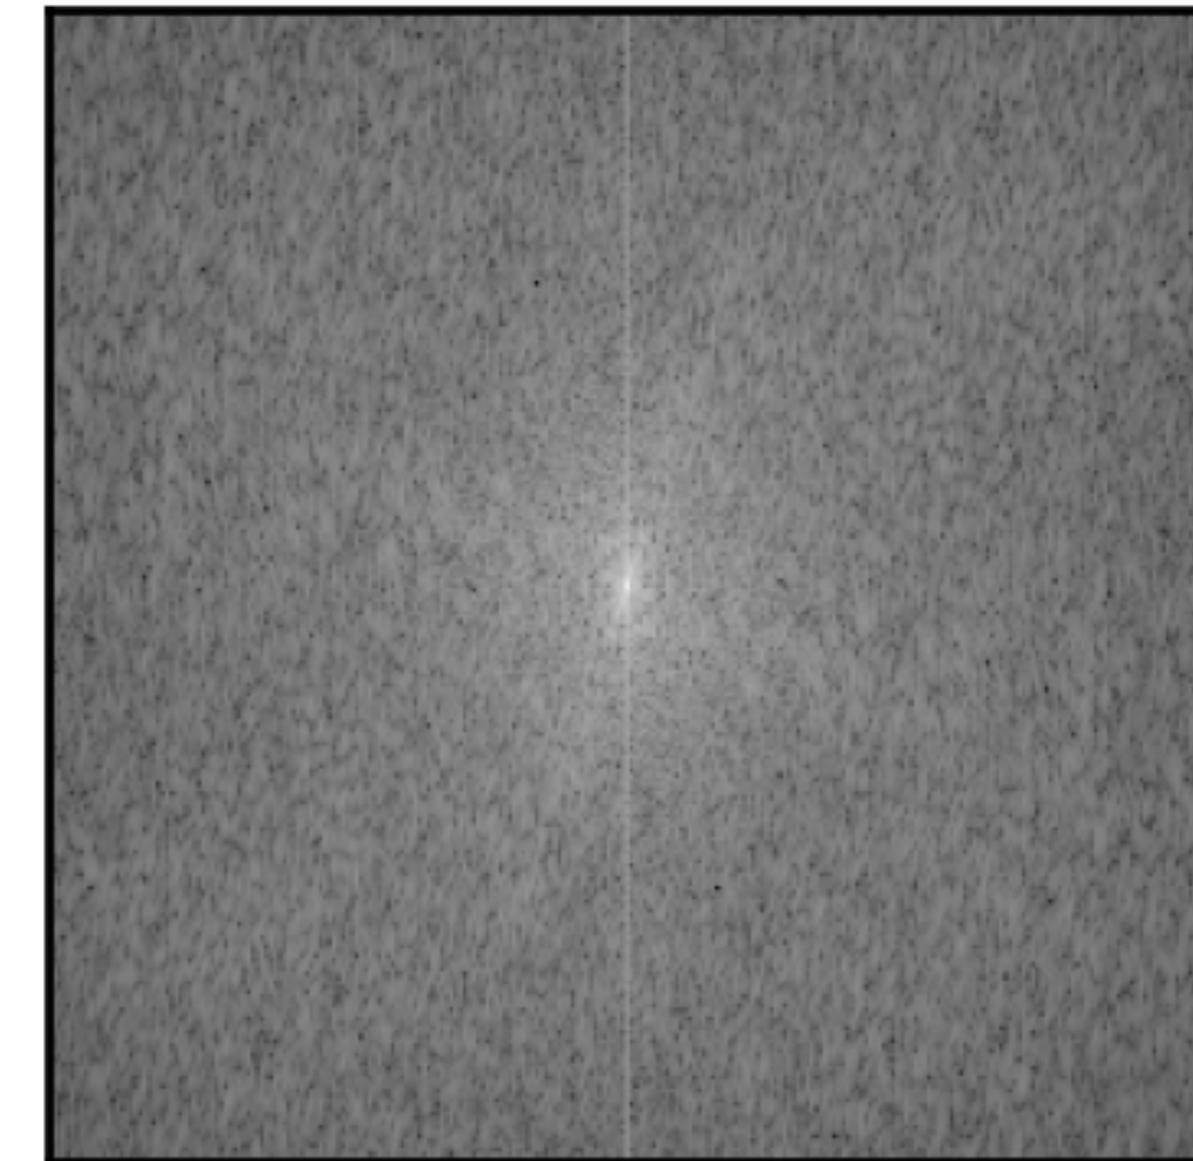
MLP mixer

Tolstikhin et al. NeurIPS '21

Convergence of features across architectures —
inputs drive feature learning!

What is in an image?

A Fourier perspective

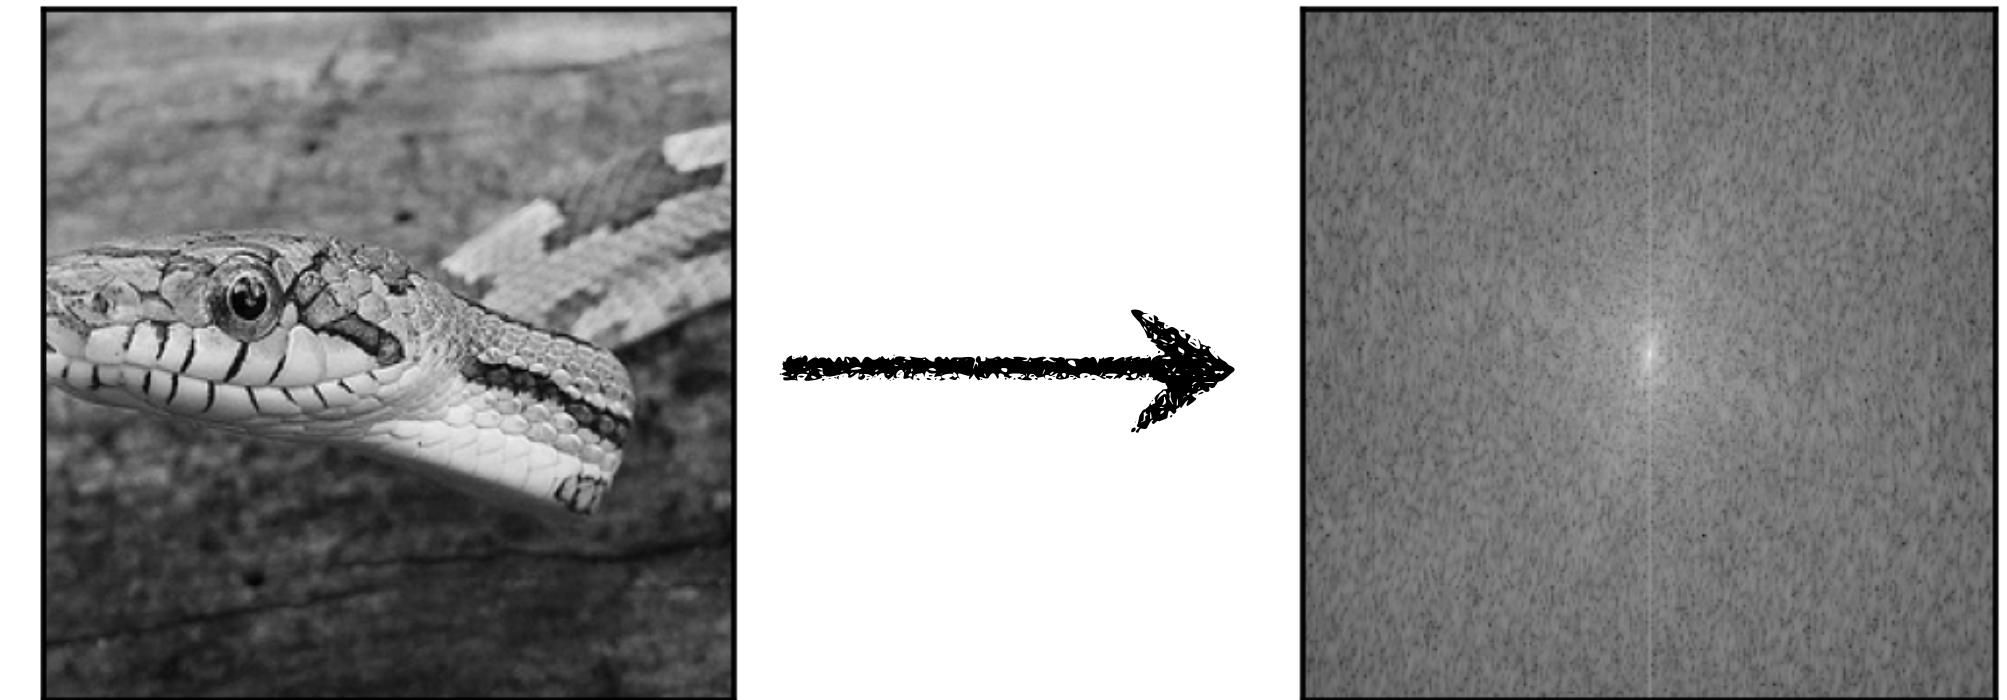


$$X_{tt'}$$

$$\tilde{X}_{kk'} = A_{kk'} \exp(i\phi_{kk'})$$

What is in an image?

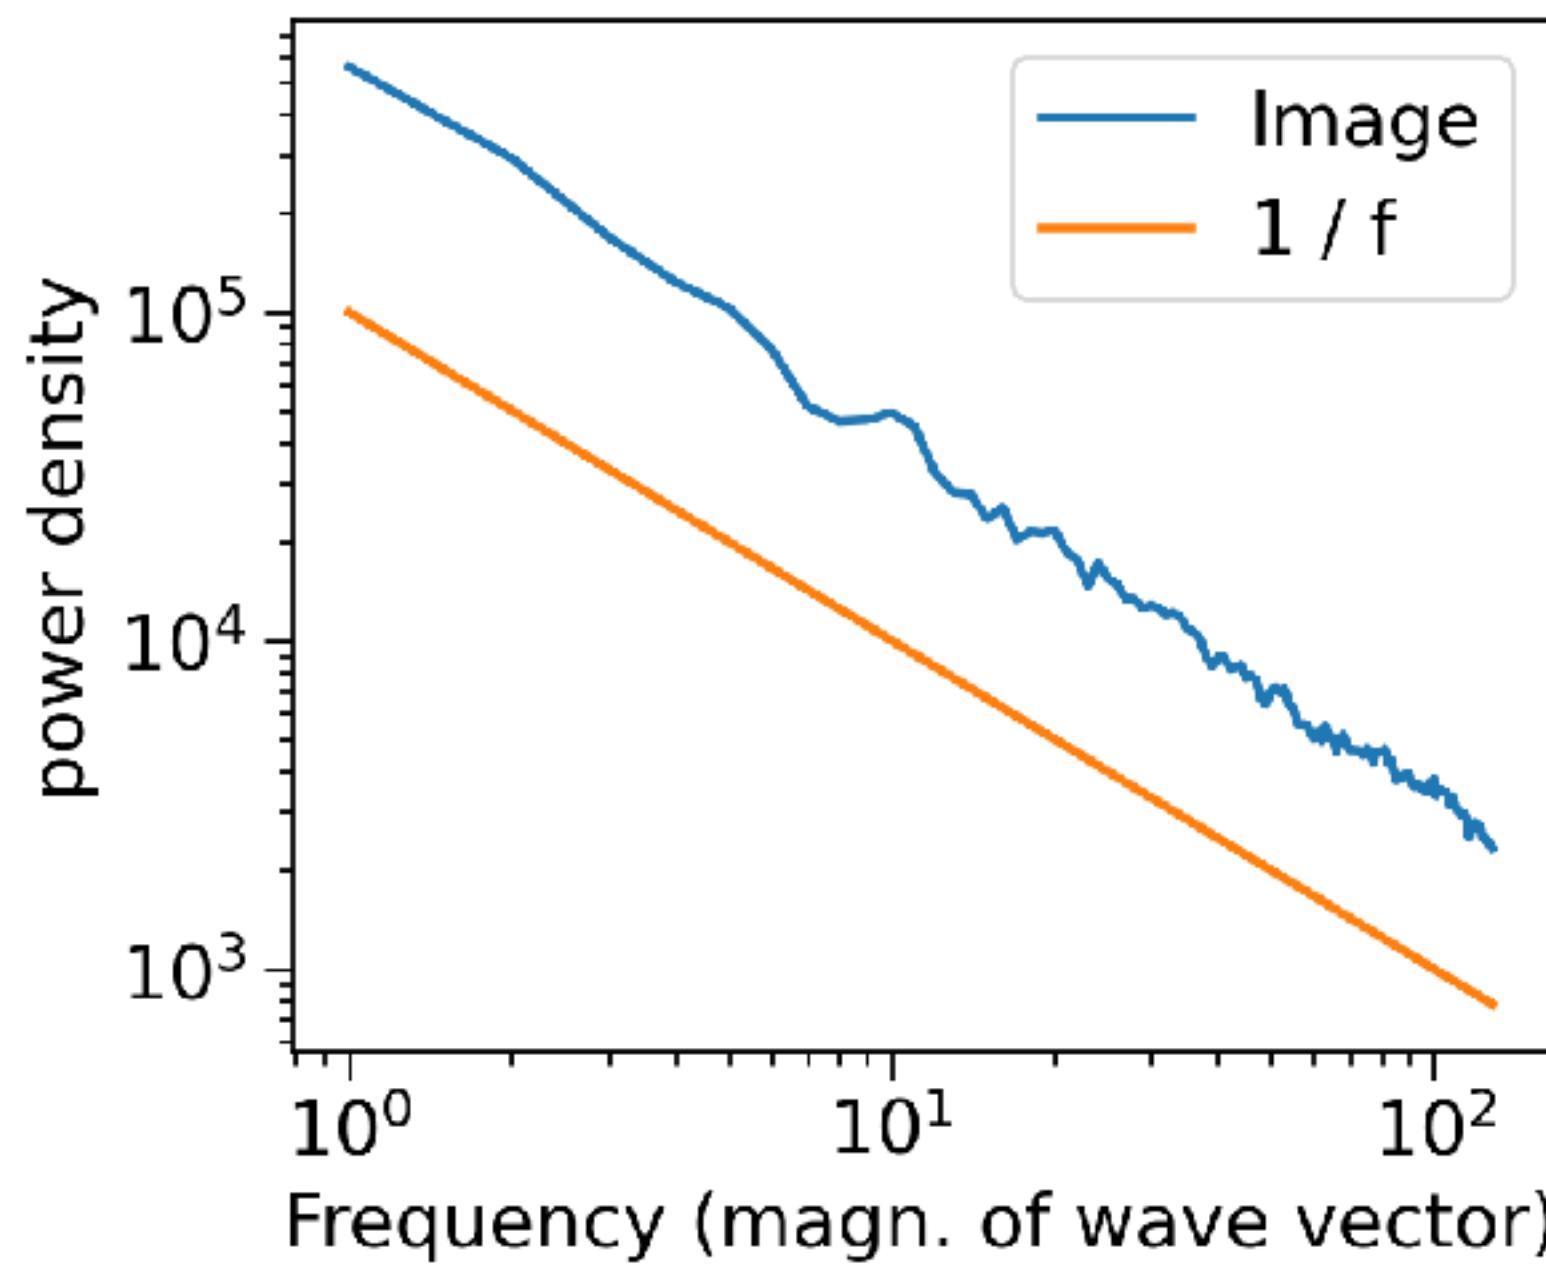
A Fourier perspective



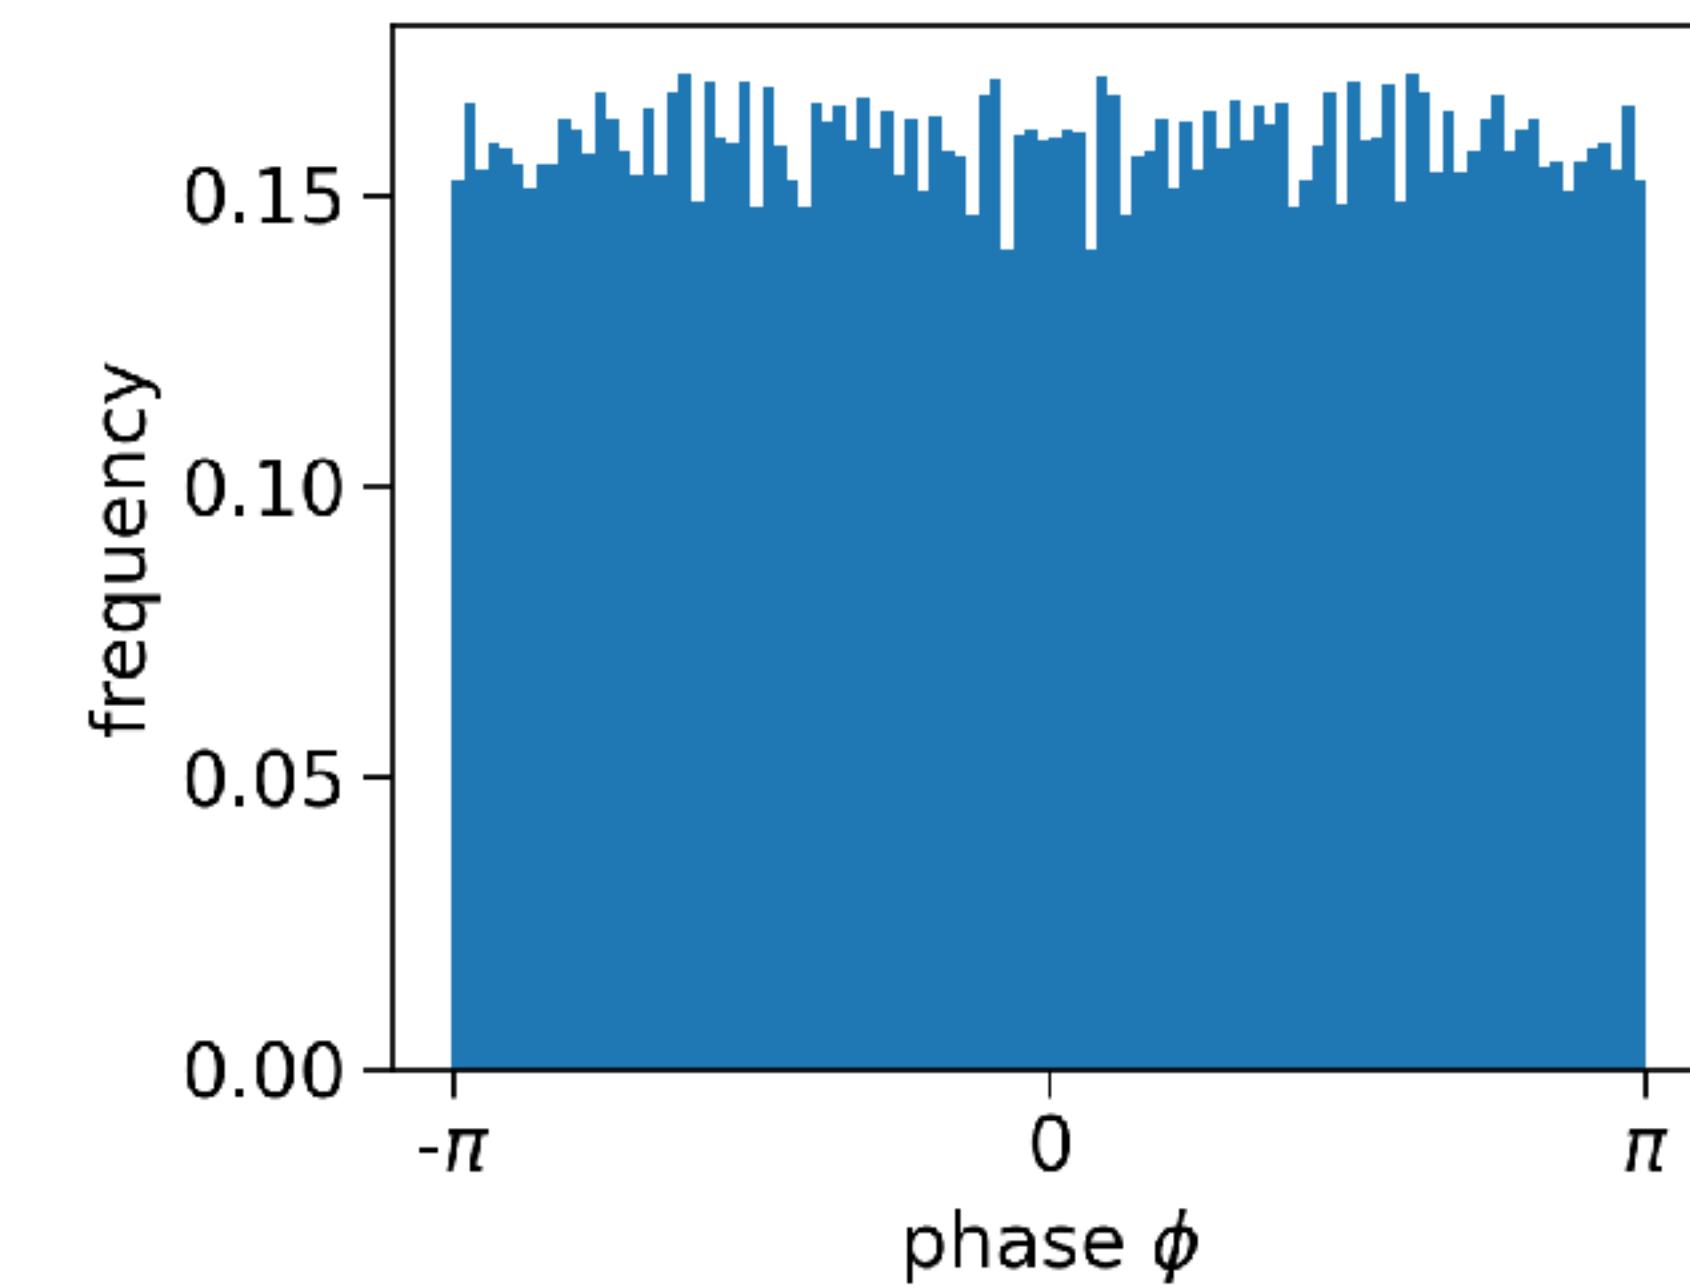
$$X_{tt'}$$

$$\tilde{X}_{kk'} = A_{kk'} \exp(i\phi_{kk'})$$

Amplitudes $A_{kk'}$ encode
pair-wise correlations



Phases $\phi_{kk'}$ encode
higher-order correlations



What matters in an image?

Let's do an experiment to find out! (Piotrowski & Campbell '82)

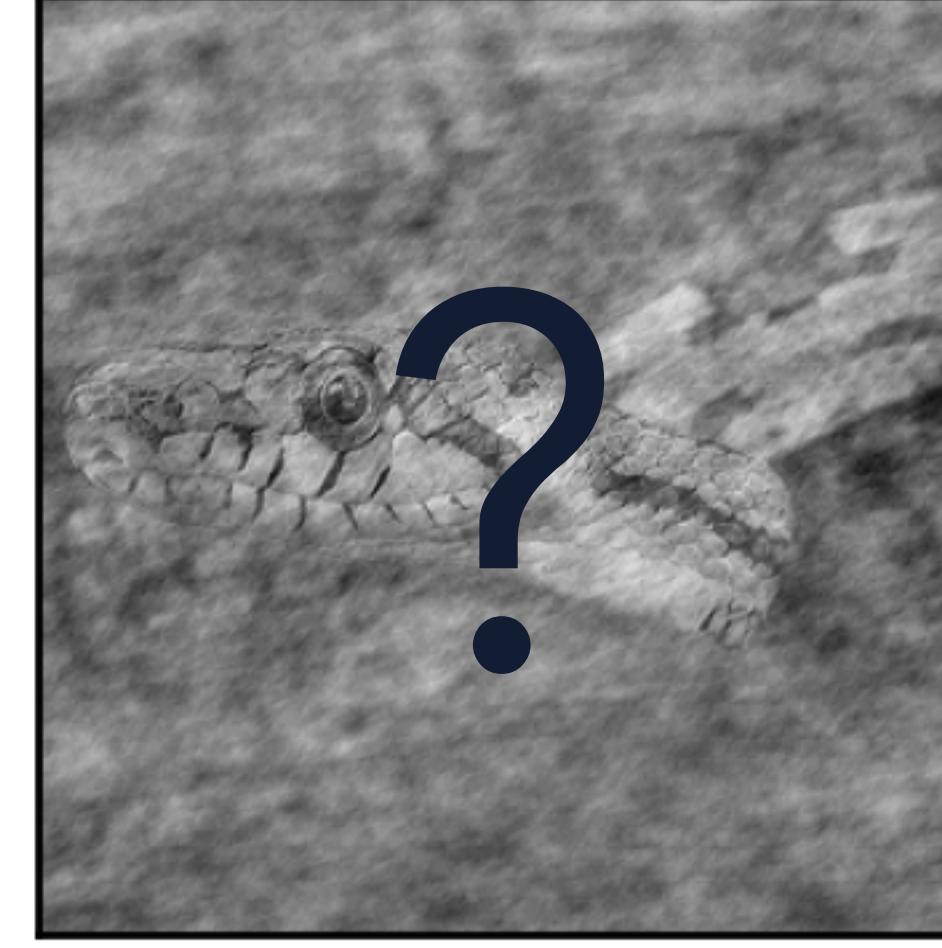
$$\tilde{X}_{kk'} = A_{kk'} \exp(i\phi_{kk'})$$



$$\tilde{X}_{kk'} = A_{kk'} \exp(i\phi_{kk'})$$



$$\tilde{X}_{kk'} = A_{kk'} \exp(i\phi_{kk'})$$



$$\tilde{X}_{kk'} = A_{kk'} \exp(i\phi_{kk'})$$

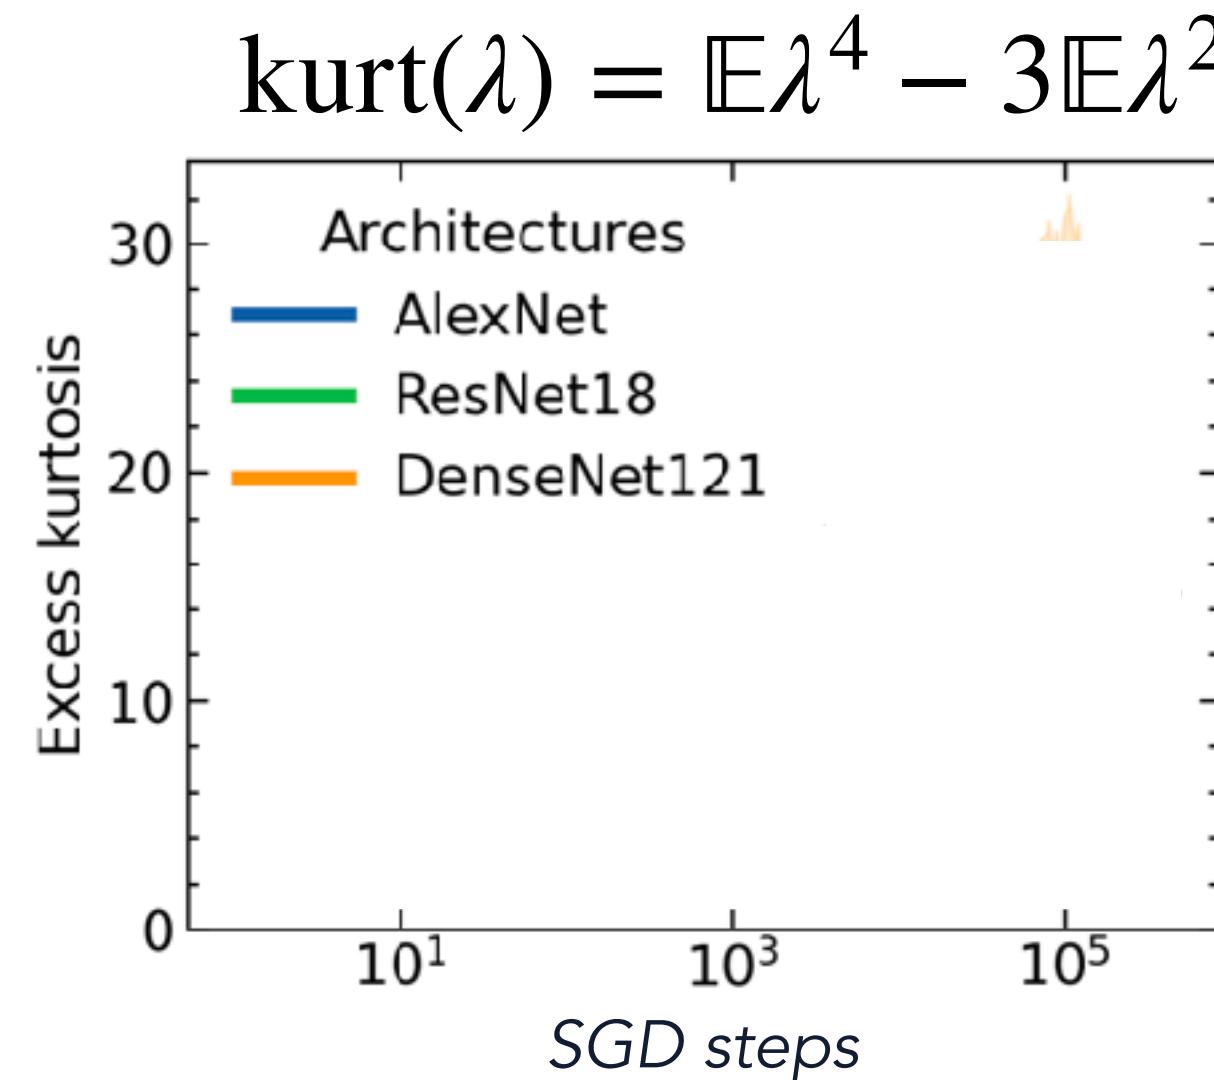
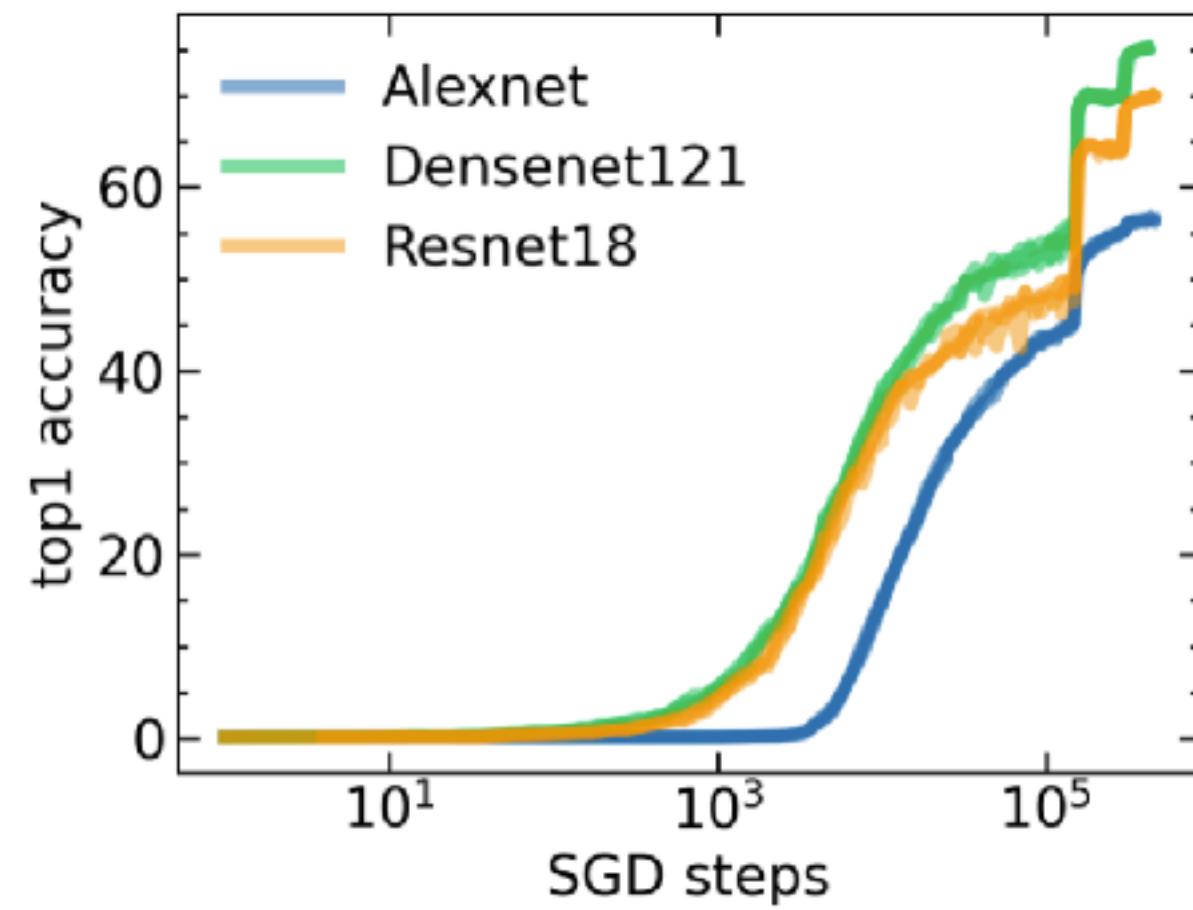


Higher-order correlations are perceptually more important!

Oppenheim & Lim (1981); Piotrowski & Campbell (1982); Wichmann et al. (2005)

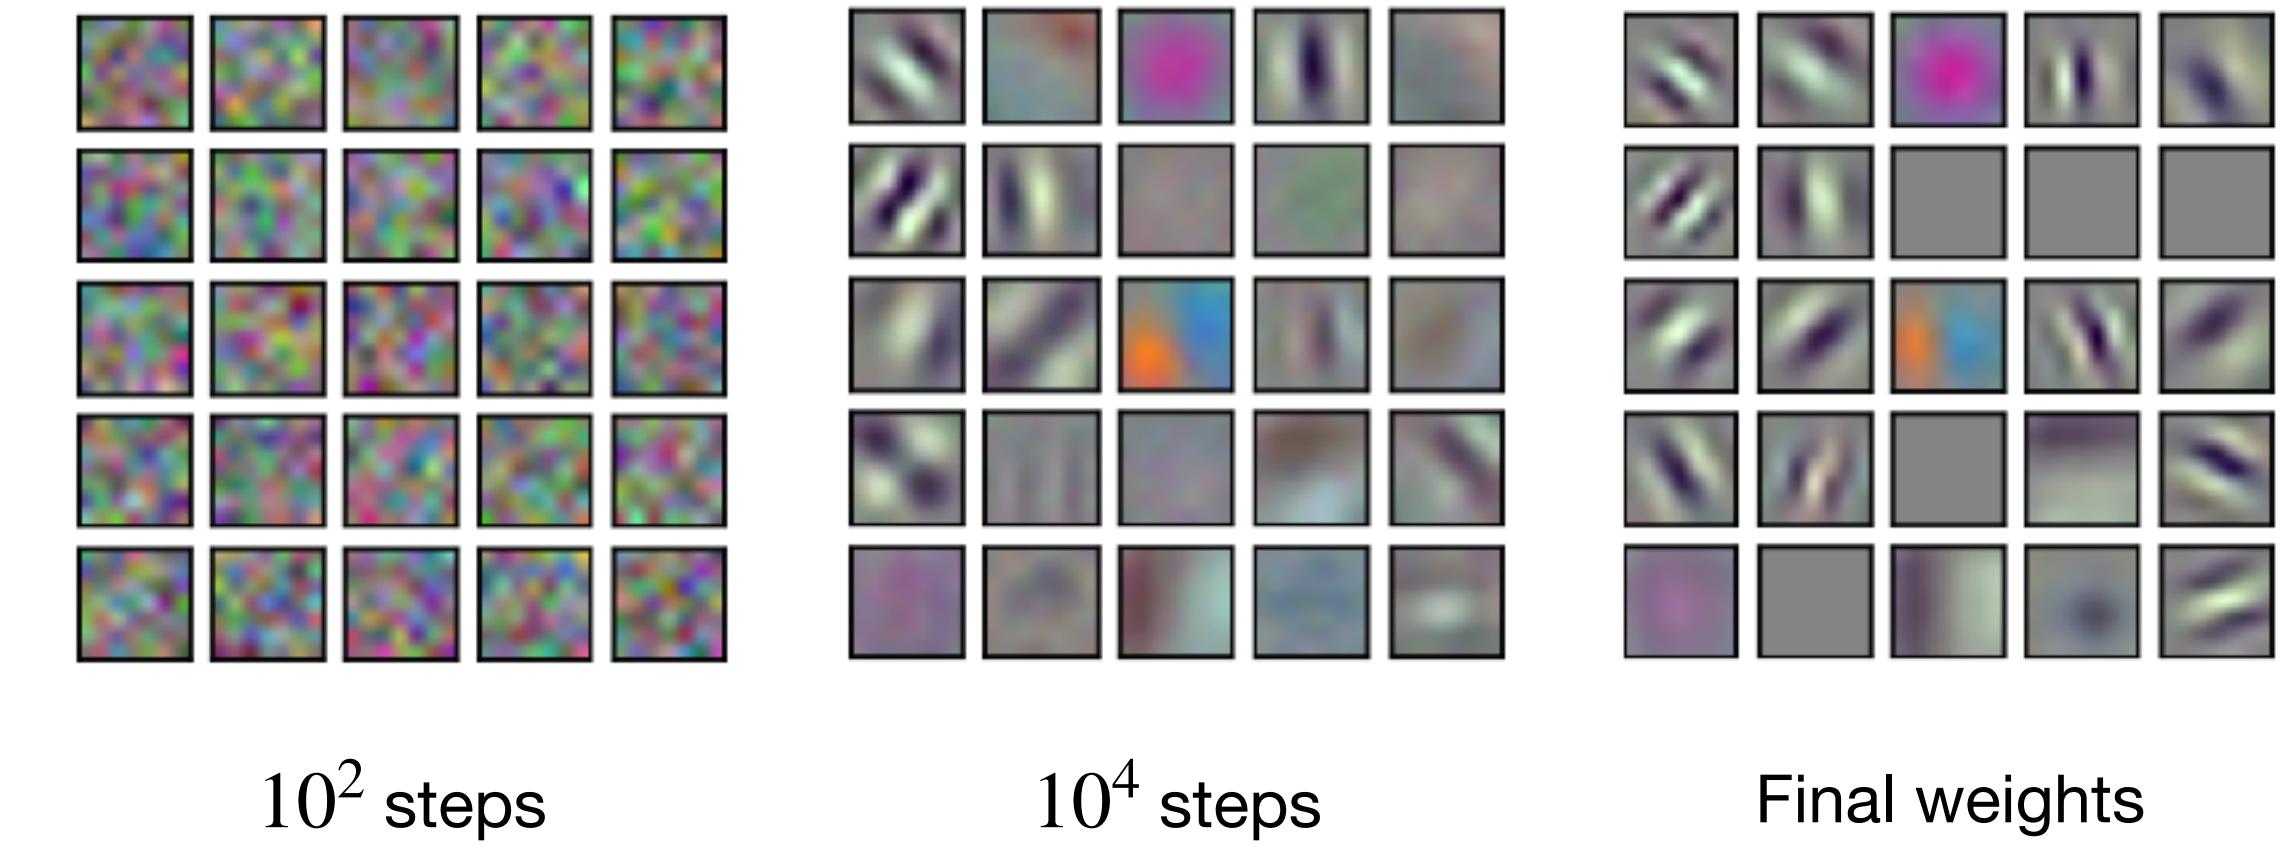
HOCs shape neural representations

First layer filters relate to strongly non-Gaussian directions



$$\lambda = w^{(1)} \cdot x$$

First-layer filters
(Densenet121)



**Neural networks learn features from
non-Gaussian input fluctuations.**

How can we **analyse** this?

A simpler model for learning

Finding “interesting” projections of data

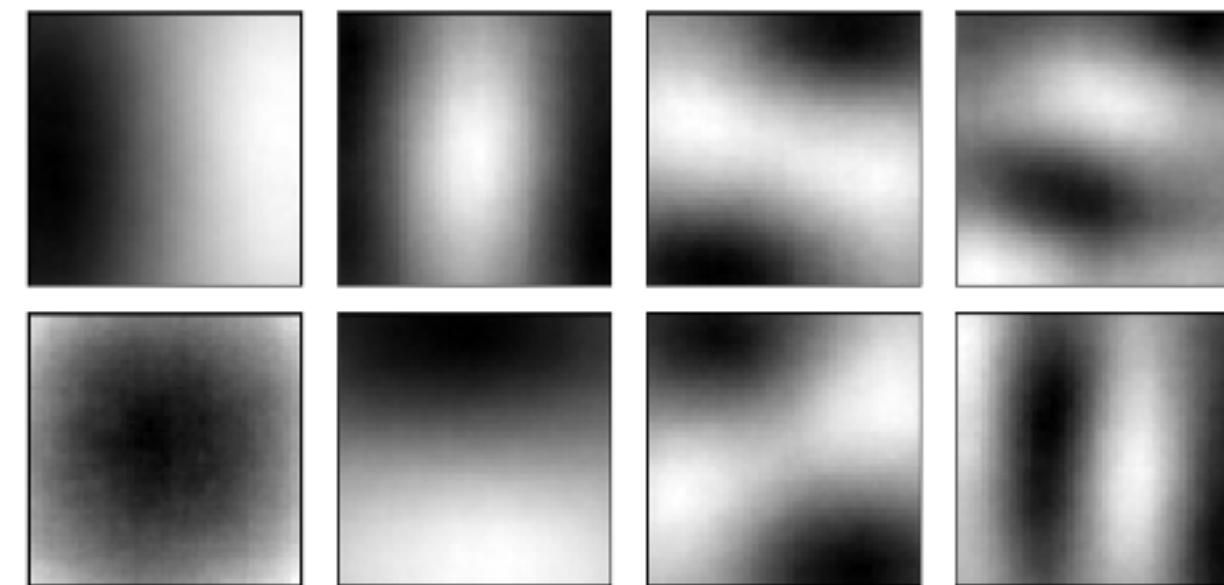
Given a dataset $\mathcal{D} = \{x_1, x_2, \dots, x_n\}$ of d -dimensional, zero-mean inputs **with identity covariance**

$$w^* := \operatorname{argmax}_{|w|=1} \mathbb{E}_{\mathcal{D}} G(w \cdot x)$$

Principal components (PCA)

Pearson 1901

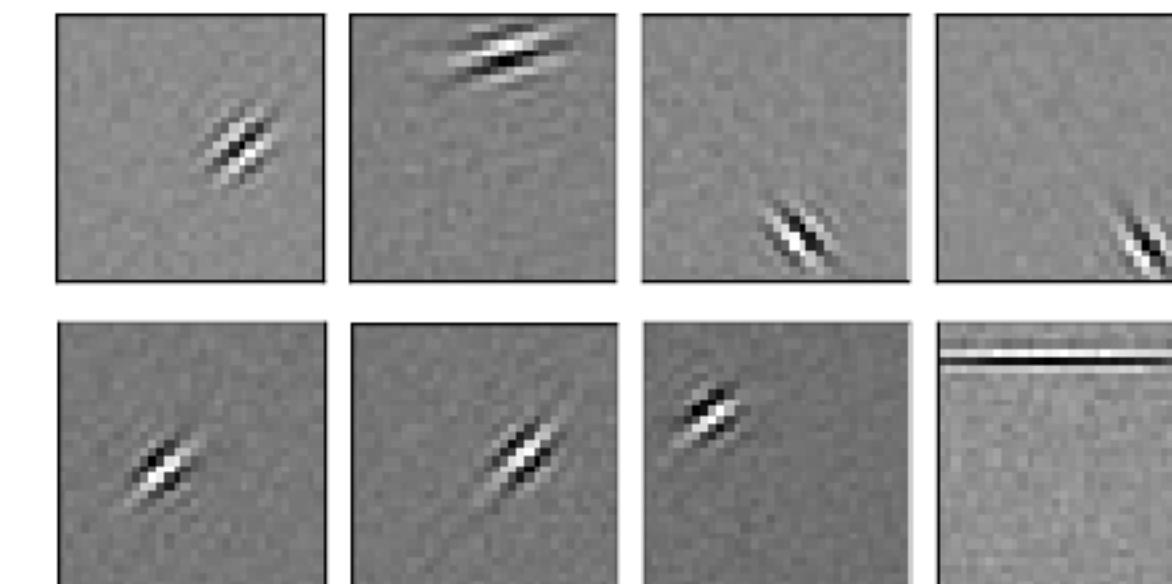
$$G(s) = s^2$$



Translation-invariance of images
=> Fourier components

Independent Components (ICA)

Comon '94; Bell & Sejnowski '95; Oja & Hyvärinen '00



The most **non-Gaussian projections**
yield **CNN-like** filters !

$$G(s) = s^4 e^{-s^2/2}$$

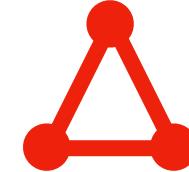
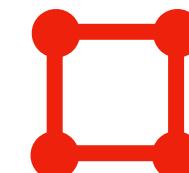
Fundamental limits of ICA

A synthetic data model gives fundamental insights

Spiked cumulant model:

$$x^\mu = \beta g^\mu \mathbf{u} + \mathbf{w}^\mu \quad g^\mu = \pm 1, \quad \mathbf{w}^\mu \sim \mathcal{N}(0, 1 - \beta \mathbf{u} \mathbf{u}^\top)$$

- $\mathbb{E}x_i = 0$ PCA is
- $\mathbb{E}x_i x_j = \delta_{ij}$ **useless!**

 $\mathbb{E}x_i x_j x_k = 0$	Goal of ICA $=$ finding \mathbf{u}!
 $\mathbb{E}x_i x_j x_k x_\ell - \mathbb{E}x_i x_j \mathbb{E}x_k x_\ell [3] \propto (u^{\otimes 4})_{ijkl}$	

How to **analyse** this problem?

$$\mathcal{L}(w) := \mathbb{E}_{\mathbb{P}}[G(w \cdot x)] = \mathbb{E}_{\mathbb{P}_0}[G(w \cdot x)\ell(v \cdot x)]$$

Likelihood ratio $\ell(s) := \frac{d\mathbb{P}}{d\mathbb{P}_0}(s)$

- Algorithmic threshold: $n \gtrsim d^2$
 (Auddy & Yuan '24, Annals Appl Prob '24
 Szekely, Bardone, Gerace & SG, NeurIPS '24)
- **How do algorithms actually perform?**

Feature learning from non-Gaussian inputs: the case of Independent Component Analysis in high dimensions

Fabiola Ricci¹ Lorenzo Bardone¹ Sebastian Goldt¹

ICML 2025
arXiv:2503.23896

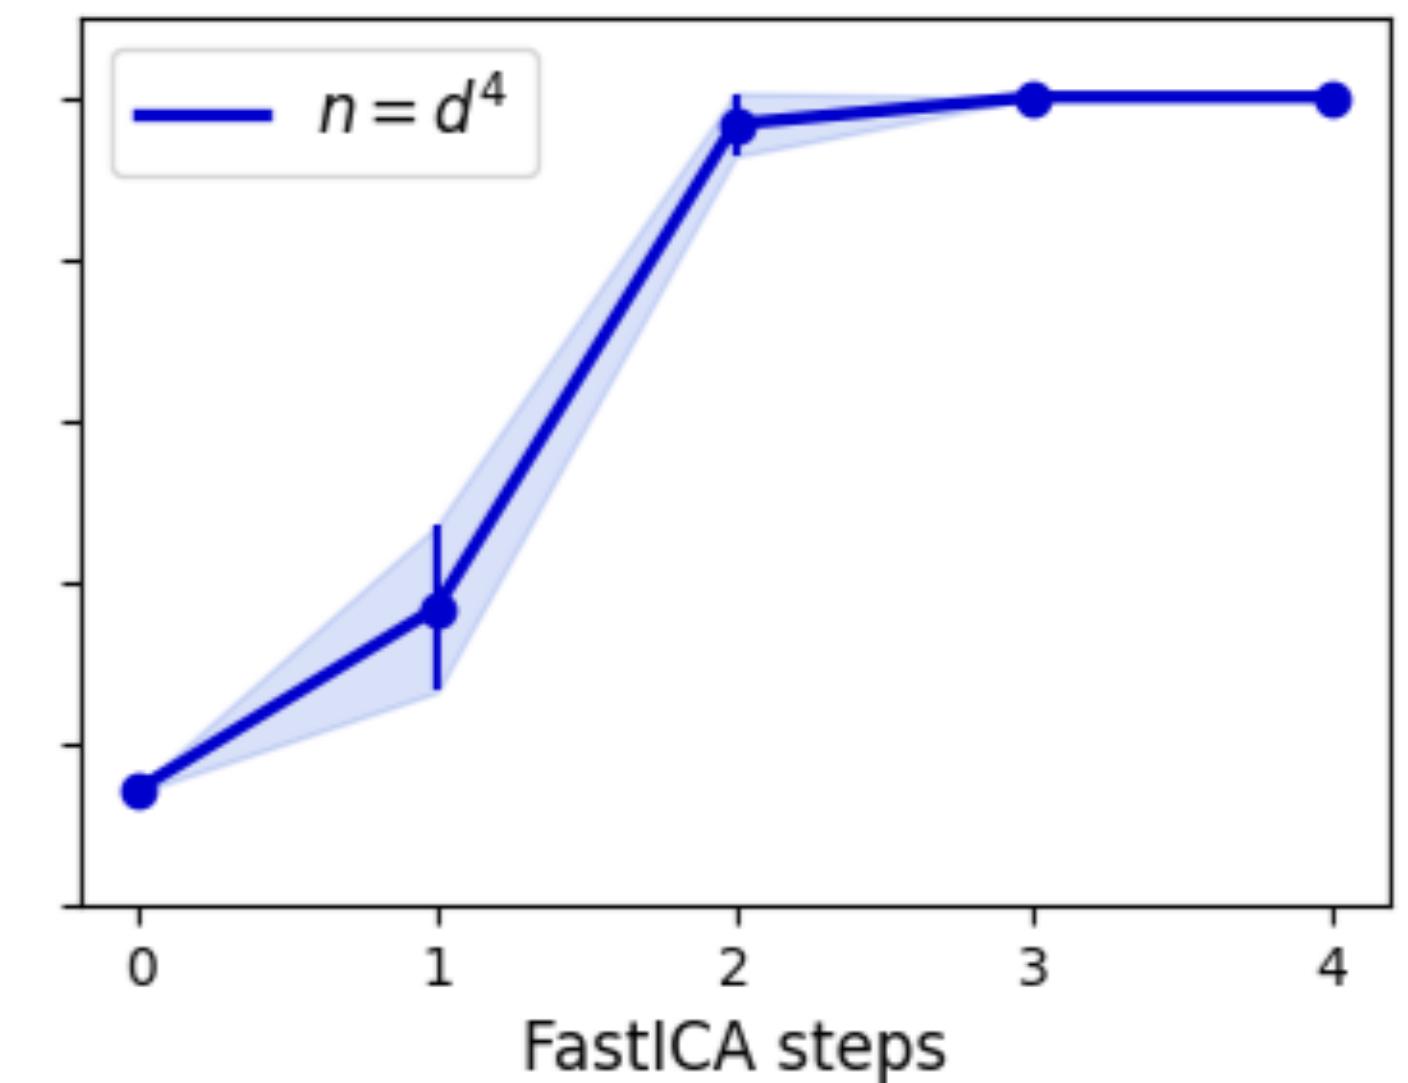
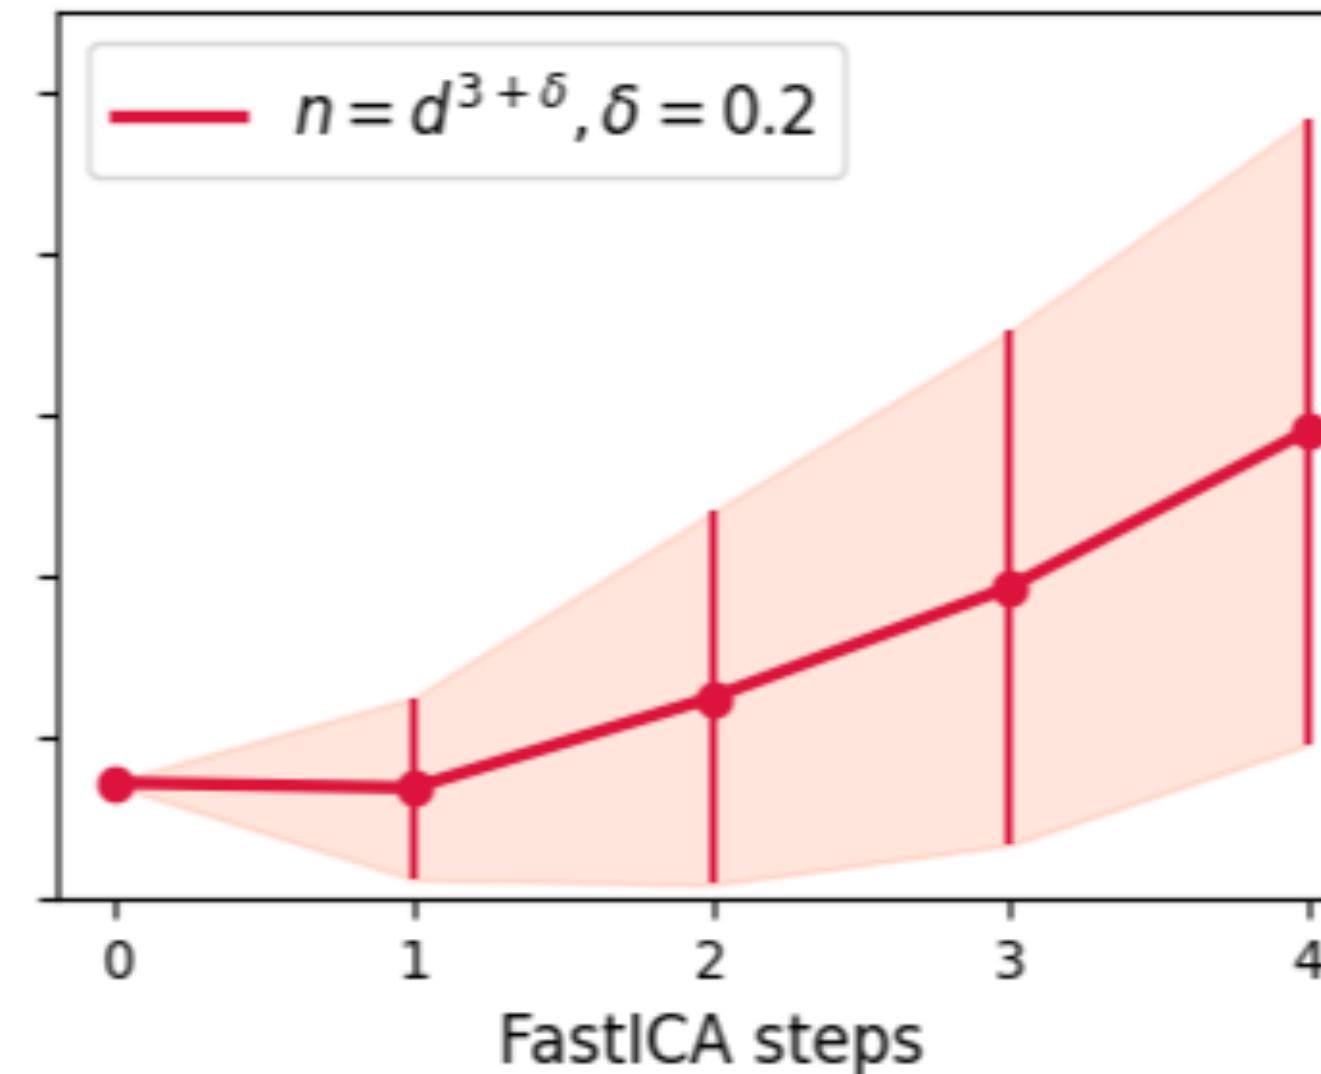
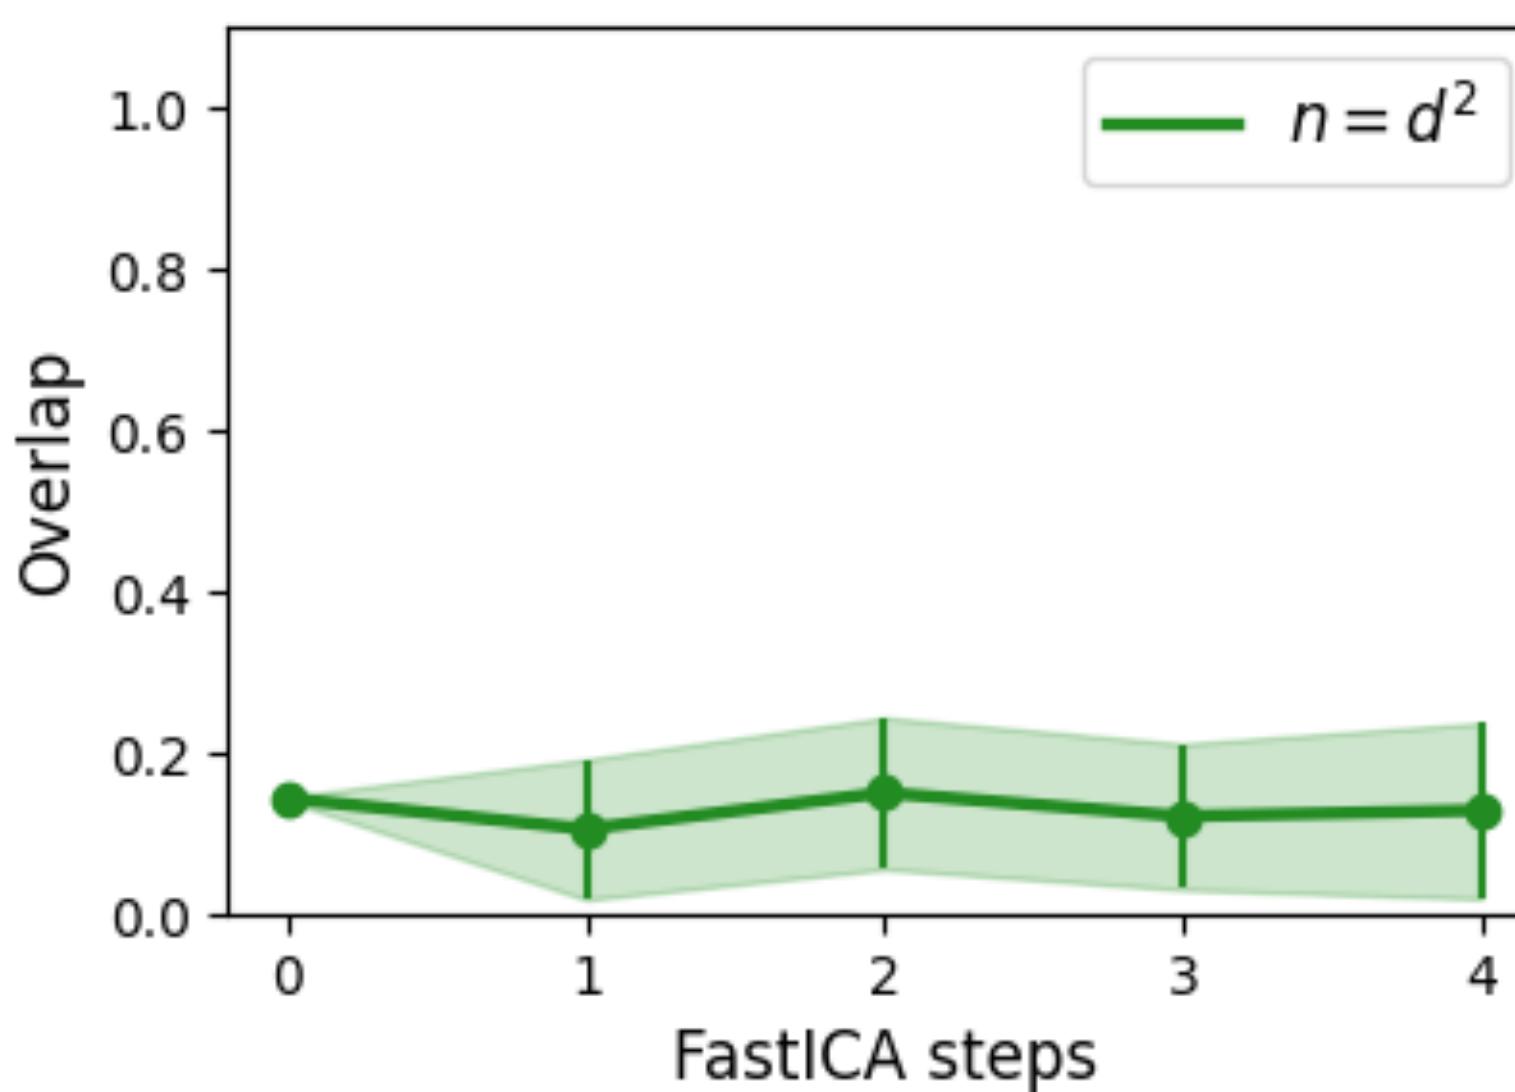


FastICA is slow in high dimensions

The most popular ICA algorithm needs a lot of data

ICA model: $x^\mu = \beta g^\mu \mathbf{u} + \mathbf{w}^\mu \quad g^\mu = \pm 1, \quad \mathbf{w}^\mu \sim \mathcal{N}(0, 1 + \beta \mathbf{u} \mathbf{u}^\top)$

FastICA Algorithm: $\begin{cases} \widetilde{\mathbf{w}}_t &= \mathbb{E}_{\mathcal{D}}[x G'(\mathbf{w}_{t-1} \cdot x)] - \mathbb{E}_{\mathcal{D}}[G''(\mathbf{w}_{t-1} \cdot x)] \mathbf{w}_{t-1}, \\ \mathbf{w}_t &= \widetilde{\mathbf{w}}_t / \|\widetilde{\mathbf{w}}_t\|. \end{cases}$ $G(s) := -e^{-s^2/2}$
 $G(s) := 1/a \log \cosh(as)$



FastICA is slow in high dimensions

The most popular ICA algorithm needs a lot of data

FastICA as a full-batch fixed point iteration:
analyse in the **giant steps framework!**

(Ba et al. '22; Damian et al. '24; Dandi et al. '24; Ben Arous et al. '21)

Theorem (informal).

Take $n = d^\vartheta$ samples. After one step of
FastICA, the overlap α scales as

$$\vartheta \leq 3$$

$$\alpha^2 = O\left(\frac{1}{d}\right)$$

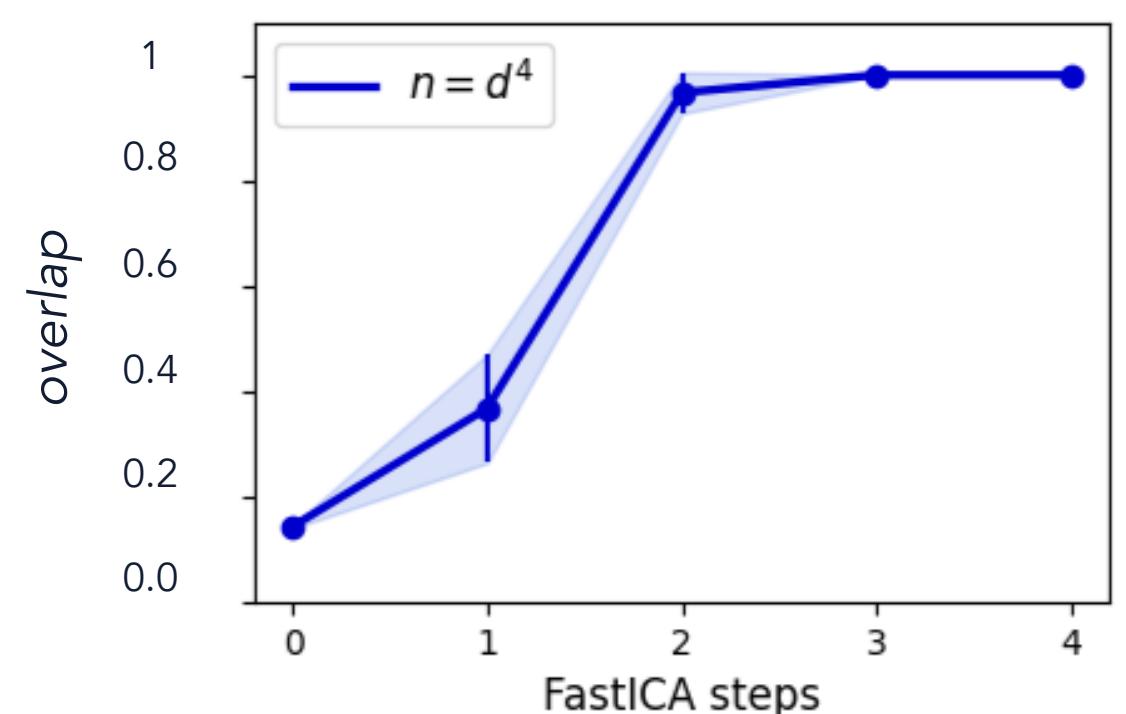
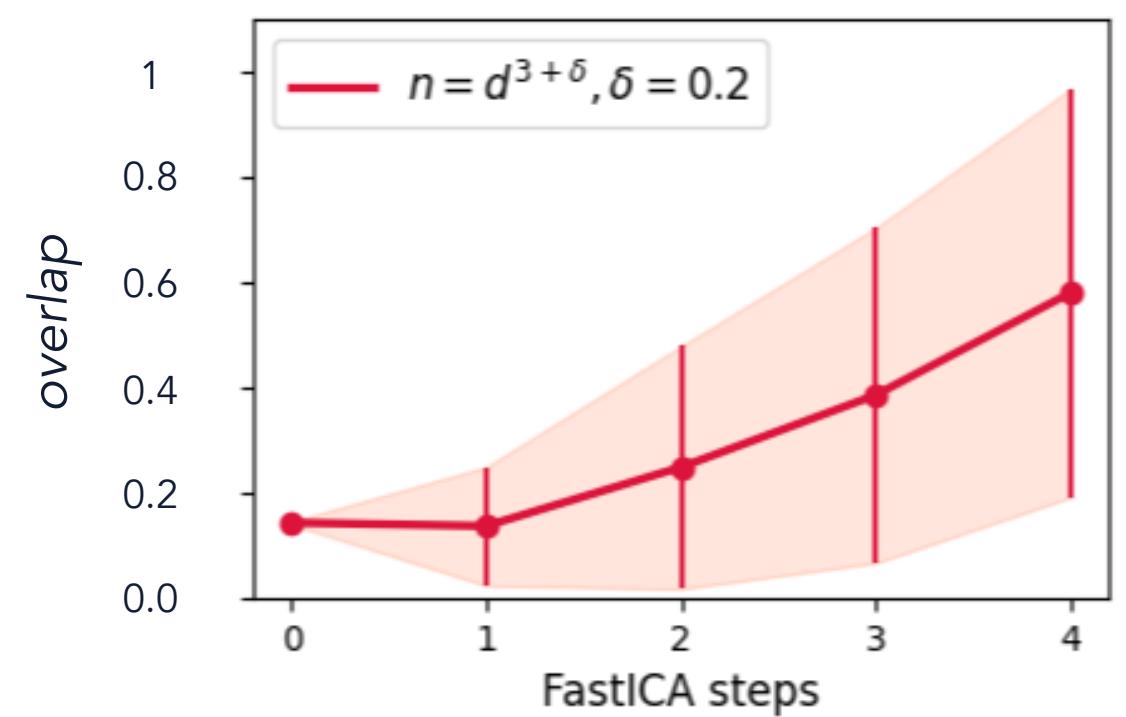
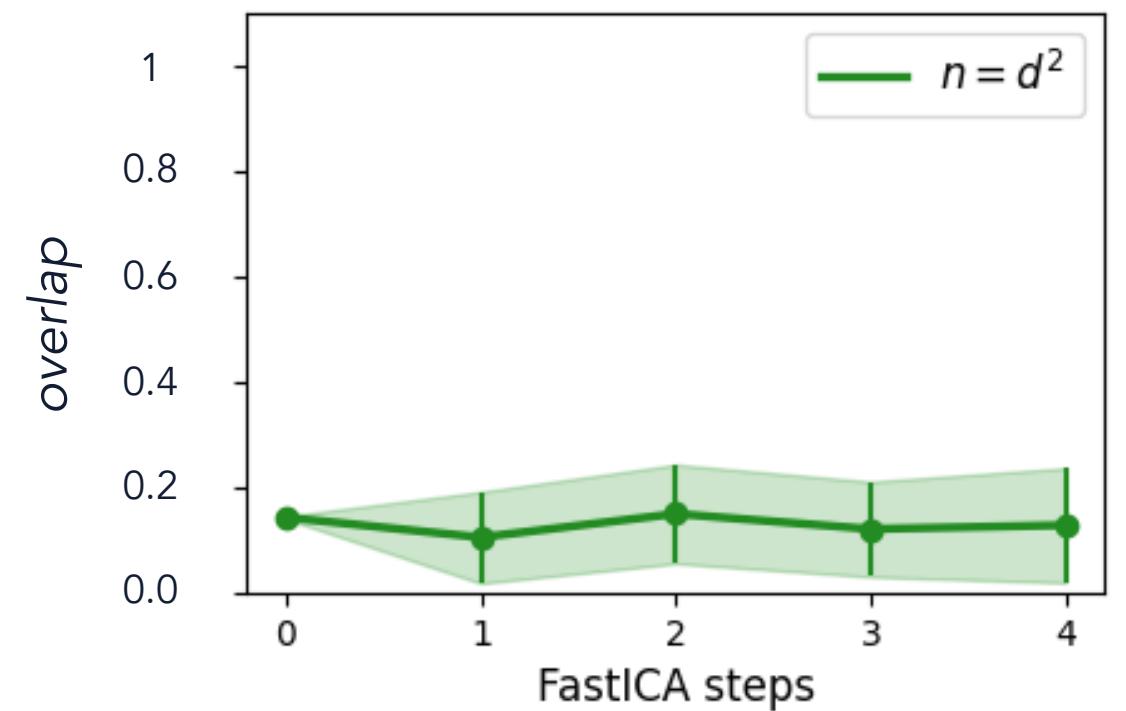
$$3 < \vartheta < 4$$



$$\alpha^2 = o(1)$$

$$4 \leq \vartheta$$

$$\alpha^2 = 1 - o(1)$$



Speeding up ICA with SGD

Smoothing the landscape is the key!

Vanilla SGD (Ben Arous et al. JMLR '21):

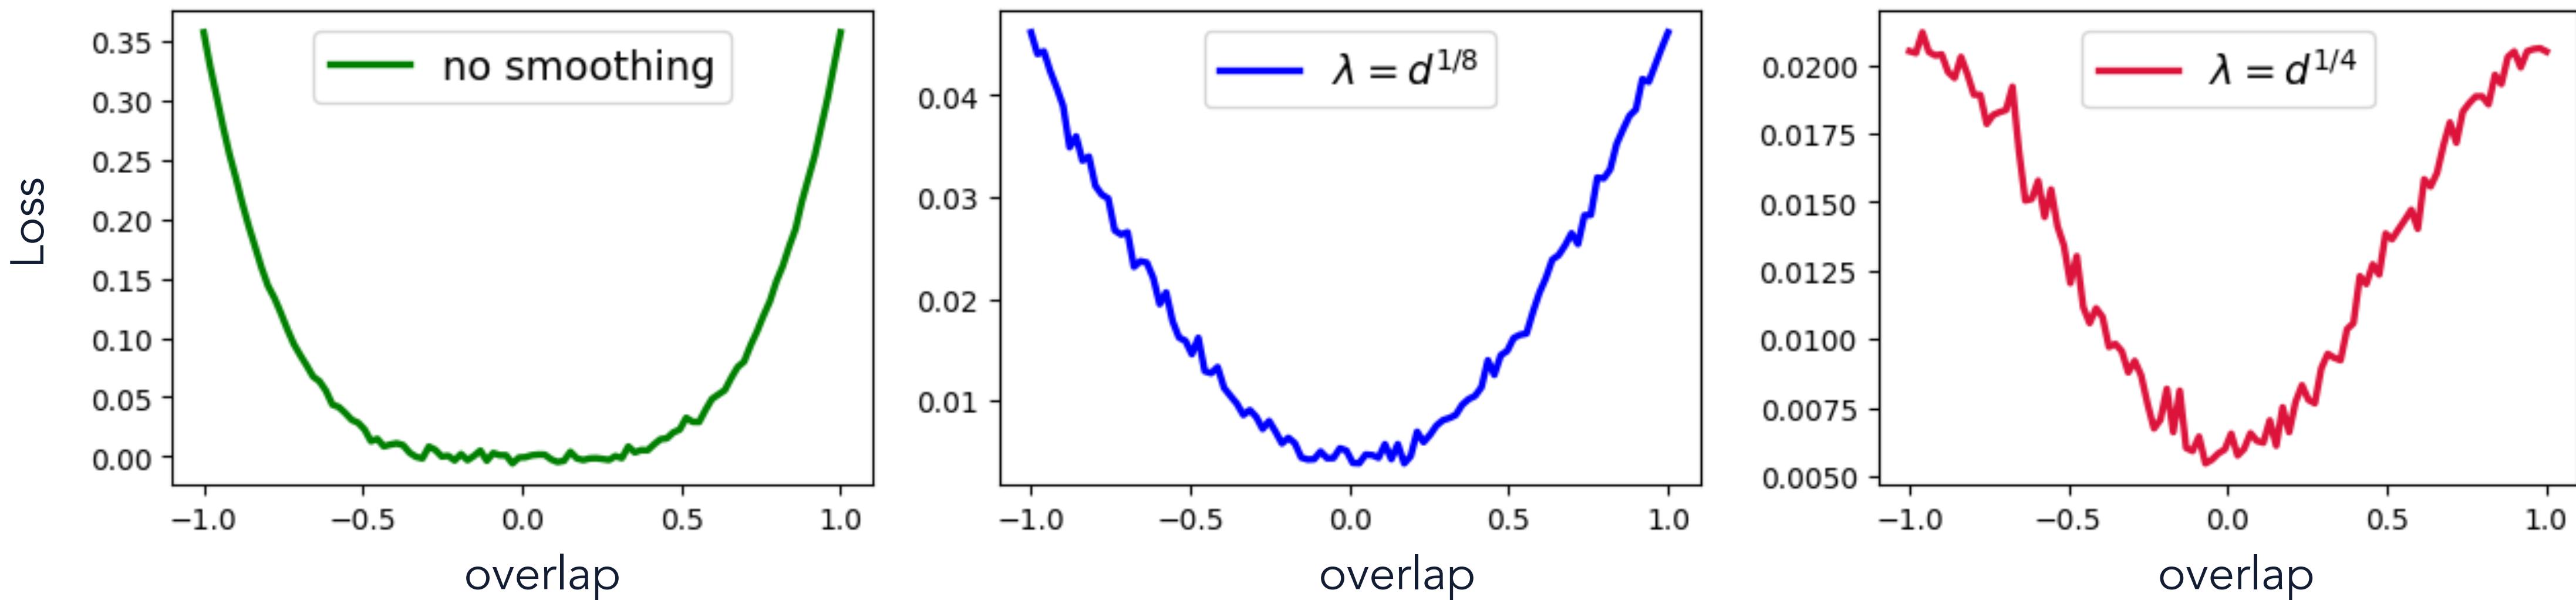
recovers the spike in $n = \Omega(d^3 \log^2 d)$ steps.

SGD on a **smoothed loss**

Biroli, Cammarota, Ricci-Tersenghi J Phys A '20

Damian et al. NeurIPS '23

$$\mathcal{L}_\lambda[G(w \cdot x)] := \mathbb{E}_{z \sim \mu_w} G\left(\frac{w + \lambda z}{\|w + \lambda z\|} \cdot x\right) \quad \lambda \geq 0$$



Speeding up ICA with SGD

Smoothing the landscape is the key!

Vanilla SGD (Ben Arous et al. JMLR '21): recovers the spike in $n = \Omega(d^3 \log^2 d)$ steps.

SGD on a **smoothed loss**

Biroli, Cammarota, Ricci-Tersenghi J Phys A '20

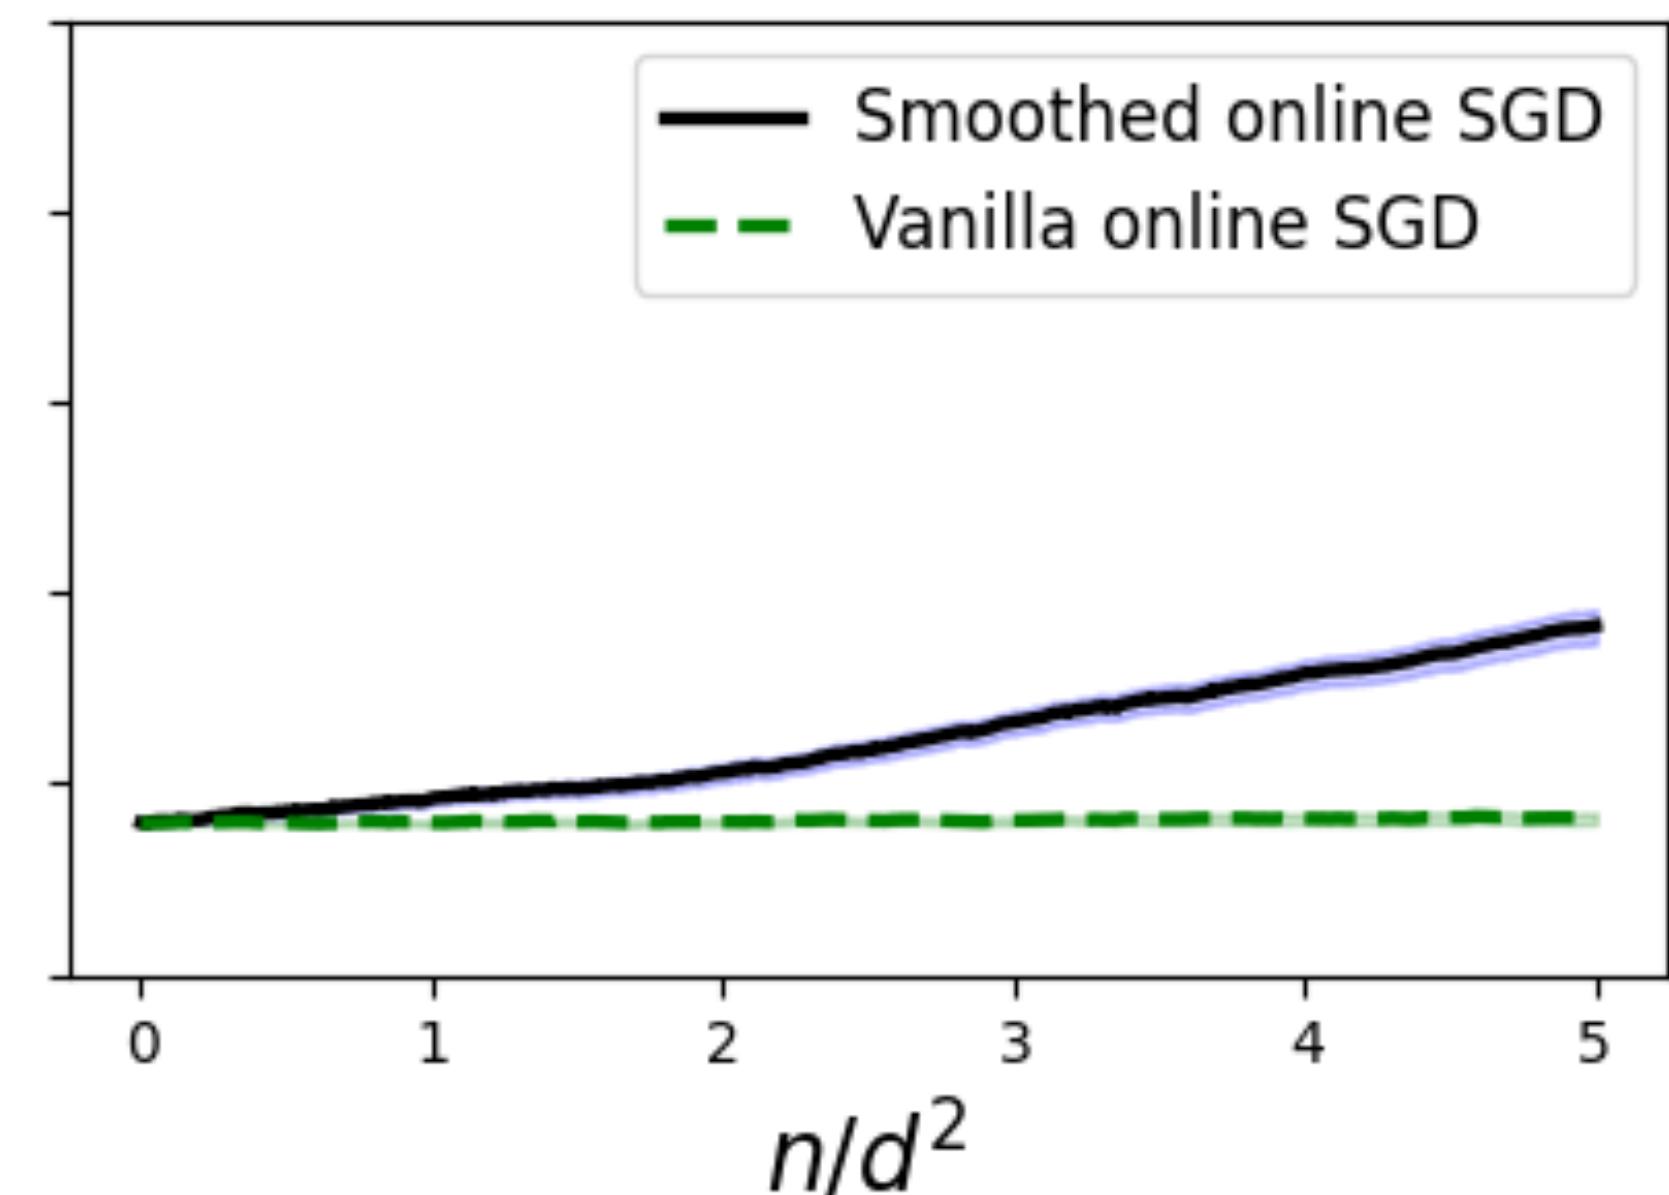
Damian et al. NeurIPS '23

Generalised ODE for the overlap:
(accounting for **data/contrast fn mismatch**)

$$m'(t) = \frac{m(t)}{d^{k_1^* - k_2^*/2}}$$

- Speed-up requires fine-tuning of “activation” function!
- Optimal choice for spiked cumulant is $\text{He}_4(s)$.
Matches LDLR bound!
- **Trade-off:** stability vs. speed!

$$\mathcal{L}_\lambda[G(w \cdot x)] := \mathbb{E}_{z \sim \mu_w} G\left(\frac{w + \lambda z}{\|w + \lambda z\|} \cdot x\right) \quad \lambda \geq 0$$

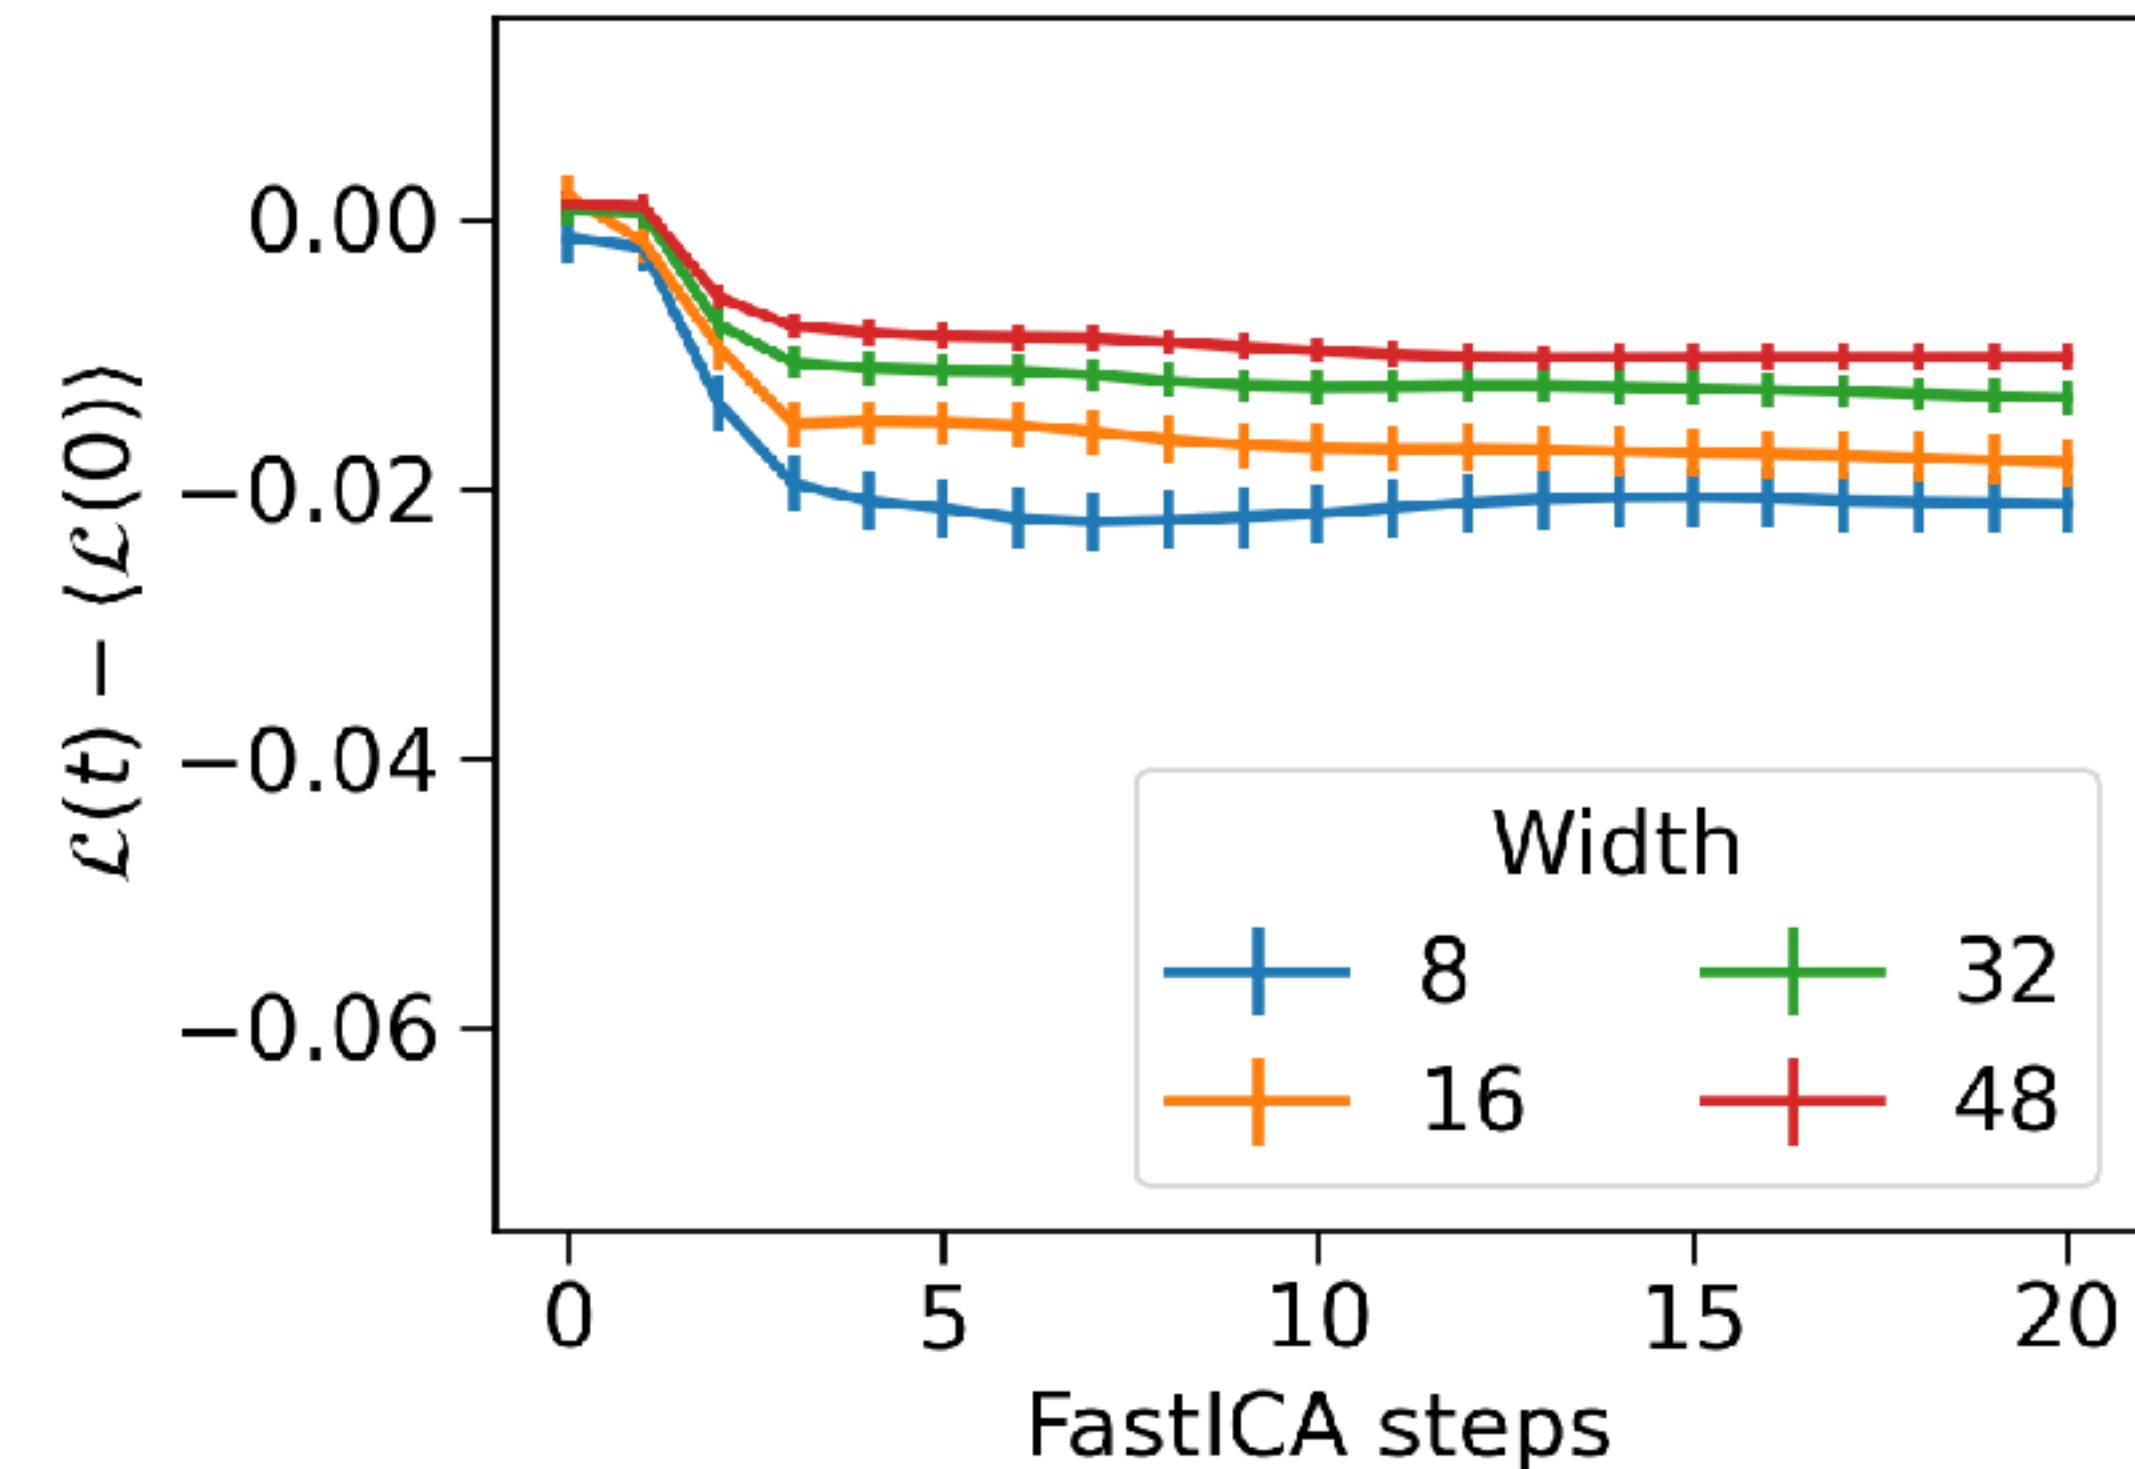


ICA is a **hard** problem in high dimensions.

So what happens on **real images** with
deep neural networks?

What about real data?

FastICA fails on real images at linear sample complexity



FastICA, logcosh activation,
 $n = 2D$, $d=D$ (left) vs. $d=32$ (right)

Reduce and conquer

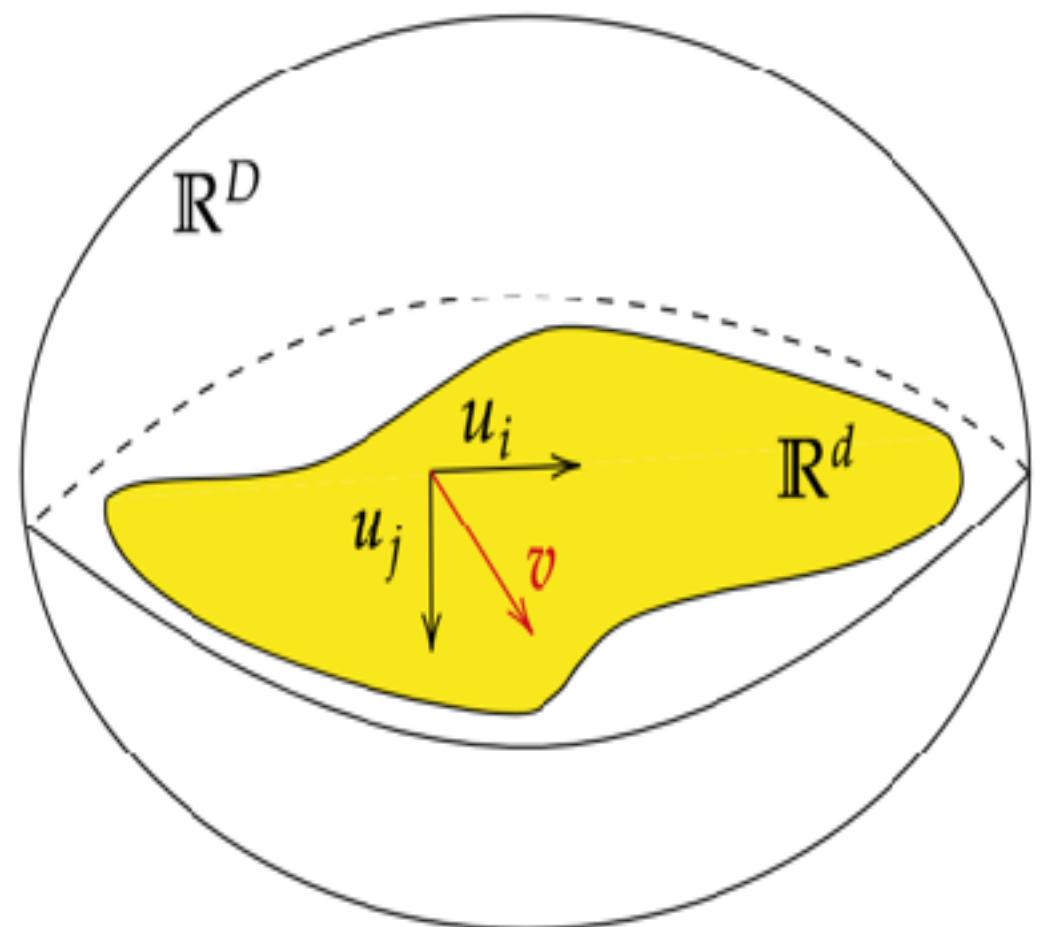
Reduce the dimension, conquer with ICA



Project inputs

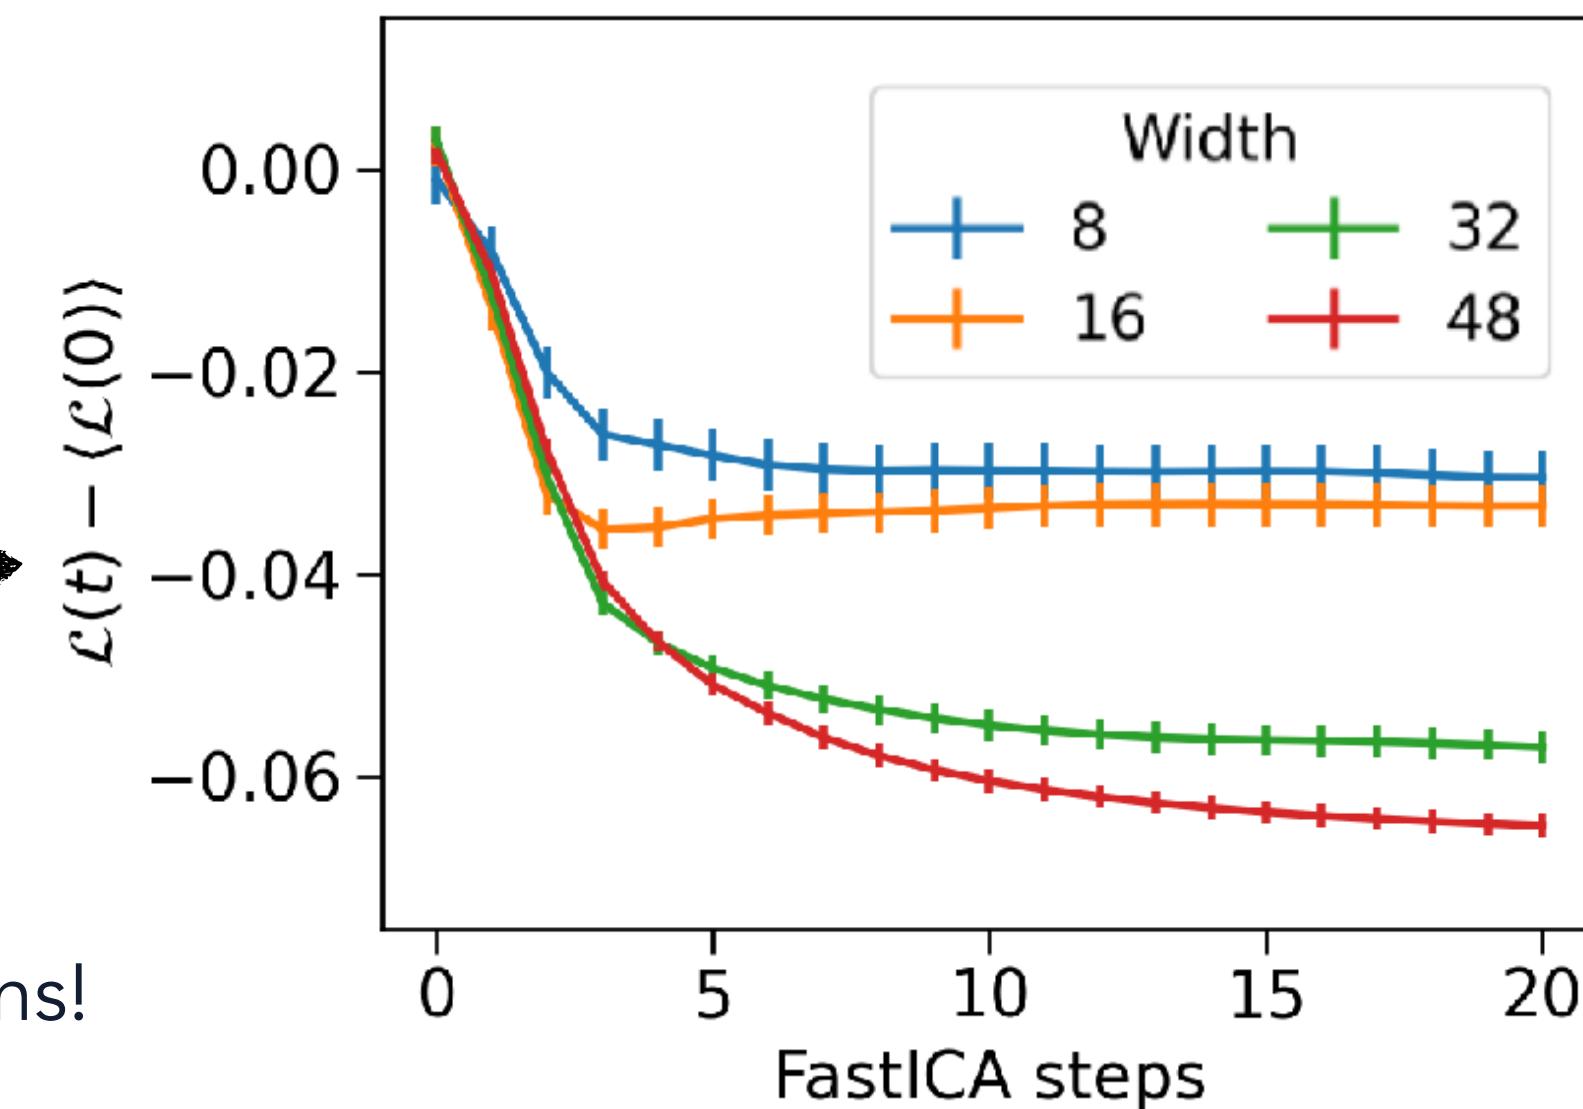
to principal
subspace

Hyvärinen '99



Run **FastICA** in
low dimensions

Loss in
high dimensions!



- Success reveals something about the **structure of the images**
- Linear Sample complexity can be proven in "**subspace model**"

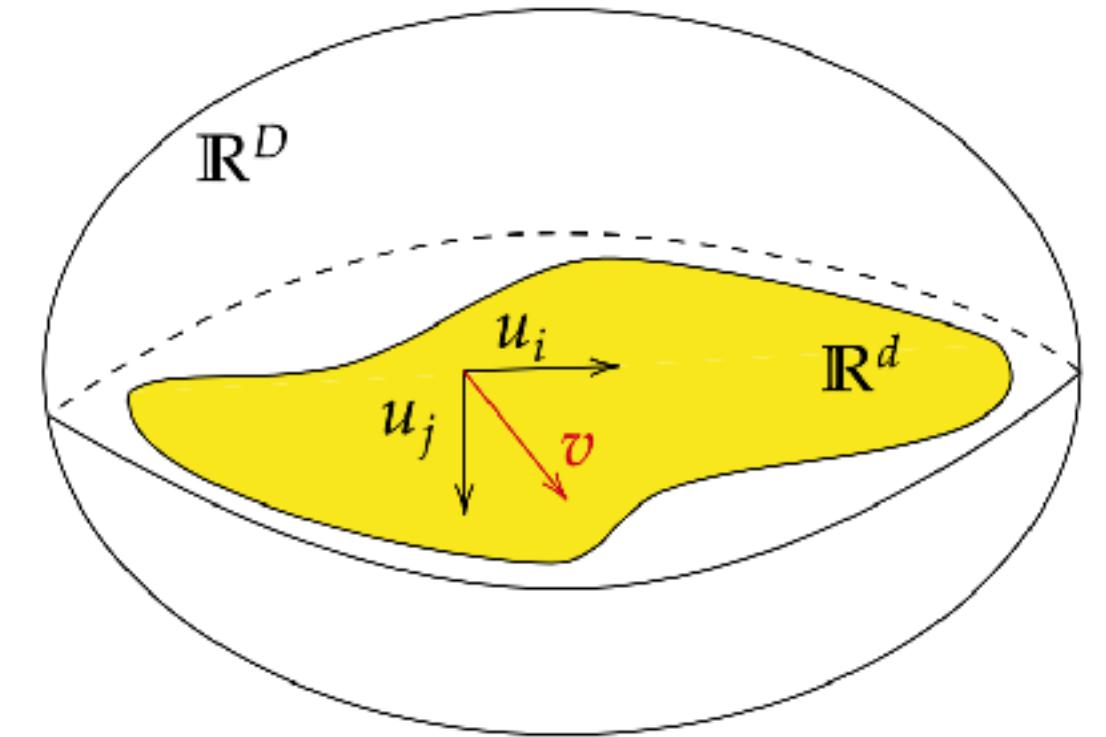
What about real data?

A mixed matrix-tensor model

Subspace model (rank-1 in Bardone & SG, ICML '24)

$$\mathbf{x}^\mu = \sum_r \beta_1 g_r^\mu \mathbf{u}_r + \beta_2 h^\mu \mathbf{v} + \mathbf{w}^\mu$$

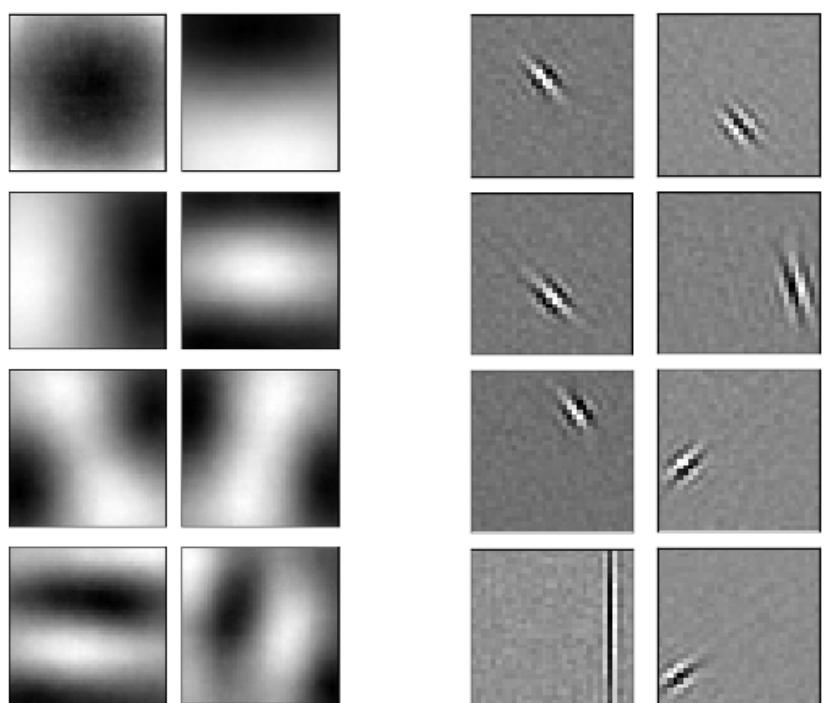
$$g_r^\mu \sim \mathcal{N}(0,1), \quad h^\mu = \pm 1 \\ \mathbf{w}^\mu \sim \mathcal{N}(0,1 - \beta_2 \mathbf{v}\mathbf{v}^\top)$$



Prove recovery by analysing GD in finite-dimensional sub-space spanned by PCs

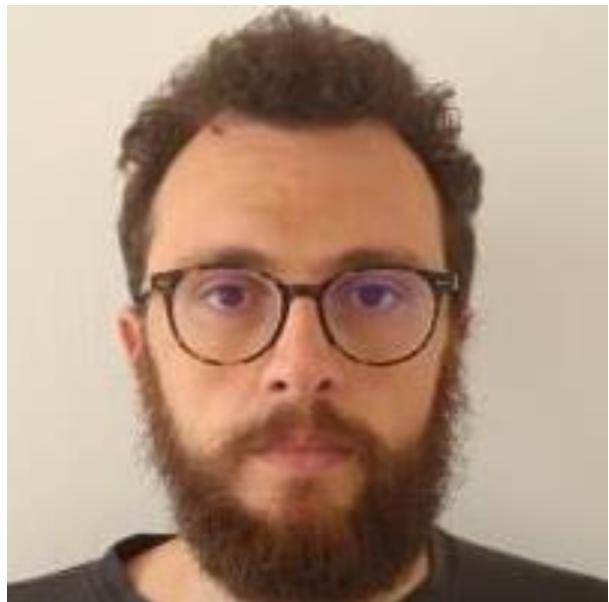
Mixed matrix-tensor models

- Richard & Montanari (NeurIPS '14). observe $X = \beta v^{\otimes p} + Z$ and $y = \beta x + z$, $\beta > 0$
- Sarao Mannelli et al. ('19a, '19b, '20) observe $M \propto vv^\top + Z_M$ and $T \propto v^{\otimes p} + Z_T$
- Asymmetric case: Tabanelli et al. arXiv:2506.02664



What about (shallow) neural networks?

Learning distributions of increasing complexity



A. Ingrosso



vs.



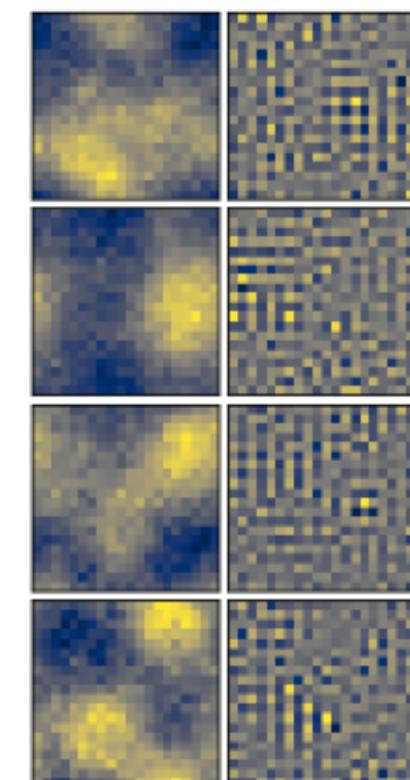
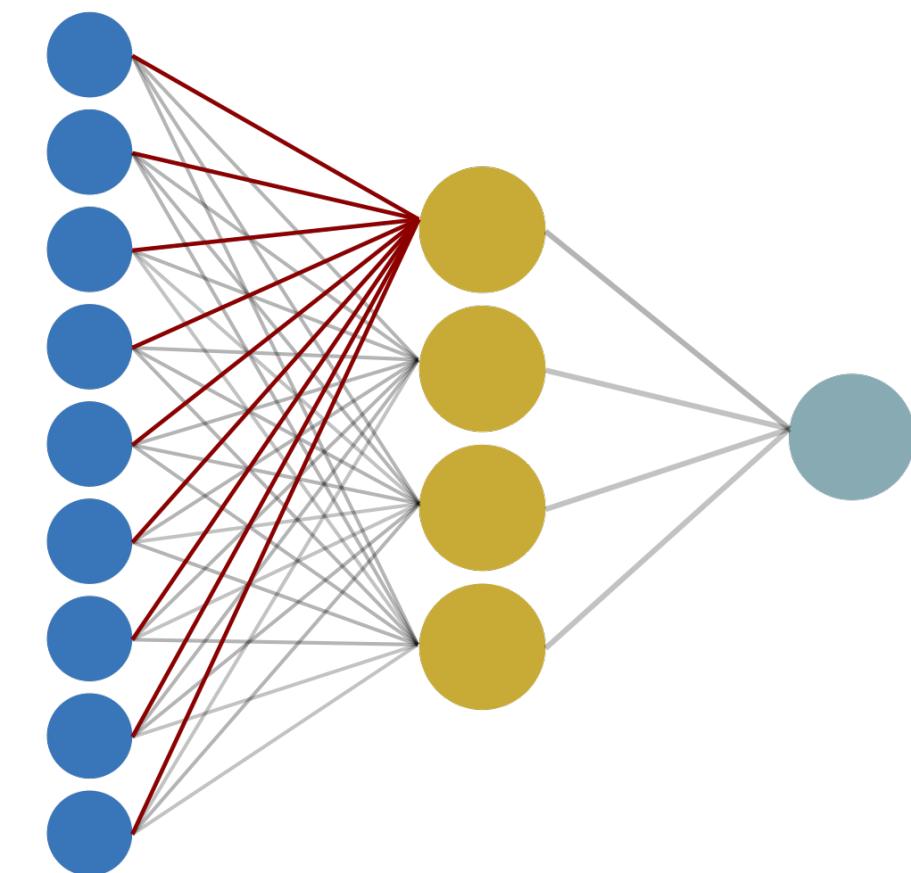
Translation-invariance:

$$\mathbb{E} z_k^\pm z_l^\pm = \exp(-|k-l|/\xi^\pm)$$

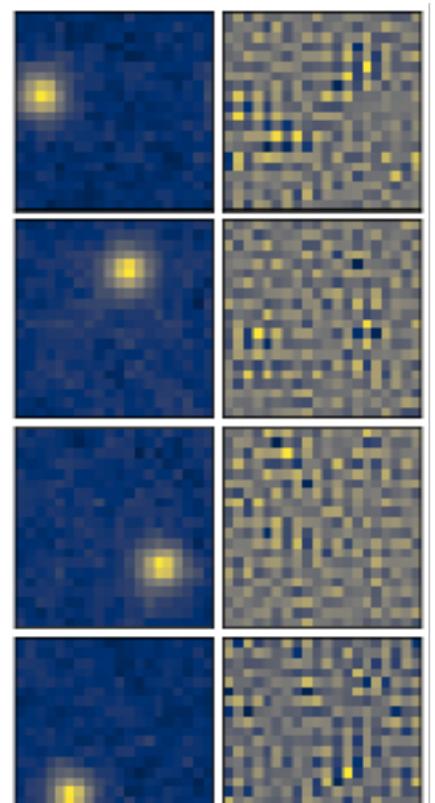
Sharp edges:

(from saturating non-linearity)

$$x_j^\pm \propto \text{erf}\left(gz_j^\pm\right)$$



$t \approx 10^1$



$t \approx 10^4$

- **Early** in training: neurons \approx Fourier modes, doing **PCA**
- **Later** in training, neurons become localised, doing \approx **ICA**

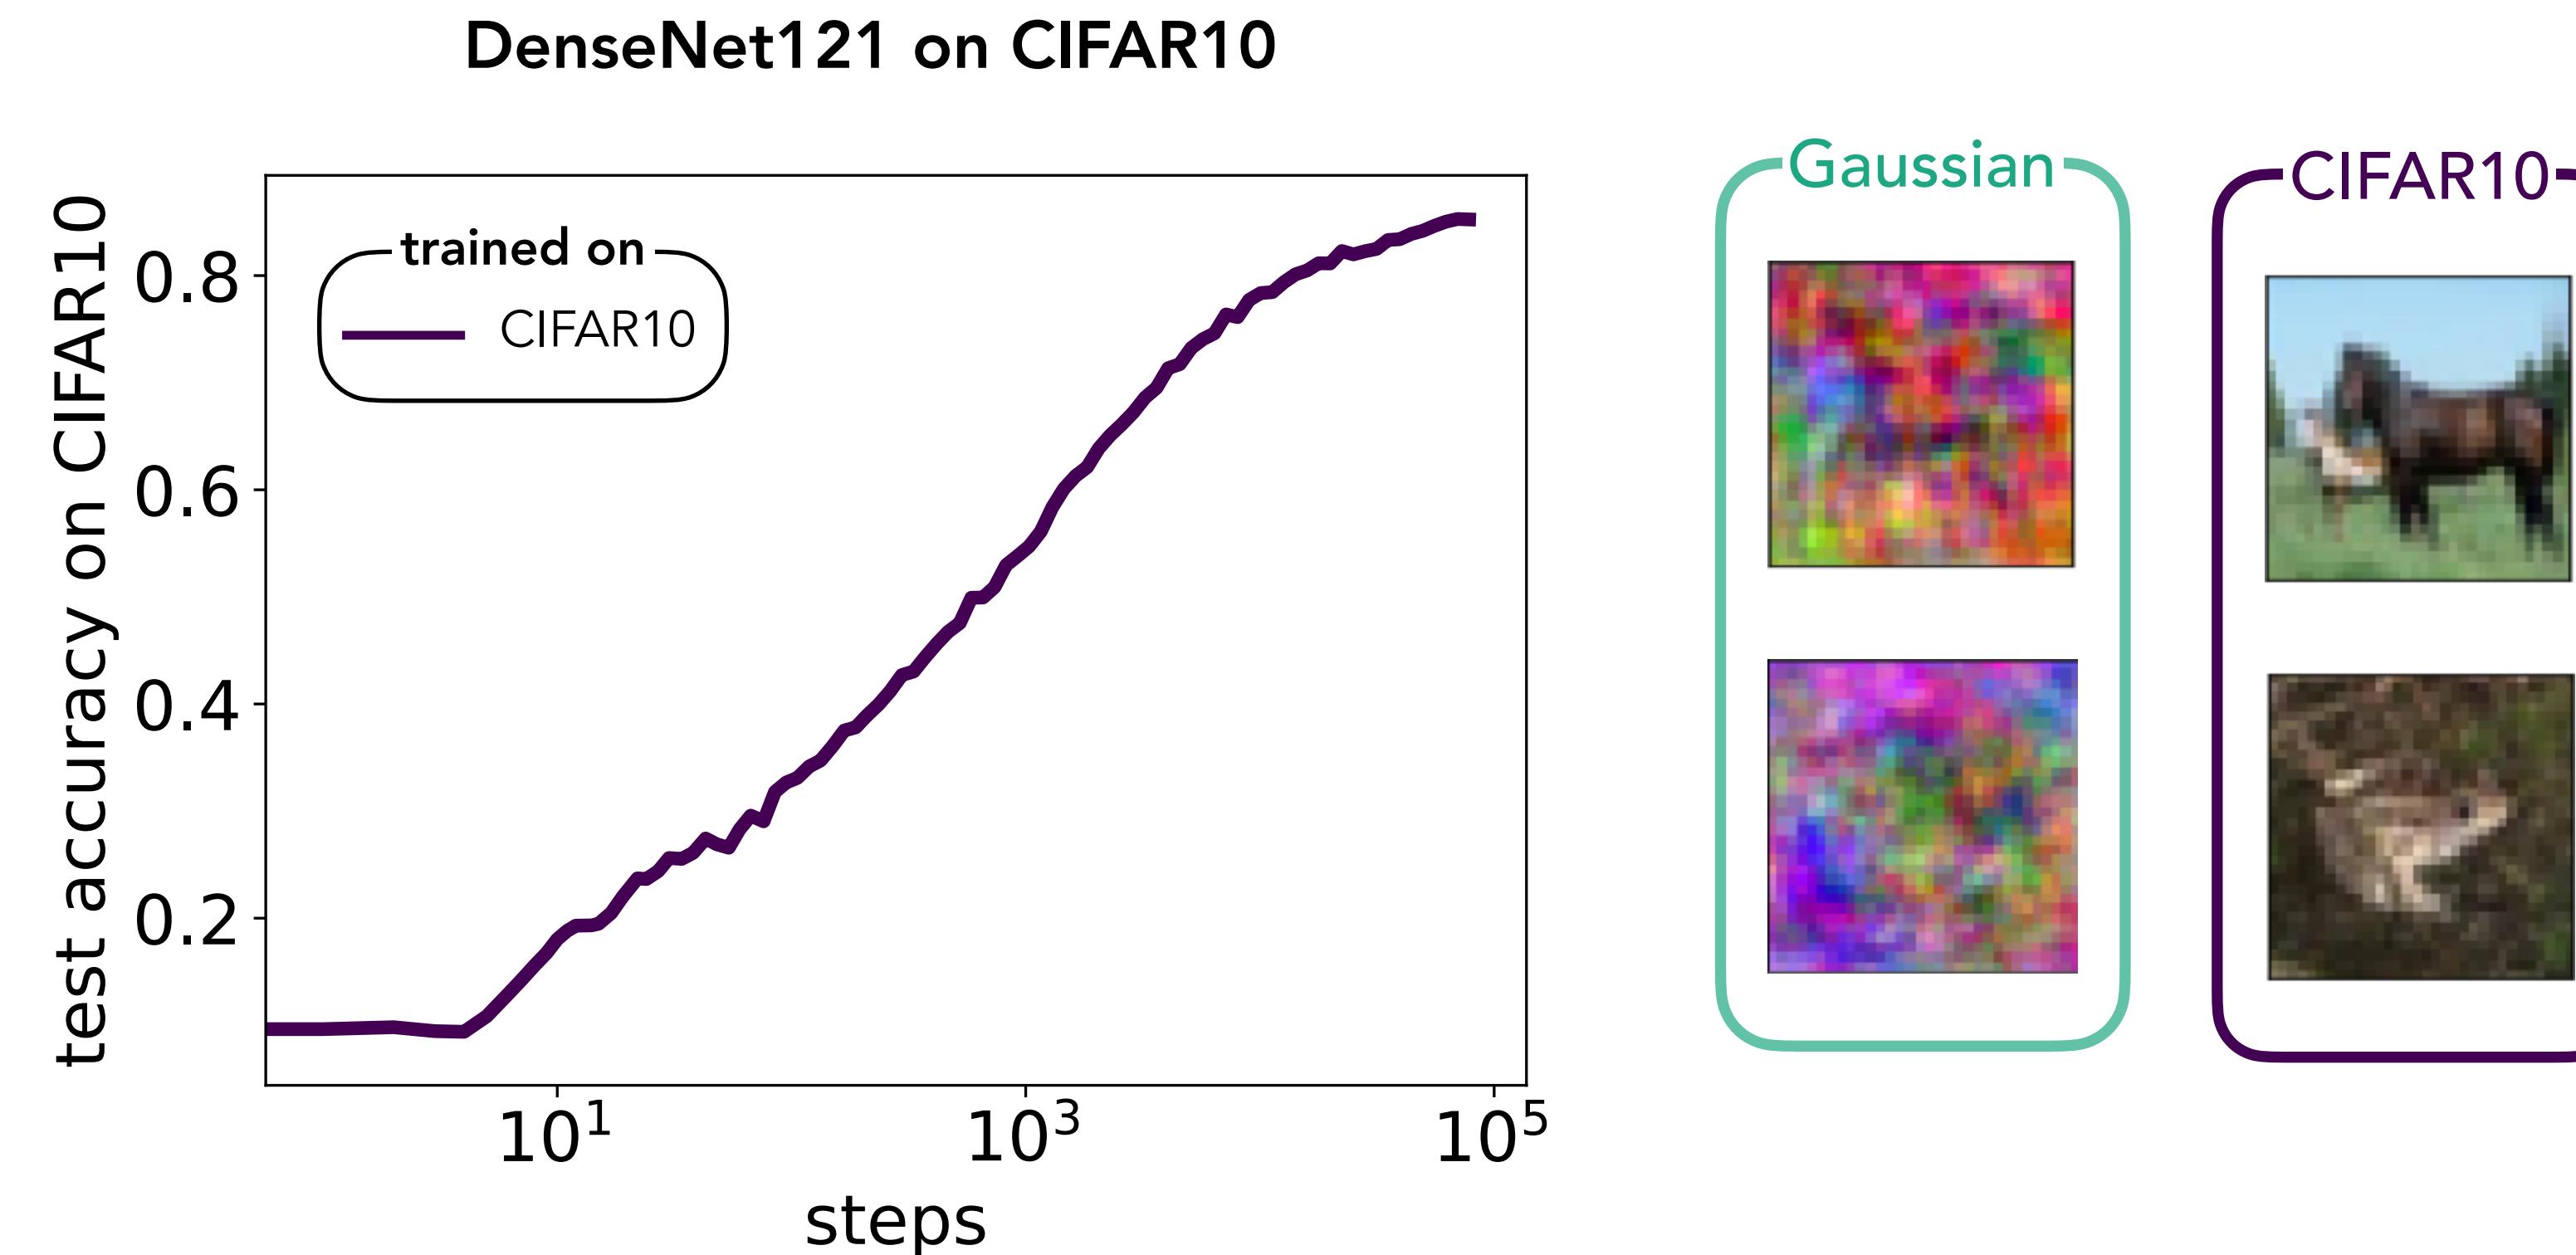
**Sequential
learning !**

What about deep networks?

Learning distributions of increasing complexity

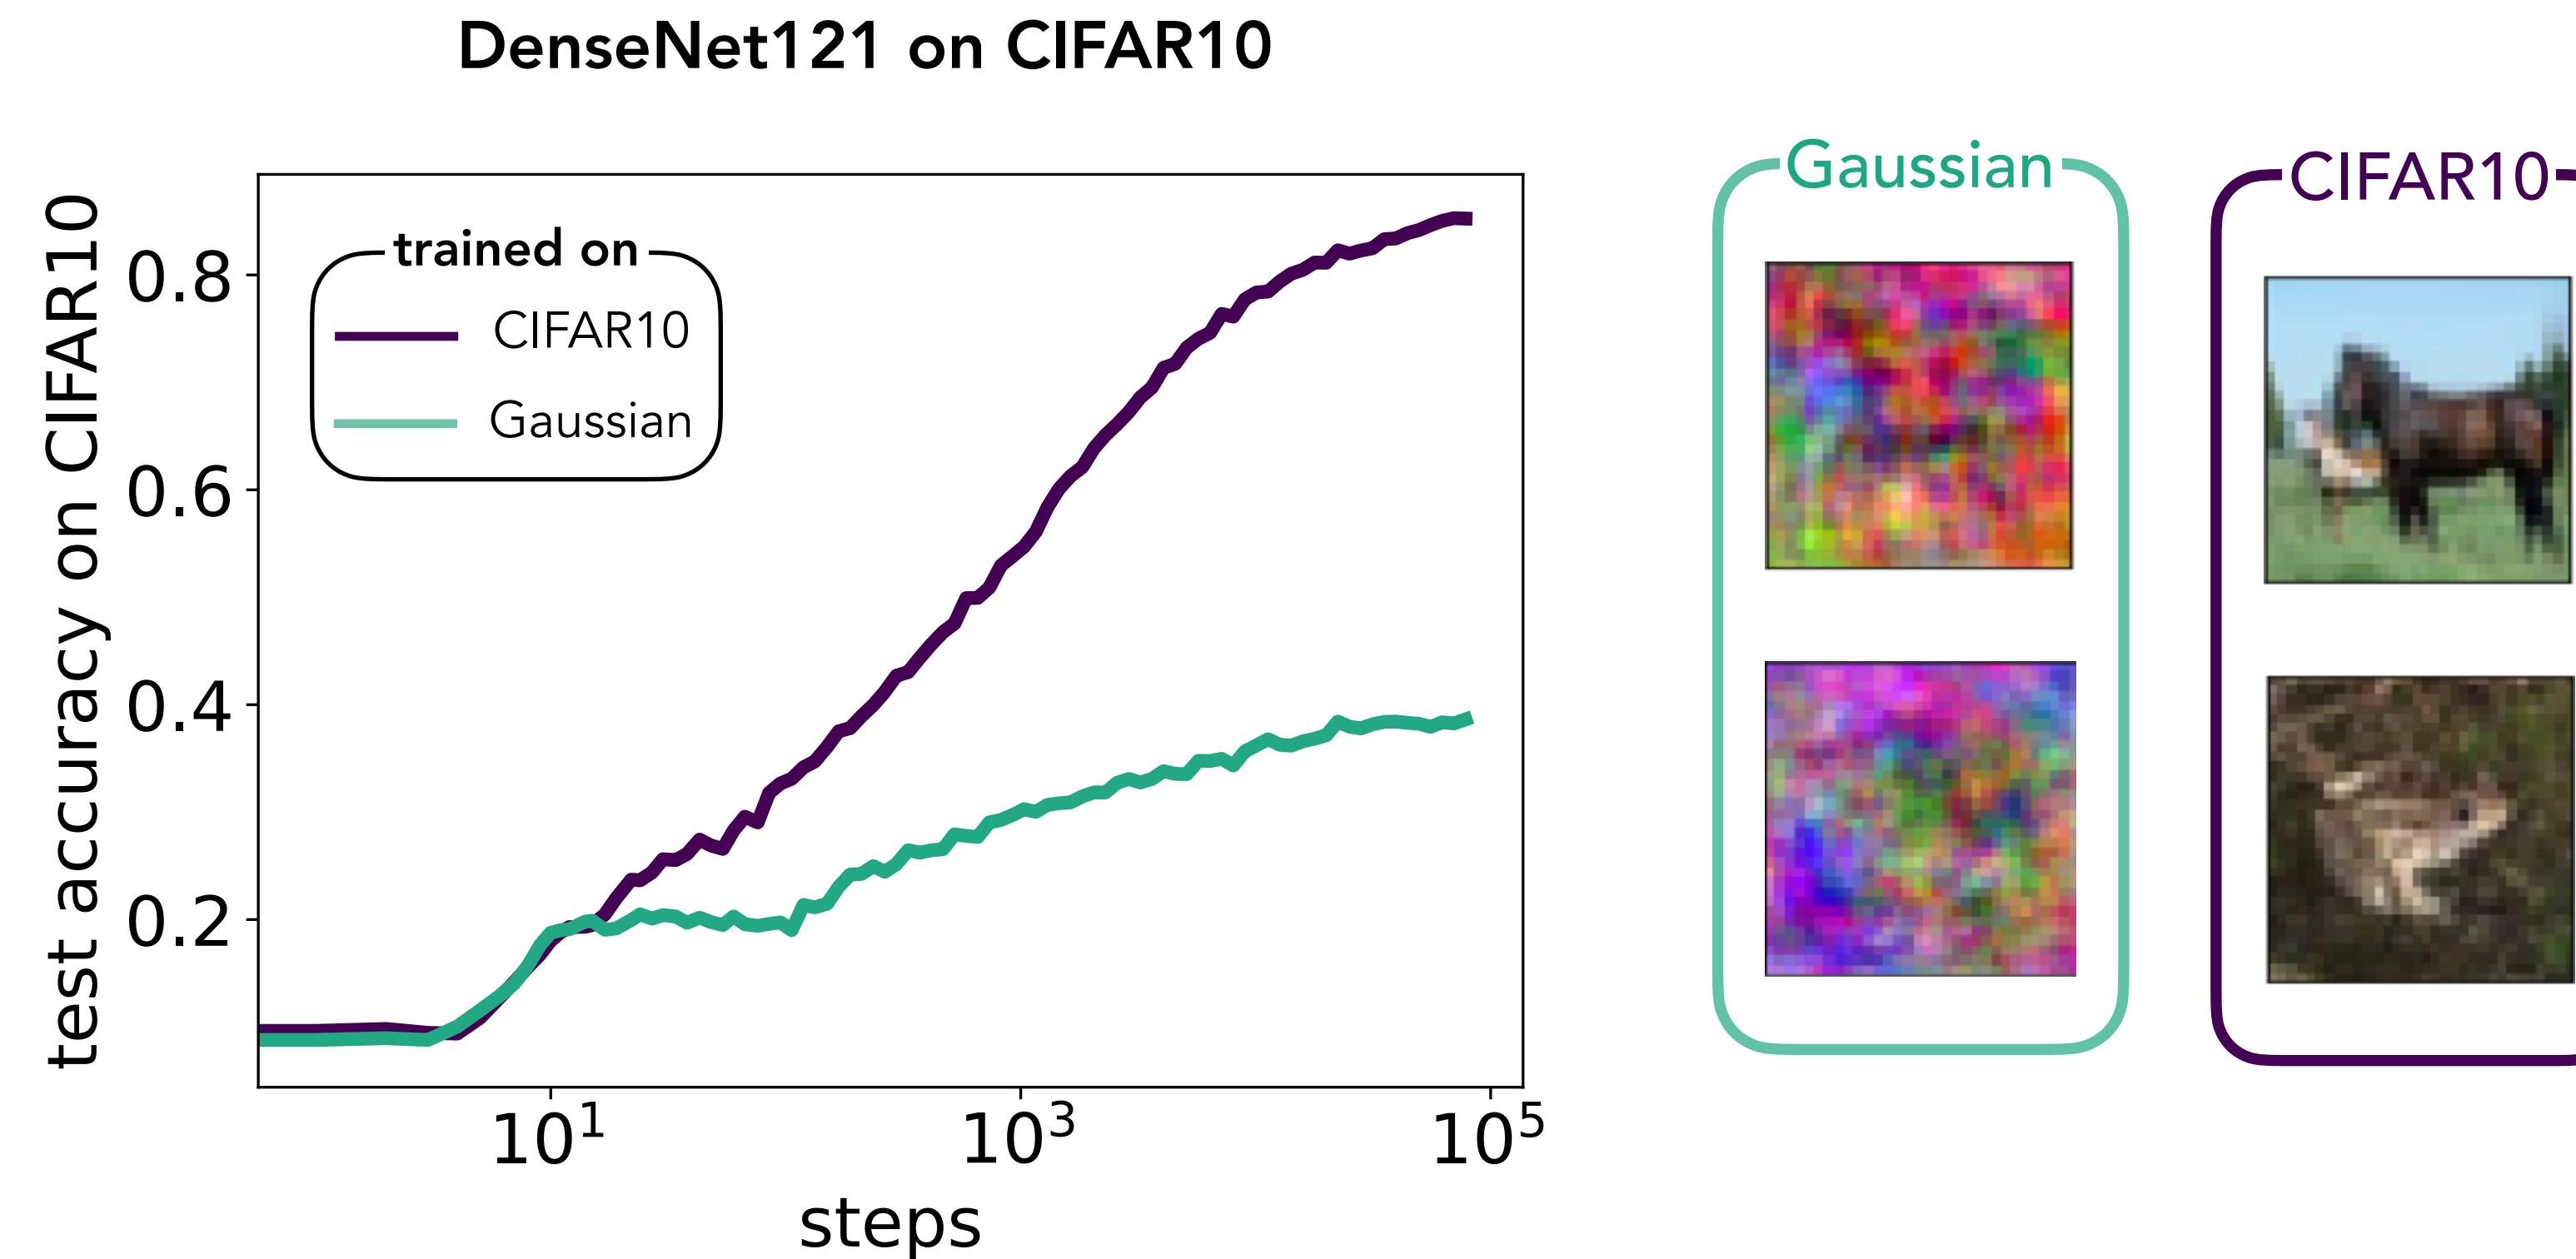
What about deep networks?

Learning distributions of increasing complexity



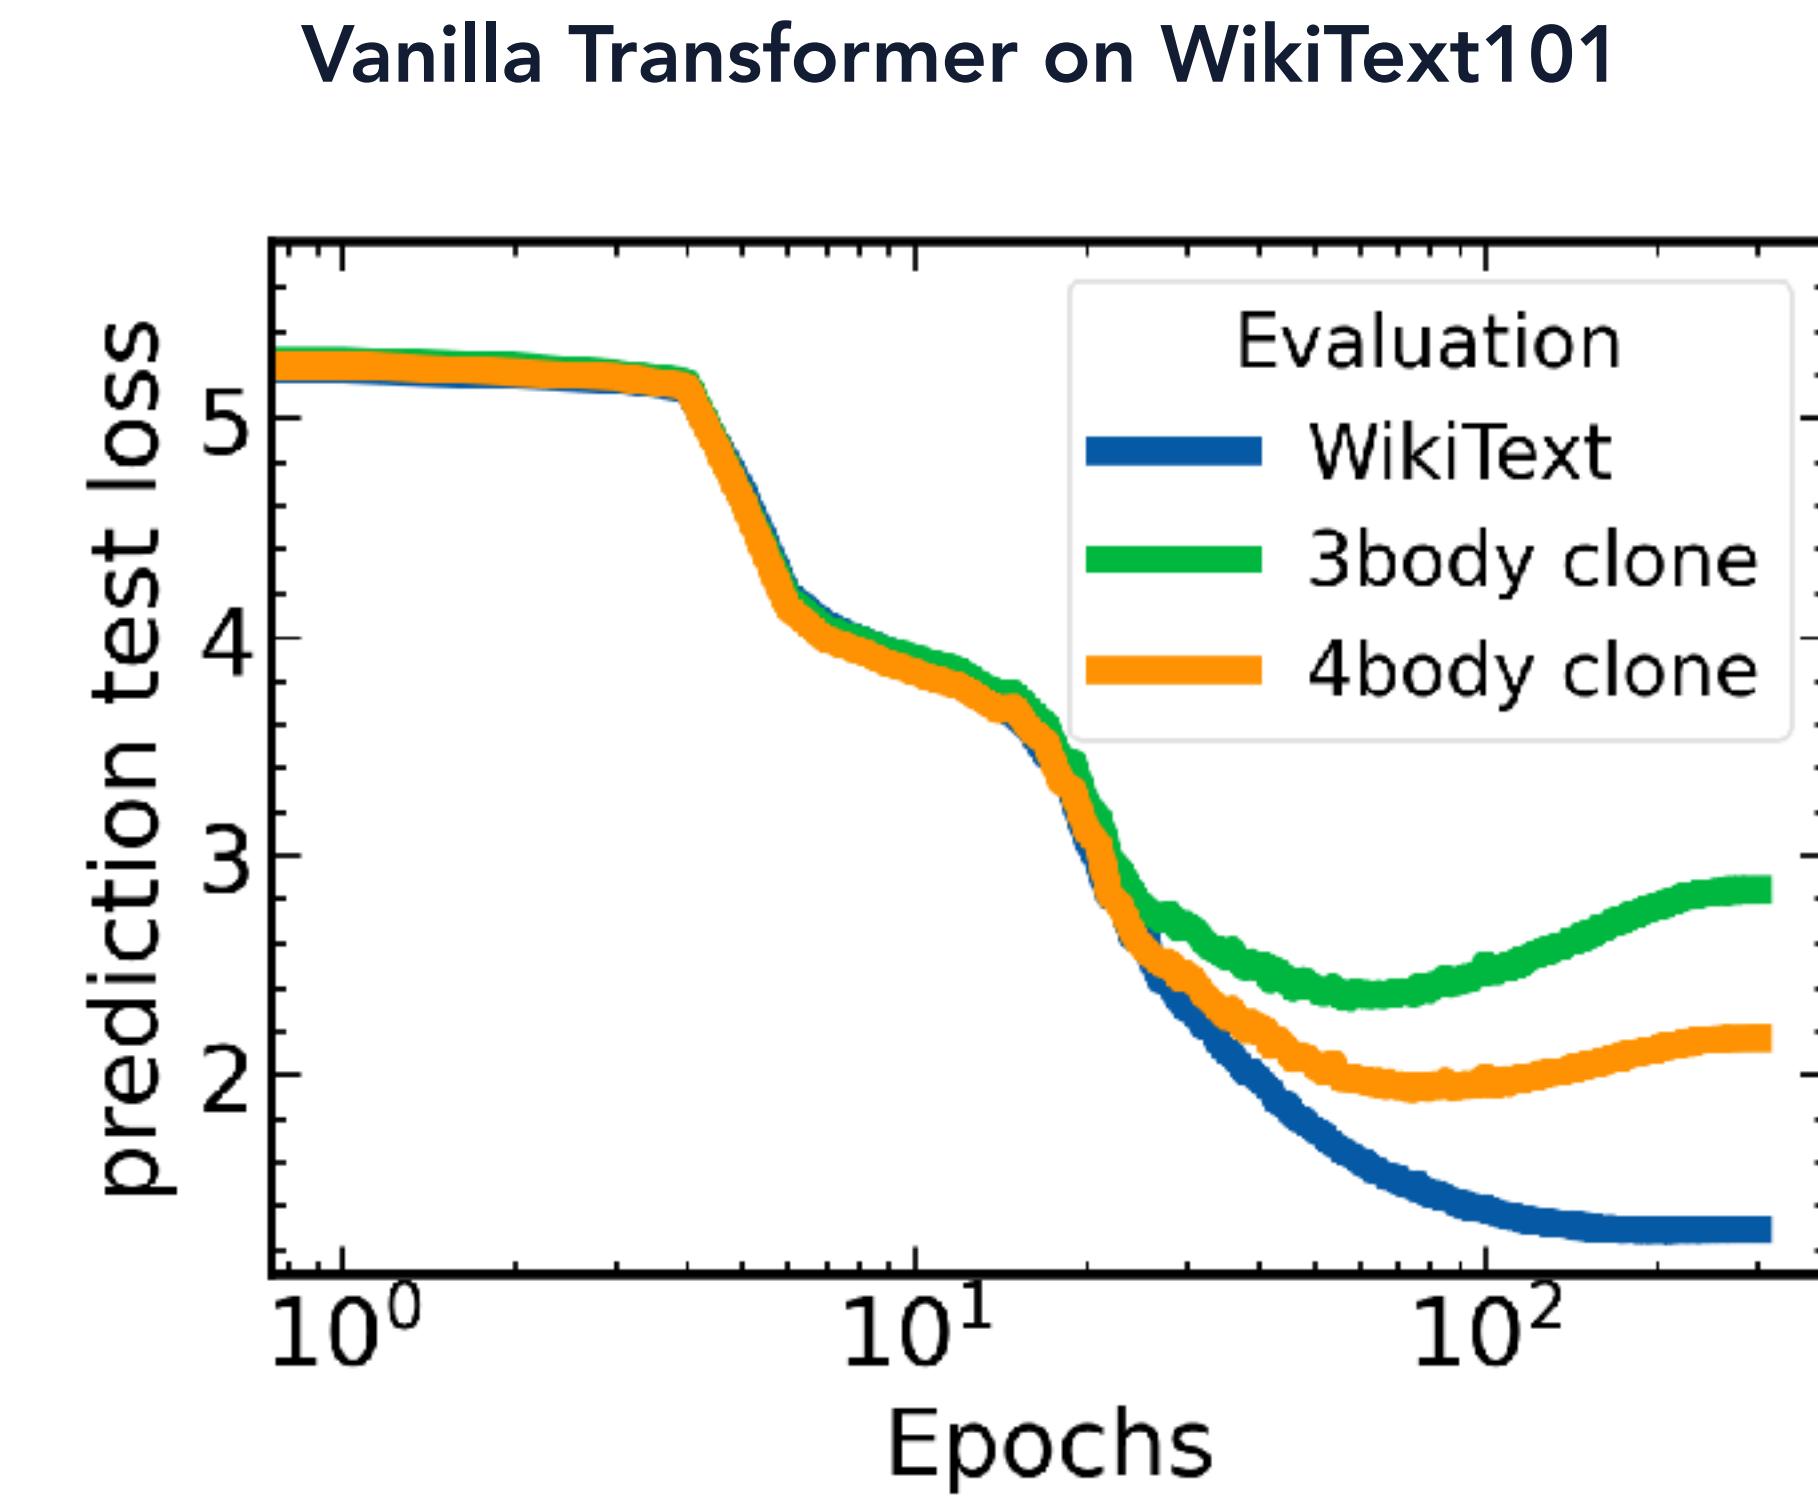
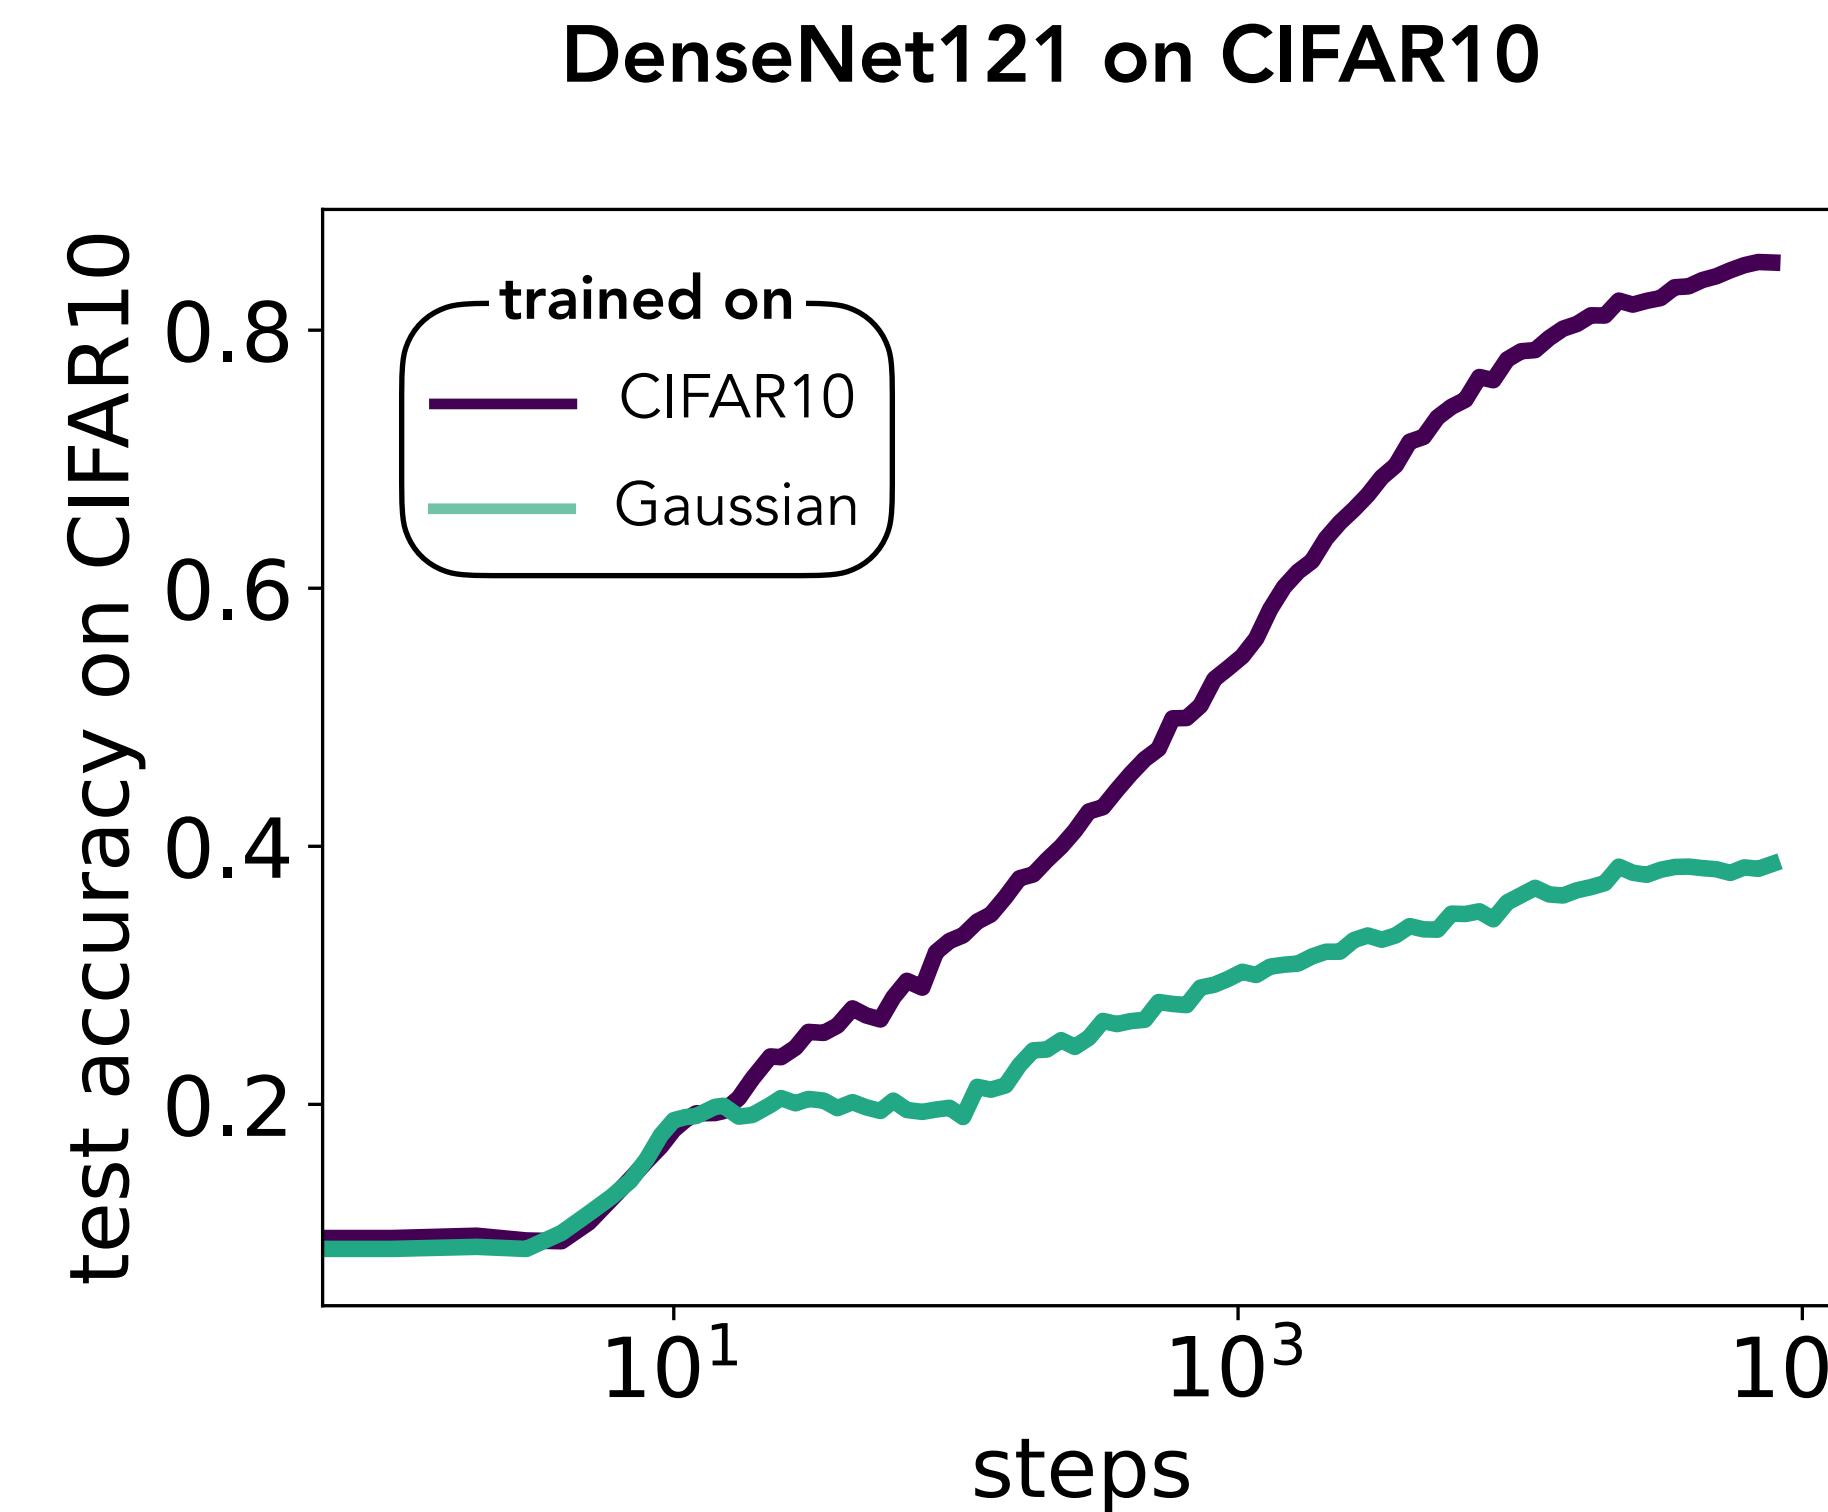
What about deep networks?

Learning distributions of increasing complexity



What about deep networks?

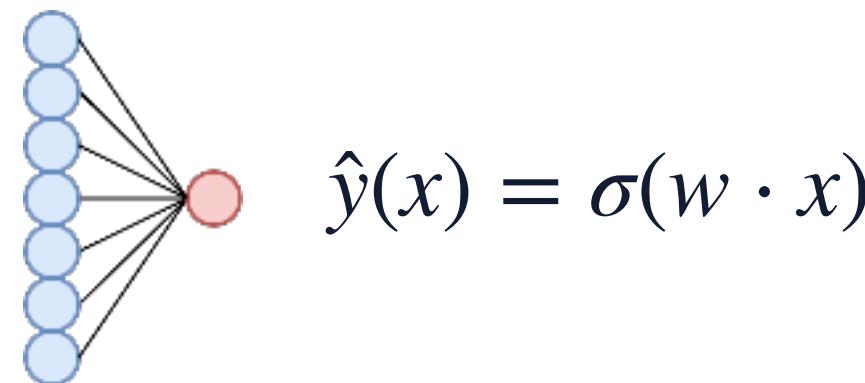
Learning distributions of increasing complexity



Connected subspaces accelerate neural networks

Rigorous analysis for a spherical perceptron

Spherical perceptron



$$\begin{cases} w_0 \sim \text{Unif}(\mathbb{S}^{d-1}) \\ \tilde{w}_t = w_{t-1} - \frac{\delta}{d} \nabla_{\text{sph}} (\mathcal{L}(w, (x_t, y_t))) \\ w_t = \frac{\tilde{w}_t}{\|\tilde{w}_t\|}. \end{cases}$$

Correlation loss

$$\mathcal{L}(w, (x, y)) = 1 - yf(w, x).$$

Mixed cumulant model:

Two overlaps:

$$\mathbf{x}^\mu = \mathbf{w}^\mu \quad \text{vs.} \quad \mathbf{x}^\mu = \beta_1 g^\mu \mathbf{u} + S(\beta_2 h^\mu \mathbf{v} + \mathbf{w}^\mu)$$

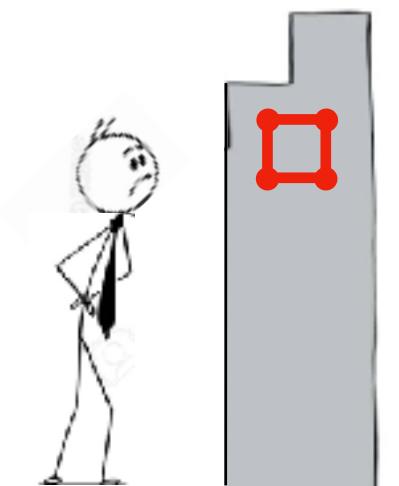
$$m_u = u \cdot w \quad m_v = v \cdot w$$

Disconnected subspaces:

Ben Arous et al. '21

$$m'_v(t) \approx 4c_{04}m_v^3$$

$$n_v \gg d^3$$



Correlated latents (=connected subspaces)

$$m'_v(t) \approx c_{11}m_u + 4c_{04}m_v^3$$

$$n_v \gg d$$



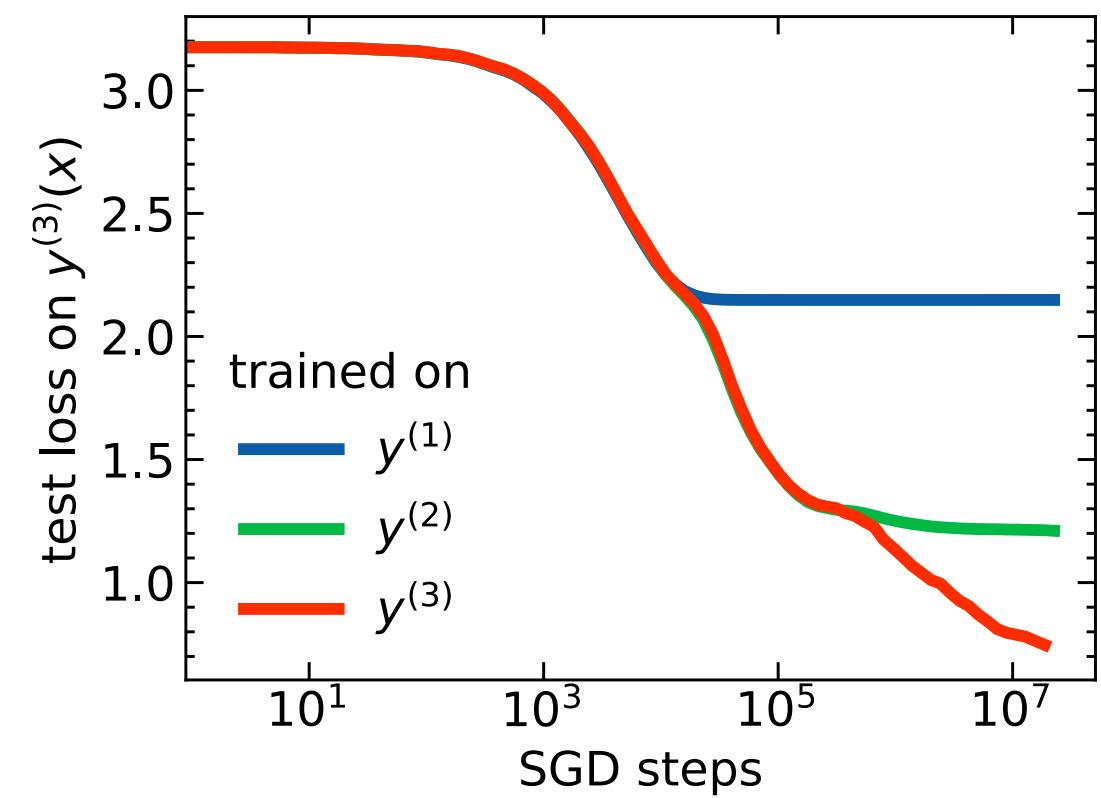
Relation to teacher-student models

Staircases, staircases everywhere!



Teacher-student model: $x \sim \mathcal{N}(0, 1_d)$

$$y^*(x) = h_1(m \cdot x) + h_2(u \cdot x) + h_4(v \cdot x)$$



Abbé '21, '22, '23; Jacot et al. '21;
Boursier et al. '22; Dandi et al. '23;
Damian et al. '23; Bietti et al. '23;
Mousavi-Hosseini et al. '24

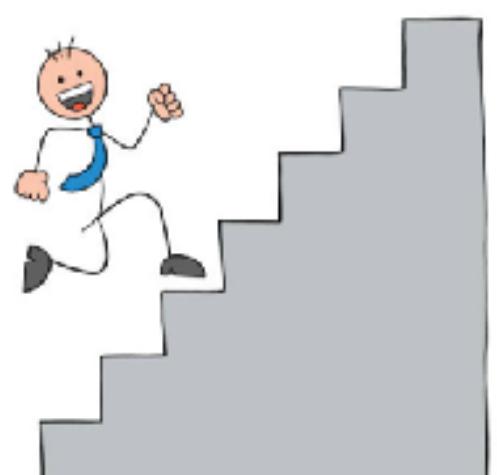
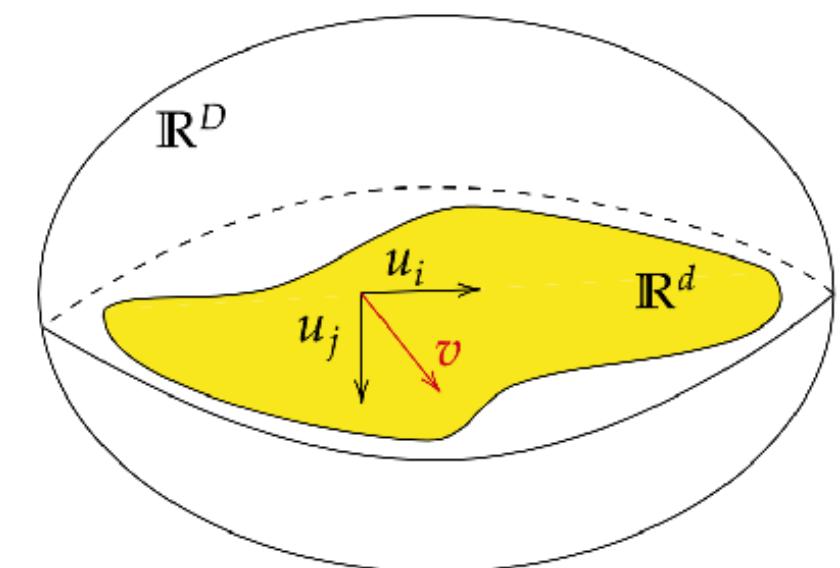
Key difference between spiked cumulants and teacher-student:

- **Generative exponent** [Damian et al. '24] of any **polynomial** is at most 2, so $y^*(x)$ can be learnt at linear sample complexity (e.g. by repeating batches [Dandi et al. '24])
- The generative exponent of the **spiked cumulant model** is at least four, since for binary labels, there is no transform T such that $\mathbb{E} [h_{1/2/3}(x) | T(y)] \neq 0$.

Concluding perspectives

How do neural networks learn from their data, efficiently?

- Neural networks learn features from **higher-order correlations**.
- ICA as a model system reveals the crucial role of **sequential learning** to access HOCs.
- We find similar behaviour in **deep CNNs**.
- Key challenge: towards more realistic models of **unsupervised learning**!?



Acknowledgements



Lorenzo Bardone
(SISSA)



Fabiola Ricci
(SISSA)



We're hiring!



WE WANT YOU



datascience.sissa.it



Finanziato
dall'Unione europea
NextGenerationEU



European Research Council
Established by the European Commission



@sebastiangoldt

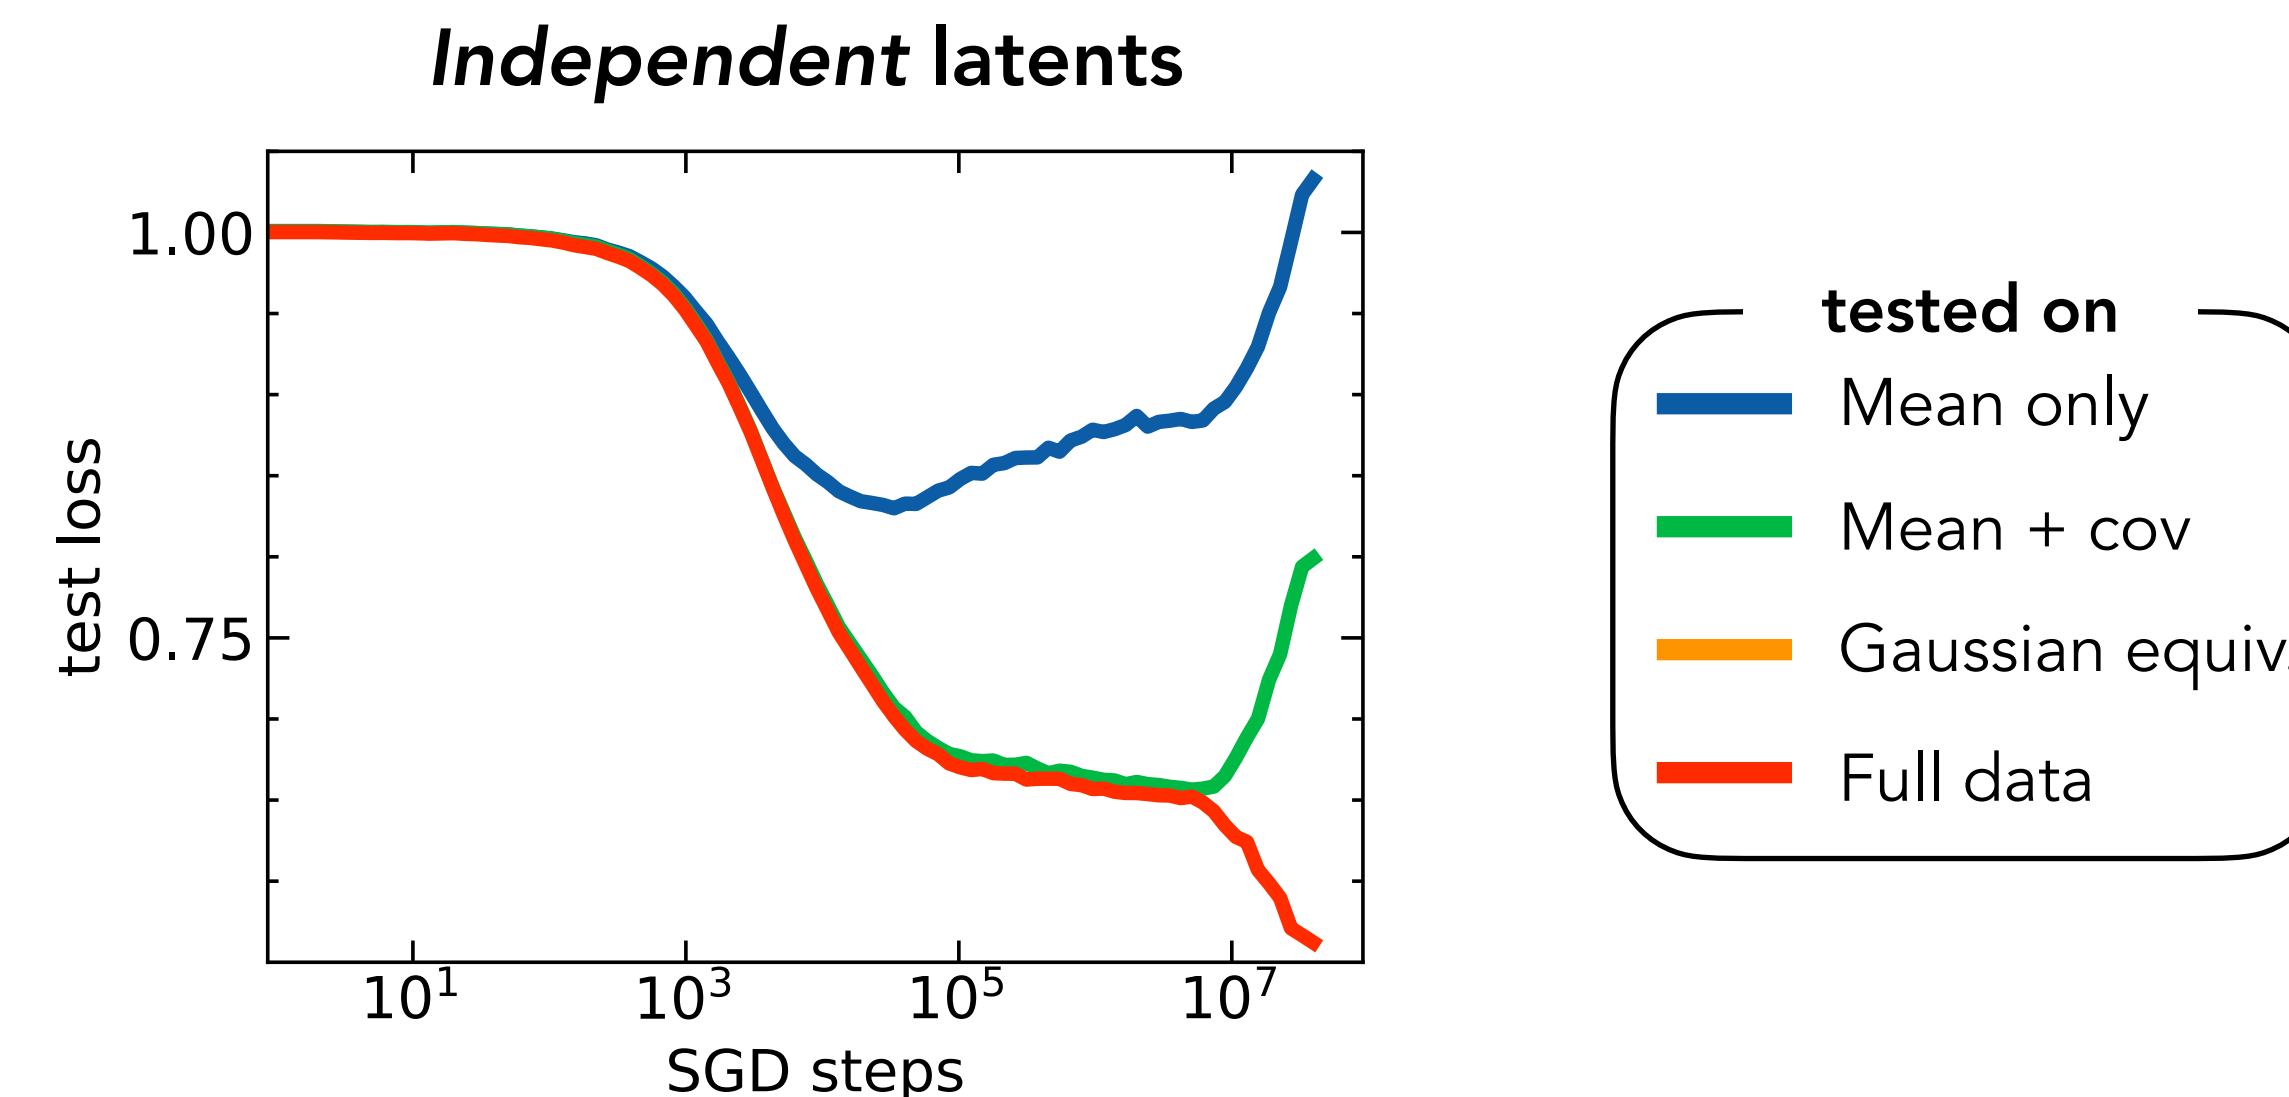
Connected subspaces accelerate neural networks

Two-layer neural networks exploit correlations between subspaces

Classification task: $\mathbf{x}^\mu = \mathbf{w}^\mu$ vs. $\mathbf{x}^\mu = \beta_0 \mathbf{m} + \beta_1 g^\mu \mathbf{u} + \beta_2 h^\mu \mathbf{v} + \mathbf{w}^\mu$

Three spikes: \mathbf{m} , \mathbf{u} , \mathbf{v}

Two latent variables: $g^\mu \sim \mathcal{N}(0,1)$, $h^\mu = \pm 1$



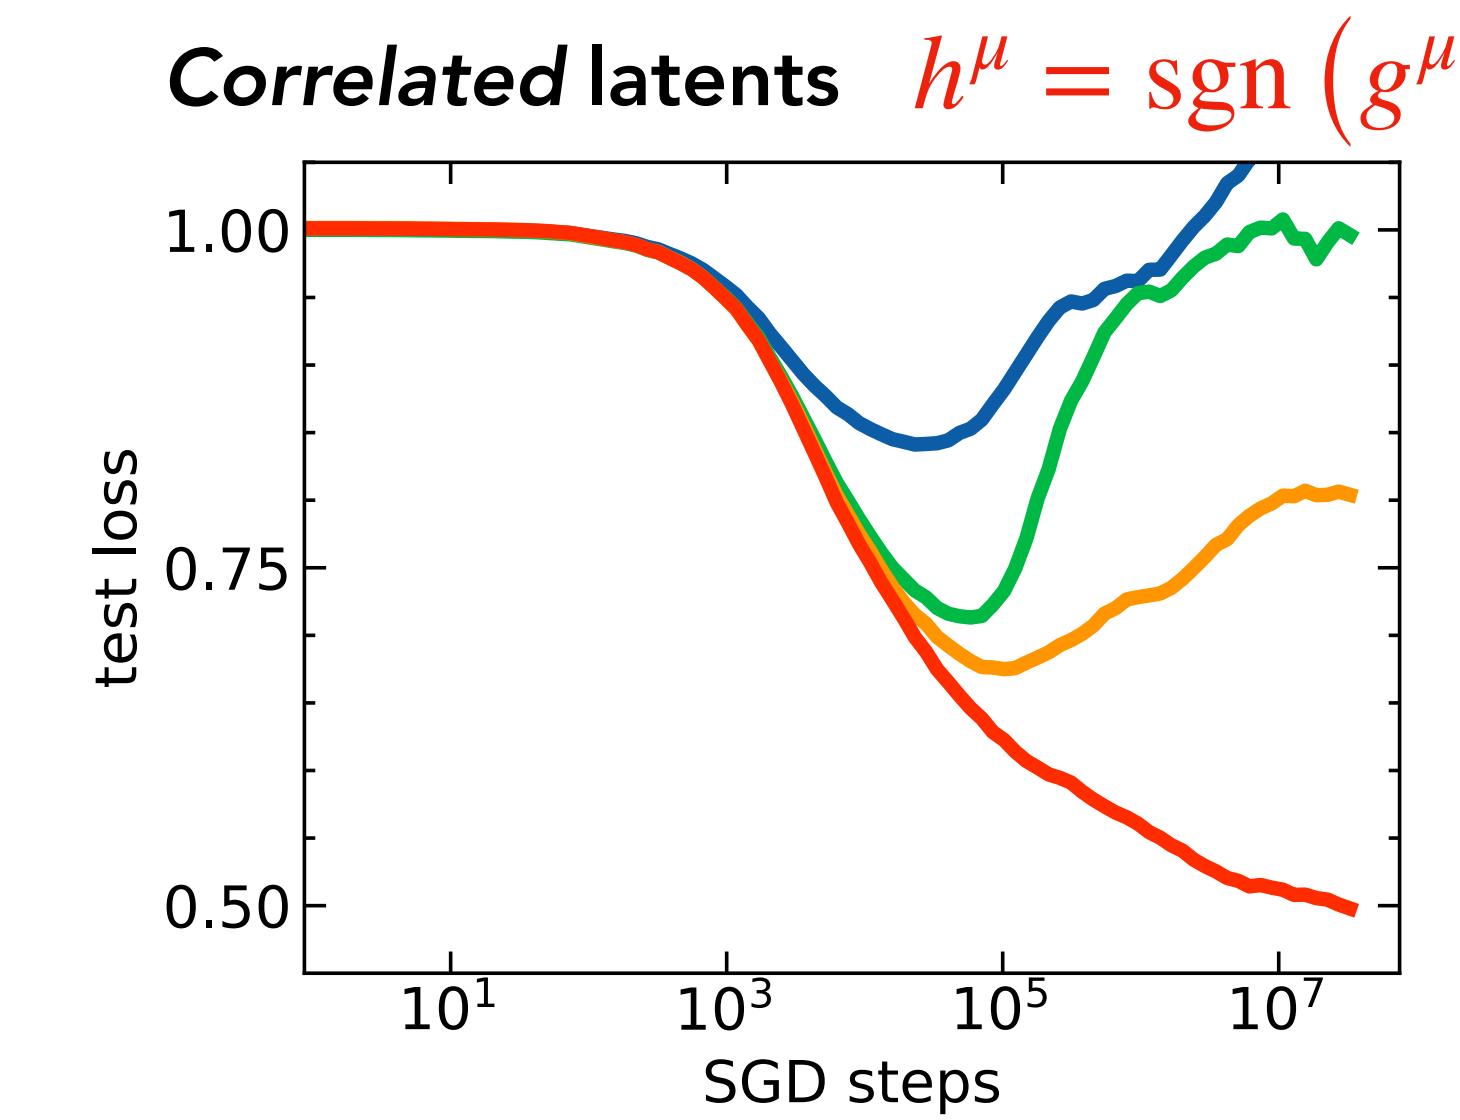
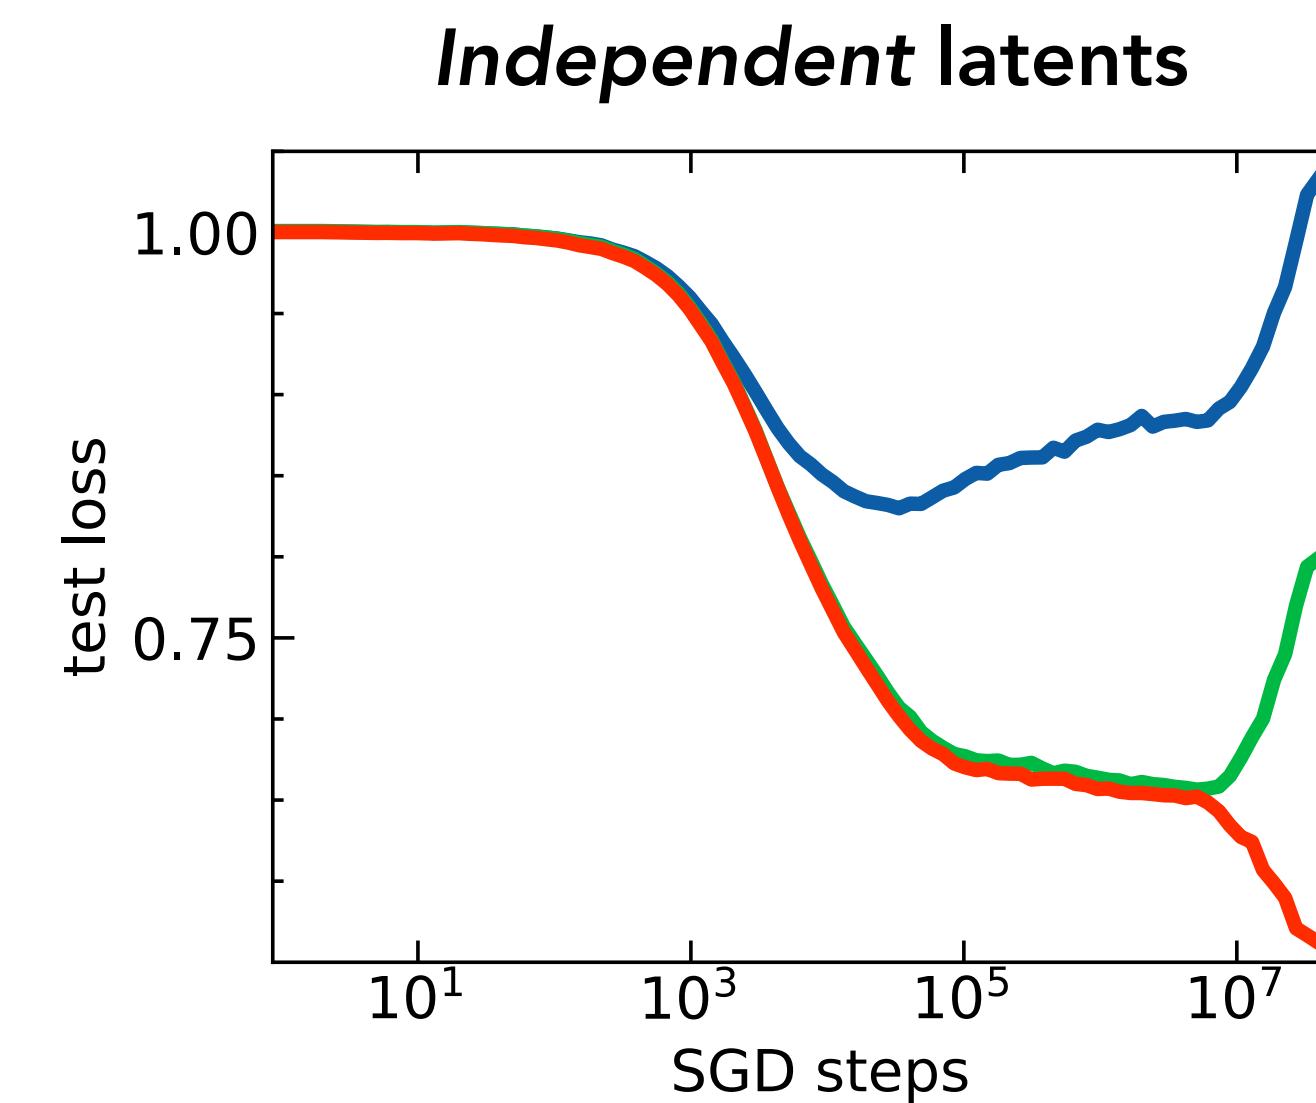
Connected subspaces accelerate neural networks

Two-layer neural networks exploit correlations between subspaces

Classification task: $\mathbf{x}^\mu = \mathbf{w}^\mu$ vs. $\mathbf{x}^\mu = \beta_0 \mathbf{m} + \beta_1 g^\mu \mathbf{u} + \beta_2 h^\mu \mathbf{v} + \mathbf{w}^\mu$

Three spikes: \mathbf{m} , \mathbf{u} , \mathbf{v}

Two latent variables: $g^\mu \sim \mathcal{N}(0,1)$, $h^\mu = \pm 1$

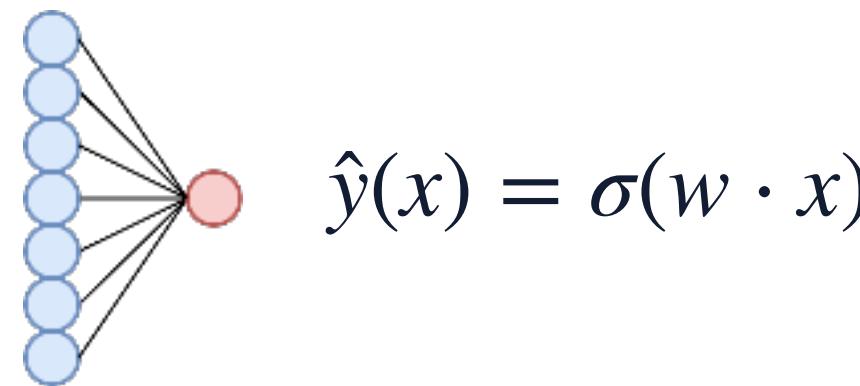


- tested on
- Mean only
- Mean + cov
- Gaussian equiv.
- Full data

Connected subspaces accelerate neural networks

Rigorous analysis for a spherical perceptron

Spherical perceptron



$$\begin{cases} w_0 \sim \text{Unif}(\mathbb{S}^{d-1}) \\ \tilde{w}_t = w_{t-1} - \frac{\delta}{d} \nabla_{\text{sph}} (\mathcal{L}(w, (x_t, y_t))) \\ w_t = \frac{\tilde{w}_t}{\|\tilde{w}_t\|}. \end{cases}$$

Correlation loss

$$\mathcal{L}(w, (x, y)) = 1 - yf(w, x).$$

Mixed cumulant model:

Two overlaps:

Cumulant spike only:

Ben Arous et al. '21

Correlated latents:

$$\mathbf{x}^\mu = \mathbf{w}^\mu \quad \text{vs.} \quad \mathbf{x}^\mu = \beta_1 g^\mu \mathbf{u} + S(\beta_2 h^\mu \mathbf{v} + \mathbf{w}^\mu)$$

$$\alpha_u = u \cdot w, \quad \alpha_v = v \cdot w$$

$$g^\mu = 0; \quad h^\mu = \pm 1 \quad \Rightarrow \quad n_v \gg d^3$$

$$\begin{cases} \dot{\alpha}_u(t) = 2c_{20}\alpha_u + c_{11}\alpha_v + O(\eta^2) \\ \dot{\alpha}_v(t) = \color{red}c_{11}\alpha_u + 4c_{04}\alpha_v^3 - 2c_{20}\alpha_u^2\alpha_v + O(\eta^4) \end{cases}$$

For correlated latents $\mathbb{E}\lambda\nu > 0$, the coefficient $c_{11} > 0$ and $n_v \gg d$

