

The Computational Advantage of Depth: Learning High-Dimensional Hierarchical Functions with Gradient Descent




The Computational Advantage of Depth: Learning High-Dimensional
Hierarchical Functions with Gradient Descent

Yatin Dandi^{1,2}, Luca Pesce¹, Lenka Zdeborová², and Florent Krzakala¹

¹Ecole Polytechnique Fédérale de Lausanne, Information, Learning and Physics Laboratory. CH-1015 Lausanne, Switzerland.

²Ecole Polytechnique Fédérale de Lausanne, Statistical Physics of Computation Laboratory. CH-1015 Lausanne, Switzerland.

Understanding deep learning is also a job*




Comment | Published: 26 May 2020

Understanding deep learning is also a job for physicists

[Lenka Zdeborová](#)✉

Nature Physics 16, 602–604 (2020) | [Cite this article](#)

Understanding deep learning is also a job*



*Till AGI


Comment | Published: 26 May 2020

Understanding deep learning is also a job for physicists

[Lenka Zdeborová](#)✉

Nature Physics **16**, 602–604 (2020) | [Cite this article](#)

Understanding deep learning is also a job*



*Till AGI

Comment | Published: 26 May 2020


Understanding deep learning is also a job for physicists

[Lenka Zdeborová](#)✉

Nature Physics 16, 602–604 (2020) | [Cite this article](#)

Understanding deep learning is also a job*

*Till AGI




It's called *deep* learning for a reason.

It's called *deep* learning for a reason.

VGG

VERY DEEP CONVOLUTIONAL NETWORKS FOR LARGE-SCALE IMAGE RECOGNITION

Karen Simonyan* & Andrew Zisserman⁺
Visual Geometry Group, Department of Engineering Science, University of Oxford
`{karen,az}@robots.ox.ac.uk`



Resnet


by the number of stacked layers (depth). Recent evidence [41, 44] reveals that network depth is of crucial importance, and the leading results [41, 44, 13, 16] on the challenging ImageNet dataset [36] all exploit “very deep” [41] models, with a depth of sixteen [41] to thirty [16]. Many other non-

It's called *deep* learning for a reason.

VGG

VERY DEEP CONVOLUTIONAL NETWORKS FOR LARGE-SCALE IMAGE RECOGNITION

Karen Simonyan* & Andrew Zisserman⁺
Visual Geometry Group, Department of Engineering Science, University of Oxford
{karen,az}@robots.ox.ac.uk



Exponential gains in
approximation capacity

The Power of Depth for Feedforward Neural Networks

Ronen Eldan
Weizmann Institute of Science
ronen.eldan@weizmann.ac.il

Ohad Shamir
Weizmann Institute of Science
ohad.shamir@weizmann.ac.il

Abstract

We show that there is a simple (approximately radial) function on \mathbb{R}^d , expressible by a small 3-layer feedforward neural networks, which cannot be approximated by any 2-layer network, to more than a certain constant accuracy, unless its width is exponential in the dimension. The result holds for virtually

Resnet

by the number of stacked layers (depth). Recent evidence [41, 44] reveals that network depth is of crucial importance, and the leading results [41, 44, 13, 16] on the challenging ImageNet dataset [36] all exploit “very deep” [41] models, with a depth of sixteen [41] to thirty [16]. Many other non-

Deep vs. Shallow Networks: an Approximation Theory Perspective

by

Hrushikesh N. Mhaskar¹ and Tomaso Poggio²

1. Department of Mathematics, California Institute of Technology, Pasadena, CA 91125
Institute of Mathematical Sciences, Claremont Graduate University, Claremont, CA 91711.
hrushikesh.mhaskar@cgu.edu


2. Center for Brains, Minds, and Machines, McGovern Institute for Brain Research,
Massachusetts Institute of Technology, Cambridge, MA, 02139.
tp@mit.edu

What do we expect to show?

What do we expect to show?

The quest for adaptivity

Posted on June 17, 2021 by Francis Bach




What do we expect to show?

The quest for adaptivity

Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**





What do we expect to show?

The quest for adaptivity

Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**




What do we expect to show?


The quest for adaptivity

Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**



**Lower-sample Complexity
with adaptive algorithms**




What do we expect to show?

The quest for adaptivity


Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**



**Lower-sample Complexity
with adaptive algorithms**

**Gaussian data+ single, multi-
index models**





What do we expect to show?

The quest for adaptivity


Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**



**Lower-sample Complexity
with adaptive algorithms**

**Gaussian data+ single, multi-
index models**




Hierarchical structure

What do we expect to show?

The quest for adaptivity



Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**




**Lower-sample Complexity
with adaptive algorithms**

**Gaussian data+ single, multi-
index models**



Hierarchical structure




What do we expect to show?

The quest for adaptivity



Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**




**Lower-sample Complexity
with adaptive algorithms**

**Gaussian data+ single, multi-
index models**



Hierarchical structure




**Lower-sample Complexity
With deep/multiple levels
of adaptation.**

What do we expect to show?

The quest for adaptivity



Posted on June 17, 2021 by Francis Bach

**Low-dimensional
structure**




**Lower-sample Complexity
with adaptive algorithms**

**Gaussian data+ single, multi-
index models**



Hierarchical structure




**Lower-sample Complexity
With deep/multiple levels
of adaptation.**

?


What's “hierarchical”?

What's “hierarchical”?



What's “hierarchical”?


**Large search
space for target**



**Large
effective dimension**


What's “hierarchical”?

**Large search
space for target**



**Large
effective dimension**

What's “hierarchical”?




**Large search
space for target**

**Large
effective dimension**

**Small search space
for target**

**small effective
dimension**

What's “hierarchical”?



Large search space for target

Small search space for target

Large effective dimension


Filtering of features of increasing complexity

small effective dimension



Akin to Coarse-graining/Renormalization (Gilio's talk)

What's “hierarchical”?



Large search
space for target

Large
effective dimension

Filtering of features of
increasing complexity


See also Cagnetta et al.
2025, Mossel et al., 2016,
Bruna & Mallat 2013

Small search space
for target

small effective
dimension

Akin to Coarse-graining/Renormalization (Gulio's talk)

What's “hierarchical”?



Large search
space for target

See also Cagnotta et al.
2025, Mossel et al., 2016,
Bruna & Mallat 2013

Small search space
for target


Large
effective dimension

Filtering of features of
increasing complexity


small effective
dimension

Akin to Coarse-graining/Renormalization (Gilio's talk)

Can we understand this in some
analyzable setting?



Recap of two-layer NNs for AGI



Recap of two-layer NNs for AGI

 **Jason Lee** 
@jasondeanlee

At the [@SimonsInstitute](#) working on AGI (Artificial Gaussian Intelligence)

5:12 PM · Nov 15, 2024 · 11.5K Views

Neural Networks without Feature Learning

Random features


$$\hat{y} = \hat{f}(\mathbf{x}) = \hat{\mathbf{a}} \cdot \sigma(W\mathbf{x})$$

[Balcan,Blum, Vempala '06, Rahimi-Recht '17...]

No training of the first layer: W is fixed

$$\hat{y} = \hat{f}(\mathbf{x}) = \sum_{i=1}^p \hat{a}_i \sigma_i(\langle \mathbf{w}_i, \mathbf{x} \rangle) = \sum_{i=1}^p \hat{a}_i \Phi_{CK}(\mathbf{x})$$

Computationally easy (linear regression)



Very popular setting among theoreticians

Equivalent to Neural Tangent Kernel/Lazy Regime/Kernel methods/ etc..

[Jacot, Gabriel, Hongler '18; Lee, Jaehoon, et al. 18; Chizat, Bach '19,...]


Recap of Sample Complexity

Theorem (Informal) [Mei, Misiakiewicz, Montanari '22]

In absence of feature learning (i.e. at initialization) one can only learn a **polynomial approximation of f^* of degree κ** with $\min(n, p) = O(d^\kappa)$

$$f^*(\mathbf{x}) = \text{cst} + \sum_i \mu_i^{(1)} h_i^* + \sum_{ij} \mu_{ij}^{(2)} h_i^* h_j^* + \sum_{ijk} \mu_{ijk}^{(3)} h_i^* h_j^* h_k^* + \dots$$

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (f^*(\mathbf{x}) - \hat{\mathbf{f}}(\mathbf{x}))^2$$



See also [El Karaoui '10; Mei-Montanari '19; Gerace '20; Jacot, Simsek, Spadaro, Hongler, Gabriel '20; Hu, Lu, '20; Dhifallah, Lu '20; Loureiro, Gerbelot, Cui, Goldt '21; Montanari & Saeed '22; Xiao, Hu, Misiakiewicz, Lu, Pennington '22; Dandi '23; Aguirre-López, Franz, Pastore '24]


Recap of Sample Complexity

Theorem (Informal) [Mei, Misiakiewicz, Montanari '22]

In absence of feature learning (i.e. at initialization) one can only learn a **polynomial approximation of f^* of degree κ** with $\min(n, p) = O(d^\kappa)$

$$f^*(\mathbf{x}) = \text{cst} + \sum_i \mu_i^{(1)} h_i^* + \sum_{ij} \mu_{ij}^{(2)} h_i^* h_j^* + \sum_{ijk} \mu_{ijk}^{(3)} h_i^* h_j^* h_k^* + \dots$$
$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (f^*(\mathbf{x}) - \hat{\mathbf{f}}(\mathbf{x}))^2$$

\vdots
 \vdots
 \vdots
 \vdots
 \vdots
 \vdots
 $(n, p) = O(d)$



See also [El Karaoui '10; Mei-Montanari '19; Gerace '20; Jacot, Simsek, Spadaro, Hongler, Gabriel '20; Hu, Lu, '20; Dhifallah, Lu '20; Loureiro, Gerbelot, Cui, Goldt '21; Montanari & Saeed '22; Xiao, Hu, Misiakiewicz, Lu, Pennington '22; Dandi '23; Aguirre-López, Franz, Pastore '24]


Recap of Sample Complexity

Theorem (Informal) [Mei, Misiakiewicz, Montanari '22]

In absence of feature learning (i.e. at initialization) one can only learn a **polynomial approximation of f^* of degree κ** with $\min(n, p) = O(d^\kappa)$

$$f^*(\mathbf{x}) = \text{cst} + \sum_i \mu_i^{(1)} h_i^* + \sum_{ij} \mu_{ij}^{(2)} h_i^* h_j^* + \sum_{ijk} \mu_{ijk}^{(3)} h_i^* h_j^* h_k^* + \dots$$
$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (f^*(\mathbf{x}) - \hat{\mathbf{f}}(\mathbf{x}))^2$$

$(n, p) = O(d)$ $(n, p) = O(d^2)$



See also [El Karaoui '10; Mei-Montanari '19; Gerace '20; Jacot, Simsek, Spadaro, Hongler, Gabriel '20; Hu, Lu, '20; Dhifallah, Lu '20; Loureiro, Gerbelot, Cui, Goldt '21; Montanari & Saeed '22; Xiao, Hu, Misiakiewicz, Lu, Pennington '22; Dandi '23; Aguirre-López, Franz, Pastore '24]

Recap of Sample Complexity

Theorem (Informal) [Mei, Misiakiewicz, Montanari '22]

In absence of feature learning (i.e. at initialization) one can only learn a **polynomial approximation of f^* of degree κ** with $\min(n, p) = O(d^\kappa)$


$$f^*(\mathbf{x}) = \text{cst} + \sum_i \mu_i^{(1)} h_i^* + \sum_{ij} \mu_{ij}^{(2)} h_i^* h_j^* + \sum_{ijk} \mu_{ijk}^{(3)} h_i^* h_j^* h_k^* + \dots$$

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (f^*(\mathbf{x}) - \hat{\mathbf{f}}(\mathbf{x}))^2$$

$$(n, p) = O(d)$$

$$(n, p) = O(d^2)$$

$$(n, p) = O(d^3)$$



See also [El Karaoui '10; Mei-Montanari '19; Gerace '20; Jacot, Simsek, Spadaro, Hongler, Gabriel '20; Hu, Lu, '20; Dhifallah, Lu '20; Loureiro, Gerbelot, Cui, Goldt '21; Montanari & Saeed '22; Xiao, Hu, Misiakiewicz, Lu, Pennington '22; Dandi '23; Aguirre-López, Franz, Pastore '24]

Recap of Sample Complexity

Theorem (Informal) [Mei, Misiakiewicz, Montanari '22]

In absence of feature learning (i.e. at initialization) one can only learn a **polynomial approximation of f^* of degree κ** with $\min(n, p) = O(d^\kappa)$

$$f^*(\mathbf{x}) = \text{cst} + \sum_i \mu_i^{(1)} h_i^* + \sum_{ij} \mu_{ij}^{(2)} h_i^* h_j^* + \sum_{ijk} \mu_{ijk}^{(3)} h_i^* h_j^* h_k^* + \dots$$

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (f^*(\mathbf{x}) - \hat{\mathbf{f}}(\mathbf{x}))^2$$


$$(n, p) = O(d)$$

$$(n, p) = O(d^2)$$

$$(n, p) = O(d^3)$$

•
•
•
•
•
•

No adaptivity \implies searching in a $\mathcal{O}(d^k)$ dimensional subspace of polynomials*



See also [El Karaoui '10; Mei-Montanari '19; Gerace '20; Jacot, Simsek, Spadaro, Hongler, Gabriel '20; Hu, Lu, '20; Dhifallah, Lu '20; Loureiro, Gerbelot, Cui, Goldt '21; Montanari & Saeed '22; Xiao, Hu, Misiakiewicz, Lu, Pennington '22; Dandi '23; Aguirre-López, Franz, Pastore '24]

Recap of Sample Complexity


Theorem (Informal) [Mei, Misiakiewicz, Montanari '22]

In absence of feature learning (i.e. at initialization) one can only learn a **polynomial approximation of f^* of degree κ** with $\min(n, p) = O(d^\kappa)$

$$f^*(\mathbf{x}) = \text{cst} + \sum_i \mu_i^{(1)} h_i^* + \sum_{ij} \mu_{ij}^{(2)} h_i^* h_j^* + \sum_{ijk} \mu_{ijk}^{(3)} h_i^* h_j^* h_k^* + \dots$$
$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (f^*(\mathbf{x}) - \hat{\mathbf{f}}(\mathbf{x}))^2$$

$(n, p) = O(d)$ $(n, p) = O(d^2)$ $(n, p) = O(d^3)$

No adaptivity \implies searching in a $\mathcal{O}(d^k)$ dimensional subspace of polynomials*



*possibly excluding a few “special” polynomials.


See also [El Karaoui '10; Mei-Montanari '19; Gerace '20; Jacot, Simsek, Spadaro, Hongler, Gabriel '20; Hu, Lu, '20; Dhifallah, Lu '20; Loureiro, Gerbelot, Cui, Goldt '21; Montanari & Saeed '22; Xiao, Hu, Misiakiewicz, Lu, Pennington '22; Dandi '23; Aguirre-López, Franz, Pastore '24]

Advantage of Feature Learning

Ben Arous et al. 2021, Abbe et al. 2022, 2023, Damian et al. 2022, Dandi et al. 2023, 2024, etc.


Advantage of Feature Learning

$$y = f^*(\mathbf{x}) = \mathbf{g}^*(\mathbf{x}_* = \mathbf{W}^* \mathbf{x})$$



Advantage of Feature Learning


$$y = f^*(\mathbf{x}) = \mathbf{g}^*(\mathbf{x}_* = \mathbf{W}^* \mathbf{x})$$



**Information/
generative exponent
 ≤ 2**

Advantage of Feature Learning

$$y = f^*(\mathbf{x}) = \mathbf{g}^*(\mathbf{x}_* = \mathbf{W}^* \mathbf{x})$$





**Information/
generative exponent
 ≤ 2**



Advantage of Feature Learning

$$y = f^*(\mathbf{x}) = \mathbf{g}^*(\mathbf{x}_* = \mathbf{W}^* \mathbf{x})$$





$$\hat{y} = \hat{\mathbf{a}} \cdot \sigma(\hat{W} \mathbf{x})$$

**Information/
generative exponent
 ≤ 2**



Advantage of Feature Learning

$$y = f^*(\mathbf{x}) = \mathbf{g}^*(\mathbf{x}_* = \mathbf{W}^* \mathbf{x})$$



$$\hat{y} = \hat{\mathbf{a}} \cdot \sigma(\hat{\mathbf{W}} \mathbf{x})$$

**Information/
generative exponent
 ≤ 2**





With $\tilde{O}(d)$ samples, GD/SGD on $\hat{\mathbf{W}}$ recovers

$$\hat{\mathbf{W}} \approx Z_1 \mathbf{W}^* + Z_2$$

Advantage of Feature Learning

$$y = f^*(\mathbf{x}) = \mathbf{g}^*(\mathbf{x}_* = \mathbf{W}^* \mathbf{x})$$



$$\hat{y} = \hat{\mathbf{a}} \cdot \sigma(\hat{\mathbf{W}} \mathbf{x})$$

**Information/
generative exponent
 ≤ 2**



With $\tilde{O}(d)$ samples, GD/SGD on $\hat{\mathbf{W}}$ recovers



$$\hat{\mathbf{W}} \approx Z_1 W^* + Z_2$$

Then the neural net is equivalent to

$$\hat{y} \approx \hat{\mathbf{a}} \cdot \sigma((Z_1 W^* + Z_2) \mathbf{x}) = \hat{\mathbf{a}} \cdot \sigma(Z_1 \mathbf{x}^* + Z_3)$$

Advantage of Feature Learning

$$y = f^*(\mathbf{x}) = \mathbf{g}^*(\mathbf{x}_* = \mathbf{W}^* \mathbf{x})$$



$$\hat{y} = \hat{\mathbf{a}} \cdot \sigma(\hat{\mathbf{W}} \mathbf{x})$$

**Information/
generative exponent
 ≤ 2**



With $\tilde{O}(d)$ samples, GD/SGD on $\hat{\mathbf{W}}$ recovers


$$\hat{\mathbf{W}} \approx \mathbf{Z}_1 \mathbf{W}^* + \mathbf{Z}_2$$

Then the neural net is equivalent to

$$\hat{y} \approx \hat{\mathbf{a}} \cdot \sigma((\mathbf{Z}_1 \mathbf{W}^* + \mathbf{Z}_2) \mathbf{x}) = \hat{\mathbf{a}} \cdot \sigma(\mathbf{Z}_1 \mathbf{x}^* + \mathbf{Z}_3)$$

Random feature in (finite)
reduced space
 $d \rightarrow d^{\text{eff}} = r = O(1)$


Dimensionality reduction



Sample Complexity Reduction
 $\mathcal{O}(d^k) \rightarrow \tilde{\mathcal{O}}(d)$

Dimensionality reduction


searching in a $\mathcal{O}(d^k)$
dimensional subspace of
polynomials



Sample Complexity Reduction
 $\mathcal{O}(d^k) \rightarrow \tilde{\mathcal{O}}(d)$

Dimensionality reduction

searching in a $\mathcal{O}(d^k)$
dimensional subspace of
polynomials




Sample Complexity Reduction
 $\mathcal{O}(d^k) \rightarrow \tilde{\mathcal{O}}(d)$

Dimensionality reduction

searching in a $\mathcal{O}(d^k)$
dimensional subspace of
polynomials



searching in a $\mathcal{O}(r^k)$
dimensional subspace of
polynomials




Sample Complexity Reduction
 $\mathcal{O}(d^k) \rightarrow \tilde{\mathcal{O}}(d)$

Dimensionality reduction

searching in a $\mathcal{O}(d^k)$
dimensional subspace of
polynomials




searching in a $\mathcal{O}(r^k)$
dimensional subspace of
polynomials



Too easy...no further
advantage of depth

Sample Complexity Reduction
 $\mathcal{O}(d^k) \rightarrow \tilde{\mathcal{O}}(d)$




**Can we generalize this picture to
arbitrary depth?**

A SIGHT beyond Two Layers (Single-Index Gaussian Hierarchical Targets)

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

P_k : polynomial of degree k
 g^\star : non-linearity




high-dimensional Gaussian
 $\mathbf{x} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_d)$

A SIGHT beyond Two Layers (Single-Index Gaussian Hierarchical Targets)

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

P_k : polynomial of degree k
 g^\star : non-linearity




A SIGHT beyond Two Layers (Single-Index Gaussian Hierarchical Targets)

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\epsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

P_k : polynomial of degree k
 g^\star : non-linearity


Subspace dimension d^ϵ
for $0 < \epsilon < 1$



A SIGHT beyond Two Layers (Single-Index Gaussian Hierarchical Targets)

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

P_k : polynomial of degree k
 g^\star : non-linearity



Subspace dimension d^ϵ
for $0 < \epsilon < 1$

A SIGHT beyond Two Layers (Single-Index Gaussian Hierarchical Targets)

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\epsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

P_k : polynomial of degree k
 g^\star : non-linearity

$$f^\star(\mathbf{x}) = g^\star(h^\star)$$




\mathbf{a}^\star

$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\epsilon}} \in \mathbb{R}$$

Subspace dimension d^ϵ
for $0 < \epsilon < 1$

$$\mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\epsilon}$$


$$\mathbf{x} \in \mathbb{R}^d$$



MIGHT (Multi-Index Gaussian Hierarchical Targets)

$$f^\star(\mathbf{x}) = g^\star(h_1^\star(\mathbf{x}), \dots, h_r^\star(\mathbf{x})), \mathbf{x} \in \mathbb{R}^d$$

$$h_m^\star(\mathbf{x}) = \frac{1}{\sqrt{\mathbf{d}^\varepsilon}} \mathbf{a}_m^{\star\top} \mathbf{P}_{k,m} (\mathbf{W}_m^\star \mathbf{x}), m = 1, \dots, r$$



Key features of SIGHTS/MIGHTS

Key features of SIGHTS/MIGHTS

- SIGHT/MIGHT are simply three-layer neural networks.

Key features of SIGHTS/MIGHTS

- SIGHT/MIGHT are simply three-layer neural networks.
- Dimension of features reduces from d to d^ϵ to $\mathcal{O}(1)$.

Key features of SIGHTS/MIGHTS

- SIGHT/MIGHT are simply three-layer neural networks.
- Dimension of features reduces from d to d^ϵ to $\mathcal{O}(1)$.
- By CLT, h^* asymptotically Gaussian and independent.

Key features of SIGHTS/MIGHTS

- SIGHT/MIGHT are simply three-layer neural networks.
- Dimension of features reduces from d to d^ϵ to $\mathcal{O}(1)$.
- By CLT, h^* asymptotically Gaussian and independent.
- Motivated by Nichani et al. 2023:

Key features of SIGHTS/MIGHTS

- SIGHT/MIGHT are simply three-layer neural networks.
- Dimension of features reduces from d to d^ϵ to $\mathcal{O}(1)$.
- By CLT, h^* asymptotically Gaussian and independent.
- Motivated by Nichani et al. 2023:

$$f^\star(\mathbf{x}) = \mathbf{g}^\star(\mathbf{x}^\top \mathbf{A} \mathbf{x})$$

Key features of SIGHTS/MIGHTS

- SIGHT/MIGHT are simply three-layer neural networks.
- Dimension of features reduces from d to d^ϵ to $\mathcal{O}(1)$.
- By CLT, h^* asymptotically Gaussian and independent.
- Motivated by Nichani et al. 2023:

$$f^\star(\mathbf{x}) = \mathbf{g}^\star(\mathbf{x}^\top \mathbf{A} \mathbf{x}) \longrightarrow$$

Key features of SIGHTS/MIGHTS


- SIGHT/MIGHT are simply three-layer neural networks.
- Dimension of features reduces from d to d^ϵ to $\mathcal{O}(1)$.
- By CLT, h^* asymptotically Gaussian and independent.
- Motivated by Nichani et al. 2023:

$$f^\star(\mathbf{x}) = \mathbf{g}^\star(\mathbf{x}^\top \mathbf{A} \mathbf{x}) \longrightarrow$$

2-layer NN can't learn
non-linear features

1 level of dimension-
reduction


No adaptivity of first layer
required



Learning SIGHT with Two-layer NNs

Lazy learning (Random Features/NTK)

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$



Learning with a two-layer net

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\varepsilon}$$

$$\hat{y} = \hat{W}_2 \sigma(\hat{W}_1 \mathbf{x}))$$


Learning with a two-layer net

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\varepsilon}$$

parameter count of W^\star

$$\hat{y} = \hat{W}_2 \sigma(\hat{W}_1 \mathbf{x})$$



With $n = \mathcal{O}(d^{1+\varepsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$


Learning with a two-layer net

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\varepsilon}$$

parameter count of W^\star

$$\hat{y} = \hat{W}_2 \sigma(\hat{W}_1 \mathbf{x})$$



With $n = \mathcal{O}(d^{1+\varepsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$

$$\hat{y} \approx \hat{W}_2 \sigma((Z_1 W^\star + Z_2) \mathbf{x}) \approx \hat{W}_2 \sigma(Z_1 x_\star + Z_3)$$


Learning with a two-layer net

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\epsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\epsilon}} \right), \mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\epsilon}$$

parameter count of W^\star

$$\hat{y} = \hat{W}_2 \sigma(\hat{W}_1 \mathbf{x})$$



With $n = \mathcal{O}(d^{1+\epsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$

$$\hat{y} \approx \hat{W}_2 \sigma((Z_1 W^\star + Z_2) \mathbf{x}) \approx \hat{W}_2 \sigma(Z_1 x_\star + Z_3)$$

random feature/lazy learning but in reduced dimension \mathbb{R}^{d^ϵ}


Learning with a two-layer net

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\epsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\epsilon}} \right), \mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\epsilon}$$

parameter count of W^\star

$$\hat{y} = \hat{W}_2 \sigma(\hat{W}_1 \mathbf{x})$$



With $n = \mathcal{O}(d^{1+\epsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$


$$\hat{y} \approx \hat{W}_2 \sigma((Z_1 W^\star + Z_2) \mathbf{x}) \approx \hat{W}_2 \sigma(Z_1 x_\star + Z_3)$$

random feature/lazy learning but in reduced dimension \mathbb{R}^{d^ϵ}

Adaptive Learning with a two-layer net

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^{\varepsilon_1}}} \right), \mathbf{x}_\star \in \mathbb{R}^{d^{\varepsilon_1}}$$





Learning SIGHT with Three-layer NNs

Idealized scenario for layer-wise training in 3-layer NNs

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star(h^\star)$$

$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$\mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^c}$$

$$W^\star \in \mathbb{R}^{d^c \times d}$$

$$\mathbf{x} \in \mathbb{R}^d$$

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$\mathbf{w}_3^\top \in \mathbb{R}^{p_2}$$

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}$$

$$W_1 \mathbf{x} \in \mathbb{R}^{p_1}$$

$$\mathbf{x} \in \mathbb{R}^d$$

$$W_2 \in \mathbb{R}^{p_2 \times p_1}$$

$$W_1 \in \mathbb{R}^{p_1 \times d}$$

Idealized scenario for layer-wise training in 3-layer NNs

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star(h^\star)$$

$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$\mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^c}$$

$$\mathbf{x} \in \mathbb{R}^d$$

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$\mathbf{w}_3^\top \in \mathbb{R}^{p_2}$$

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}$$

$$W_1 \mathbf{x} \in \mathbb{R}^{p_1}$$

$$\mathbf{x} \in \mathbb{R}^d$$

Idealized scenario for layer-wise training in 3-layer NNs

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star(h^\star)$$

$$\mathbf{a}^\star$$

$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$W^\star \in \mathbb{R}^{d^e \times d^c}$$

$$\mathbf{x} \in \mathbb{R}^d$$

$$\mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^c}$$


$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$\mathbf{w}_3^\top \in \mathbb{R}^{p_2}$$

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}$$

$$W_1 \mathbf{x} \in \mathbb{R}^{p_1}$$

$$\mathbf{x} \in \mathbb{R}^d$$



Idealized scenario for layer-wise training in 3-layer NNs

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star(h^\star)$$

$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$


\mathbf{a}^\star

$W^\star \in \mathbb{R}^{d^\varepsilon \times d}$

$\mathbf{x} \in \mathbb{R}^d$

W_2 learns features not weights a^\star

$$\mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\varepsilon}$$



$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$\mathbf{x} \in \mathbb{R}^d$

Idealized scenario for layer-wise training in 3-layer NNs

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star(h^\star)$$

$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

\mathbf{a}^\star

$W^\star \in \mathbb{R}^{d^\varepsilon \times d}$

$\mathbf{x} \in \mathbb{R}^d$

W_2 learns features not weights a^\star

$$\mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\varepsilon}$$



$W_2 \in \mathbb{R}^{p_2 \times p_1}$

$W_1 \in \mathbb{R}^{p_1 \times d}$

$\mathbf{x} \in \mathbb{R}^d$

$$\mathbf{w}_3^\top \in \mathbb{R}^{p_2}$$

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}$$

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x}) + b_1) + b_2$$

Idealized scenario for layer-wise training in 3-layer NNs

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star(h^\star)$$

$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

\mathbf{a}^\star

$W^\star \in \mathbb{R}^{d^\varepsilon \times d}$

$$\mathbf{x}_\star = W^\star \mathbf{x} \in \mathbb{R}^{d^\varepsilon}$$

$$\mathbf{x} \in \mathbb{R}^d$$

W_2 learns features not weights a^\star



$W_2 \in \mathbb{R}^{p_2 \times p_1}$

$W_1 \in \mathbb{R}^{p_1 \times d}$

$$\mathbf{w}_3^\top \in \mathbb{R}^{p_2}$$

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}$$

$$W_1 \mathbf{x} \in \mathbb{R}^{p_1}$$


$$\mathbf{x} \in \mathbb{R}^d$$

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}$$

Idealized scenario for layer-wise training in 3-layer NNs

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^{\star\mathbf{x}})}{\sqrt{d^{\varepsilon_1}}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$



Expected behavior under idealized scenario

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star \in \mathbb{R}^{d^\epsilon}$$

$$\hat{y} = \hat{\mathbf{W}}^3 \sigma(\hat{W}_2 \sigma(\hat{W}_1 \mathbf{x}))$$

Expected behavior under idealized scenario

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star \in \mathbb{R}^{d^\epsilon}$$

$$\hat{y} = \hat{\mathbf{W}}^3 \sigma(\hat{W}_2 \sigma(\hat{W}_1 \mathbf{x}))$$

With $n = \mathcal{O}(d^{1+\epsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$

$$\hat{y} \approx \hat{\mathbf{w}}_3^\top \sigma(\hat{W}_2 \sigma(Z_1 x_\star + Z_2))$$

Expected behavior under idealized scenario

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star \in \mathbb{R}^{d^\varepsilon}$$

$$\hat{y} = \hat{\mathbf{W}}^3 \sigma(\hat{W}_2 \sigma(\hat{W}_1 \mathbf{x}))$$

With $n = \mathcal{O}(d^{1+\varepsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$

With each $n = \mathcal{O}(d^\varepsilon)$ data, $\hat{W}_2 \sigma(\cdot)$ learns a higher-degree component of P_k

$$\hat{y} \approx \hat{\mathbf{w}}_3^\top \sigma(\hat{W}_2 \sigma(Z_1 x_\star + Z_2))$$

Expected behavior under idealized scenario

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^{\varepsilon_1}}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(x_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star \in \mathbb{R}^{d^\varepsilon}$$

$$\hat{y} = \hat{\mathbf{W}}^3 \sigma(\hat{W}_2 \sigma(\hat{W}_1 \mathbf{x}))$$

With $n = \mathcal{O}(d^{1+\varepsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$

With each $n = \mathcal{O}(d^\varepsilon)$ data, $\hat{W}_2 \sigma(\cdot)$ learns a higher-degree component of P_k

$$\hat{y} \approx \hat{\mathbf{w}}_3^\top \sigma(\hat{W}_2 \sigma(Z_1 x_\star + Z_2))$$

$$\approx \hat{\mathbf{w}}_3^\top \sigma(Z_1 h^\star + Z_2)$$

Expected behavior under idealized scenario

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d$$

$$f^\star(\mathbf{x}) = g^\star(h^\star), h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$\hat{y} = \hat{\mathbf{W}}^3 \sigma(\hat{W}_2 \sigma(\hat{W}_1 \mathbf{x}))$$


With $n = \mathcal{O}(d^{1+\varepsilon})$ data, \hat{W}_1 can learn the features with GD : $\hat{W}_1 \approx Z_1 W^\star + Z_2$

With each $n = \mathcal{O}(d^\varepsilon)$ data, $\hat{W}_2 \sigma(\cdot)$ learns a higher-degree component of P_k

$$\hat{y} \approx \hat{\mathbf{w}}_3^\top \sigma(\hat{W}_2 \sigma(Z_1 x_\star + Z_2))$$


$$\approx \hat{\mathbf{w}}_3^\top \sigma(Z_1 h^\star + Z_2)$$

Two levels of dimension reduction




Two levels of dimension reduction

$$f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{W}^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$




Two levels of dimension reduction

$$f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{W}^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x} \in \mathbb{R}^d$$




Two levels of dimension reduction

$$f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{W}^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d \quad \longrightarrow \quad f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star \in \mathbb{R}^{d_\epsilon}$$




Two levels of dimension reduction

$$f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{W}^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d \rightarrow f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star \in \mathbb{R}^{d_\varepsilon}$$




Two levels of dimension reduction


$$f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{W}^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \mathbf{x} \in \mathbb{R}^d \rightarrow f^\star(\mathbf{x}) = \mathbf{g}^\star \left(\frac{\mathbf{a}^{\star\top} \mathbf{P}_k (\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \right), \mathbf{x}_\star \in \mathbb{R}^{d_\varepsilon} \rightarrow f^\star(\mathbf{x}) = g^\star(h^\star)$$



Deep version




Deep version




**Tree-structure maintains independence
of features**

Deep version




Iterative dimensionality reduction

$$d^{\epsilon_1} \rightarrow d^{\epsilon_2} \rightarrow d^{\epsilon_3}, \dots, \rightarrow 1$$



Tree-structure maintains independence of features


Deep version



Iterative dimensionality reduction


$$d^{\epsilon_1} \rightarrow d^{\epsilon_2} \rightarrow d^{\epsilon_3}, \dots, \rightarrow 1$$

L-layer network



Tree-structure maintains independence of features


Deep version



Iterative dimensionality reduction


$$d^{\epsilon_1} \rightarrow d^{\epsilon_2} \rightarrow d^{\epsilon_3}, \dots, \rightarrow 1$$

L-layer network



Tree-structure maintains independence of features

Deep version



Iterative dimensionality reduction

$$d^{\epsilon_1} \rightarrow d^{\epsilon_2} \rightarrow d^{\epsilon_3}, \dots, \rightarrow 1$$

L-layer network

Tree-structure maintains independence of features

L-1 levels of dimension reduction

Conditions for learnability

Conditions for learnability

Information exponent

$$\text{IE}(\ell) = \inf\{k : \|\mathbb{E}[(\mathbf{x})^{\otimes k} f^\star(\mathbf{x})]\|_F \neq 0\}$$

Conditions for learnability

Information exponent



$$\text{IE}(\ell) = \inf\{k : \|\mathbb{E}[(\mathbf{x})^{\otimes k} f^\star(\mathbf{x})]\|_F \neq 0\}$$

Conditions for learnability

Information exponent



Compositional
Information exponent

$$\text{IE}(\ell) = \inf\{k : \|\mathbb{E}[(\mathbf{x})^{\otimes k} f^\star(\mathbf{x})]\|_F \neq 0\}$$

$$\text{CIE}(\ell) = \inf\{k : \|\mathbb{E}[(h_\ell(\mathbf{x}))^{\otimes k} f^\star(\mathbf{x})]\|_F = \Theta(1)\}$$

Conditions for learnability

Information exponent



Compositional
Information exponent

$$\text{IE}(\ell) = \inf\{k : \|\mathbb{E}[(\mathbf{x})^{\otimes k} f^\star(\mathbf{x})]\|_F \neq 0\}$$

$$\text{CIE}(\ell) = \inf\{k : \|\mathbb{E}[(h_\ell(\mathbf{x}))^{\otimes k} f^\star(\mathbf{x})]\|_F = \Theta(1)\}$$

Requires non-trivial low-degree correlations
between intermediate features and labels

Conditions for learnability

Information exponent



Compositional
Information exponent

$$\text{IE}(\ell) = \inf\{k : \|\mathbb{E}[(\mathbf{x})^{\otimes k} f^\star(\mathbf{x})]\|_F \neq 0\}$$

$$\text{CIE}(\ell) = \inf\{k : \|\mathbb{E}[(h_\ell(\mathbf{x}))^{\otimes k} f^\star(\mathbf{x})]\|_F = \Theta(1)\}$$

Requires non-trivial low-degree correlations
between intermediate features and labels

Matches behavior of
real data

	Cal-101 (30/class)	Cal-256 (60/class)
SVM (1)	44.8 ± 0.7	24.6 ± 0.4
SVM (2)	66.2 ± 0.5	39.6 ± 0.3
SVM (3)	72.3 ± 0.4	46.0 ± 0.3
SVM (4)	76.6 ± 0.4	51.3 ± 0.1
SVM (5)	86.2 ± 0.8	65.6 ± 0.3
SVM (7)	85.5 ± 0.4	71.7 ± 0.2
Softmax (5)	82.9 ± 0.4	65.7 ± 0.5
Softmax (7)	85.4 ± 0.4	72.6 ± 0.1

Table 7. Analysis of the discriminative information contained in each layer of feature maps within our ImageNet-pretrained convnet. We train either a linear SVM or softmax on features from different layers (as indicated in brackets) from the convnet. Higher layers generally produce more discriminative features.

Full set of conditions

Essential assumptions

Technical assumptions

Full set of conditions

Essential assumptions

- g^* : information exponent 1.
 - P_k : information exponent/leap
 ≤ 2
- $$\text{CIE}(1) = IE(g^*) \times IE(P_k)$$
- expressive σ (non-zero Hermites) and regular.

Technical assumptions

Full set of conditions

Essential assumptions

- g^* : information exponent 1.
- P_k : information exponent/leap
 ≤ 2
$$\text{CIE}(1) = IE(g^*) \times IE(P_k)$$
- expressive σ (non-zero Hermites) and regular.

Technical assumptions

- $a_i^* = 1 \forall i$ (symmetric targets)
- Correlation loss
- Re-initialization of layers
- expressive (non-zero Hermites) and regular.
- P_k : information exponent $\neq 1$ (to avoid spikes)
- $\mathbb{E}[\sigma(\sigma(z))\text{He}_2(z)]\mathbb{E}[P_k(z)\text{He}_2(z)] > 0$,
 $\mathbb{E}[\sigma(\sigma(z))z] = 0$,
- $\mathbb{E}_{z \sim \mathcal{N}(0,1)}[g^*(z)\text{He}_j(z) = 0, 1 < j \leq k]$

Main theorem

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{p_3}$$

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{p_3}$$

Let $\delta > 0$ be arbitrarily small

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{p_3}$$

Let $\delta > 0$ be arbitrarily small

- First layer: $\mathcal{O}(\log d)$ steps of spherical SGD on W_1 with correlation loss, batch-size $\mathcal{O}(d^{1+\epsilon+\delta})$, vanishing step-size:

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{\mathbf{p}_3}$$

Let $\delta > 0$ be arbitrarily small

- First layer: $\mathcal{O}(\log d)$ steps of spherical SGD on W_1 with correlation loss, batch-size $\mathcal{O}(d^{1+\epsilon+\delta})$, vanishing step-size:

$$W_1 = Z(W^\star) + o_d(1), z_i \sim U(S_1)$$

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{p_3}$$

Let $\delta > 0$ be arbitrarily small

- First layer: $\mathcal{O}(\log d)$ steps of spherical SGD on W_1 with correlation loss, batch-size $\mathcal{O}(d^{1+\epsilon+\delta})$, vanishing step-size:

$$W_1 = Z(W^\star) + o_d(1), z_i \sim U(S_1)$$

- Second layer: upon reinitializing $W_2 = \mathbf{0}_{d \times d}$, one step of pre-conditioned step with batch-size $\mathcal{O}(d^{k\epsilon+\delta})$, $p_1 = \mathcal{O}(d^{k\epsilon+\delta})$:

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{\mathbf{p}_3}$$

Let $\delta > 0$ be arbitrarily small

- First layer: $\mathcal{O}(\log d)$ steps of spherical SGD on W_1 with correlation loss, batch-size $\mathcal{O}(d^{1+\epsilon+\delta})$, vanishing step-size:

$$W_1 = Z(W^\star) + o_d(1), z_i \sim U(S_1)$$

- Second layer: upon reinitializing $W_2 = \mathbf{0}_{d \times d}$, one step of pre-conditioned step with batch-size $\mathcal{O}(d^{k\epsilon+\delta})$, $p_1 = \mathcal{O}(d^{k\epsilon+\delta})$:

$$\mathbf{h}_i(\mathbf{x}) = \mathbf{W}_2 \sigma(\mathbf{W}_1 \mathbf{x}) = \mathbf{c} \mathbf{w}_3^\top \mathbf{h}^\star(\mathbf{x}) + \mathbf{o}_d(1)$$

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{p_3}$$

Let $\delta > 0$ be arbitrarily small

- First layer: $\mathcal{O}(\log d)$ steps of spherical SGD on W_1 with correlation loss, batch-size $\mathcal{O}(d^{1+\epsilon+\delta})$, vanishing step-size:

$$W_1 = Z(W^\star) + o_d(1), z_i \sim U(S_1)$$

- Second layer: upon reinitializing $W_2 = \mathbf{0}_{d \times d}$, one step of pre-conditioned step with batch-size $\mathcal{O}(d^{k\epsilon+\delta})$, $p_1 = \mathcal{O}(d^{k\epsilon+\delta})$:

$$\mathbf{h}_i(\mathbf{x}) = \mathbf{W}_2 \sigma(\mathbf{W}_1 \mathbf{x}) = \mathbf{c} \mathbf{w}_3^\top \mathbf{h}^\star(\mathbf{x}) + \mathbf{o}_d(1)$$

- Third layer: Ridge regression on \mathbf{w}_3 with samples $\mathcal{O}(d^\delta)$, $p_2 = \mathcal{O}(d^\delta)$:

Main theorem

$$\hat{f}(\mathbf{x}) = \mathbf{w}_3^\top \sigma(W_2 \sigma(W_1 \mathbf{x} + b_1) + b_2)$$

$$W_1 \in \mathbb{R}^{p_1 \times d}, W_2 \in \mathbb{R}^{p_2 \times p_1}, \mathbf{w}_3 \in \mathbb{R}^{p_3}$$

Let $\delta > 0$ be arbitrarily small

- First layer: $\mathcal{O}(\log d)$ steps of spherical SGD on W_1 with correlation loss, batch-size $\mathcal{O}(d^{1+\epsilon+\delta})$, vanishing step-size:

$$W_1 = Z(W^\star) + o_d(1), z_i \sim U(S_1)$$

- Second layer: upon reinitializing $W_2 = \mathbf{0}_{d \times d}$, one step of pre-conditioned step with batch-size $\mathcal{O}(d^{k\epsilon+\delta})$, $p_1 = \mathcal{O}(d^{k\epsilon+\delta})$:


$$\mathbf{h}_i(\mathbf{x}) = \mathbf{W}_2 \sigma(\mathbf{W}_1 \mathbf{x}) = \mathbf{c} \mathbf{w}_3^\top \mathbf{h}^\star(\mathbf{x}) + \mathbf{o}_d(1)$$

- Third layer: Ridge regression on \mathbf{w}_3 with samples $\mathcal{O}(d^\delta)$, $p_2 = \mathcal{O}(d^\delta)$:


$$\hat{f}(\mathbf{x}) = \mathbf{f}^\star(\mathbf{x}) + \mathbf{o}_d(1)$$

General Depth


General Depth



General Depth




General Depth



- Conditional on perfect spherical recovery for W_{L-2} , the same picture holds for the last two layers.
- Key idea: Features are independent by tree-structure, asymptotically Gaussian, and maintain nice tails (hypercontractivity).

General Depth



- Conditional on perfect spherical recovery for W_{L-2} , the same picture holds for the last two layers.
- Key idea: Features are independent by tree-structure, asymptotically Gaussian, and maintain nice tails (hypercontractivity).



Theorem 2. For $L \in \mathbb{N}$, let $f^\star(\mathbf{x})$ denote a target as in Eq. (6) with $r = 1$, and let δ', δ be arbitrary reals satisfying $0 < \delta < \delta' < 1$. Consider a model of the form $\hat{f}_\theta(\mathbf{x}) = \mathbf{w}_L^\top \sigma(W_{L-1} \sigma(W h_{L-1}^\star(\mathbf{x})))$ with $W \in \mathbb{R}^{p_{L-2} \times d^{\varepsilon_{L-2}}}$ having $p_{L-2} = \Theta(d^{k\varepsilon_{L-2} + \delta'})$ rows independently sampled as $\mathbf{w}_i \sim U(\mathcal{S}_{d^{\varepsilon_{L-2}}}(\mathbf{1}))$. Under Ass. 1-3, after a single step of pre-conditioned SGD on W_{L-1} with batch-size $\Theta(d^{k\varepsilon_{L-2} + \delta})$, step-size $\Theta(\sqrt{p_{L-1}})$, the pre-activations $h_{L-1}(\mathbf{x}) := W_{L-1} \sigma(W h_{L-1}^\star(\mathbf{x}))$ satisfy, for a constant $c > 0$:

$$h_{L-1}(\mathbf{x}) = c \mathbf{w}_L h_L^\star(\mathbf{x}) + o_d(1), \quad (22)$$

Example

$$f^*(\mathbf{x}) = \tanh \left(\frac{\mathbf{a}^{*\top} P_3(W^* \mathbf{x})}{\sqrt{d^{\varepsilon_1=1/2}}} \right)$$

$$P_3(x) = \text{He}_2(x) + \text{He}_3(x)$$





Overlaps in parameter and function space

$$f^*(\mathbf{x}) = \tanh \left(\frac{\mathbf{a}^{*\top} P_3(W^* \mathbf{x})}{\sqrt{d^{\varepsilon_1=1/2}}} \right)$$

$$P_3(x) = \text{He}_2(x) + \text{He}_3(x)$$

$$M_W = \frac{W_1 W^*}{\|W_1\|_2},$$

$$M_h = \frac{\mathbb{E}[\mathbf{h}(\mathbf{z})\mathbf{h}^*(\mathbf{z})]}{\sqrt{\mathbb{E}[\mathbf{h}(\mathbf{z})^2]}}.$$





Key ideas and proof sketches

Recursive Expansion of the Target

Recursive Expansion of the Target

$$f^\star(\mathbf{x}) = g^\star\left(\frac{\mathbf{a}^{\star^\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}}\right), \quad \mathbf{x} \in \mathbb{R}^d$$

Recursive Expansion of the Target

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x} \in \mathbb{R}^d \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

Recursive Expansion of the Target

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x} \in \mathbb{R}^d \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$g^\star(h_\star(\mathbf{x})) \approx \mu_1 h^\star(\mathbf{x}) + \mu_2^\star \text{He}_2(h^\star(\mathbf{x})) + \dots$$

Recursive Expansion of the Target

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}^{\star\top} P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \right), \quad \mathbf{x} \in \mathbb{R}^d \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$g^\star(h_\star(\mathbf{x})) \approx \mu_1 h^\star(\mathbf{x}) + \mu_2^\star \text{He}_2(h^\star(\mathbf{x})) + \dots$$

!!! $\text{He}_k(h^\star(\mathbf{x}))$ can contribute low-degree terms

Composition of Hermites

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{\mathbf{d}^\epsilon}} \sum_{i=1}^{\mathbf{d}^\epsilon} \text{He}_k(\langle \mathbf{w}_i^\star, \mathbf{x} \rangle)$$

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{d^\epsilon}} \sum_{i=1}^{d^\epsilon} \text{He}_k(\langle \mathbf{w}_i^*, \mathbf{x} \rangle)$$


$$\text{He}_m(h(\mathbf{x})) \approx \frac{1}{\sqrt{md^{m\epsilon_1}}} \sum_{\text{distinct subsets } s_i} \prod \text{He}_k(\langle \mathbf{w}_{s_i}^*, \mathbf{x} \rangle)$$

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{d^\epsilon}} \sum_{i=1}^{d^\epsilon} \text{He}_k(\langle \mathbf{w}_i^*, \mathbf{x} \rangle)$$

$$\text{He}_m(h(\mathbf{x})) \approx \frac{1}{\sqrt{md^{m\epsilon_1}}} \sum_{\text{distinct subsets } s_i} \prod \text{He}_k(\langle \mathbf{w}_{s_i}^*, \mathbf{x} \rangle)$$

$$g^*(h_\star(\mathbf{x})) \approx \mu_1 \mathbf{h}^*(\mathbf{x}) + \mu_2^* \text{He}_2(\mathbf{h}^*(\mathbf{x})) + \dots$$




High Hermite-degree

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{d^\epsilon}} \sum_{i=1}^{d^\epsilon} \text{He}_k(\langle \mathbf{w}_i^*, \mathbf{x} \rangle)$$

$$\text{He}_m(h(\mathbf{x})) \approx \frac{1}{\sqrt{md^{m\epsilon_1}}} \sum_{\text{distinct subsets } s_i} \prod \text{He}_k(\langle \mathbf{w}_{s_i}^*, \mathbf{x} \rangle)$$

$$g^*(h_\star(\mathbf{x})) \approx \mu_1 \mathbf{h}^*(\mathbf{x}) + \mu_2^* \text{He}_2(\mathbf{h}^*(\mathbf{x})) + \dots$$




High Hermite-degree

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{d^\epsilon}} \sum_{i=1}^{d^\epsilon} \text{He}_k(\langle \mathbf{w}_i^*, \mathbf{x} \rangle)$$

$$\text{He}_m(h(\mathbf{x})) \approx \frac{1}{\sqrt{md^{m\epsilon_1}}} \sum_{\text{distinct subsets } s_i} \prod \text{He}_k(\langle \mathbf{w}_{s_i}^*, \mathbf{x} \rangle)$$

$$g^*(h_\star(\mathbf{x})) \approx \mu_1 \mathbf{h}^*(\mathbf{x}) + \mu_2^* \text{He}_2(\mathbf{h}^*(\mathbf{x})) + \dots$$




High Hermite-degree

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{d^\epsilon}} \sum_{i=1}^{d^\epsilon} \text{He}_k(\langle \mathbf{w}_i^*, \mathbf{x} \rangle)$$

$$\text{He}_m(h(\mathbf{x})) \approx \frac{1}{\sqrt{md^{m\epsilon_1}}} \sum_{\text{distinct subsets } s_i} \prod \text{He}_k(\langle \mathbf{w}_{s_i}^*, \mathbf{x} \rangle)$$

$$g^*(h^*(\mathbf{x})) \approx \mu_1 h^*(\mathbf{x}) + \mu_2^* \text{He}_2(h^*(\mathbf{x})) + \dots$$



High Hermite-degree

- The dominant term in the expansion is along $h^*(\mathbf{x})$.

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{d^\epsilon}} \sum_{i=1}^{d^\epsilon} \text{He}_k(\langle \mathbf{w}_i^*, \mathbf{x} \rangle)$$

$$\text{He}_m(h(\mathbf{x})) \approx \frac{1}{\sqrt{md^{m\epsilon_1}}} \sum_{\text{distinct subsets } s_i} \prod_{s_i} \text{He}_k(\langle \mathbf{w}_{s_i}^*, \mathbf{x} \rangle)$$

$$g^*(h^*(\mathbf{x})) \approx \mu_1 h^*(\mathbf{x}) + \mu_2^* \text{He}_2(h^*(\mathbf{x})) + \dots$$



High Hermite-degree


- The dominant term in the expansion is along $h^*(\mathbf{x})$.
- Generalization of the “approximate Stein’s Lemma” in Nichani et al. 2023.

Composition of Hermites

$$h(\mathbf{x}) = \frac{1}{\sqrt{d^\epsilon}} \sum_{i=1}^{d^\epsilon} \text{He}_k(\langle \mathbf{w}_i^*, \mathbf{x} \rangle)$$

$$\text{He}_m(h(\mathbf{x})) \approx \frac{1}{\sqrt{md^{m\epsilon_1}}} \sum_{\text{distinct subsets } s_i} \prod_{s_i} \text{He}_k(\langle \mathbf{w}_{s_i}^*, \mathbf{x} \rangle)$$

$$g^*(h^*(\mathbf{x})) \approx \mu_1 h^*(\mathbf{x}) + \mu_2^* \text{He}_2(h^*(\mathbf{x})) + \dots$$



High Hermite-degree

- The dominant term in the expansion is along $h^*(\mathbf{x})$.
- Generalization of the “approximate Stein’s Lemma” in Nichani et al. 2023.

Recovery by the first layer

Recovery by the first layer

$$g^\star(h_\star(\mathbf{x})) \approx \mu_1 \mathbf{h}^\star(\mathbf{x}) + \mu_2^\star \mathbf{He}_2(\mathbf{h}^\star(\mathbf{x})) + \dots$$


Recovery by the first layer

$$g^\star(h_\star(\mathbf{x})) \approx \mu_1 \mathbf{h}^\star(\mathbf{x}) + \mu_2^\star \mathbf{He}_2(\mathbf{h}^\star(\mathbf{x})) + \dots$$

Updates to $W_1 \approx \text{SGD}$ on
$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

Recovery by the first layer


$$g^\star(h_\star(\mathbf{x})) \approx \mu_1 h^\star(\mathbf{x}) + \mu_2^\star \text{He}_2(h^\star(\mathbf{x})) + \dots$$



Updates to $W_1 \approx \text{SGD}$ on
$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(W^\star \mathbf{x})}{\sqrt{d\varepsilon}} \in \mathbb{R}$$

Recovery by the first layer

$$g^*(h_\star(\mathbf{x})) \approx \mu_1 h^*(\mathbf{x}) + \mu_2^* \text{He}_2(h^*(\mathbf{x})) + \dots$$




Updates to $W_1 \approx \text{SGD}$ on
$$h^* = \mathbf{a}^* \cdot \frac{P_k(W^*\mathbf{x})}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

Gradient dominated by initial direction

Recovery by the first layer

$$g^*(h_\star(\mathbf{x})) \approx \mu_1 h^*(\mathbf{x}) + \mu_2^* \text{He}_2(h^*(\mathbf{x})) + \dots$$



Gradient dominated by initial direction


Updates to $W_1 \approx \text{SGD on}$

$$h^* = \mathbf{a}^* \cdot \frac{P_k(W^* \mathbf{x})}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$\frac{d\mathbf{w}_i}{dt} \propto u_i^* + \text{noise}$$

Recovery by the first layer

$$g^*(h_\star(\mathbf{x})) \approx \mu_1 h^*(\mathbf{x}) + \mu_2^* \text{He}_2(h^*(\mathbf{x})) + \dots$$



Gradient dominated by initial direction

Updates to $W_1 \approx \text{SGD on}$


$$h^* = \mathbf{a}^* \cdot \frac{P_k(W^* \mathbf{x})}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$\frac{d\mathbf{w}_i}{dt} \propto u_i^* + \text{noise}$$

drift+martingale

Recovery by the first layer

$$g^\star(h_\star(\mathbf{x})) \approx \mu_1 h^\star(\mathbf{x}) + \mu_2^\star \text{He}_2(h^\star(\mathbf{x})) + \dots$$



Gradient dominated by initial direction

higher-order terms suppressed by vanishing specialization along w_i^* + vanishing step size.

Updates to $W_1 \approx \text{SGD on}$
$$h^\star = \mathbf{a}^\star \cdot \frac{P_k(W^\star \mathbf{x})}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$\frac{d\mathbf{w}_i}{dt} \propto u_i^\star + \text{noise}$$

drift+martingale

Recovery by the second layer

Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2\sigma(\mathbf{W}_1\mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$

Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2\sigma(W_1\mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$

Gradient updates

Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2\sigma(W_1\mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$

Gradient updates

\approx

Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2\sigma(W_1\mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$

Gradient updates

\approx

Projections on the conjugate
Kernel*

Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$

Gradient updates

\approx

Projections on the conjugate
Kernel*

$$\Delta W_2^i \propto w_i^3 \sigma(W_1 X^\top) f^\star(X) \odot \sigma'(\mathbf{h}_i^\mathbf{t}(\mathbf{X}))$$

Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$


Gradient updates

\approx

Projections on the conjugate
Kernel*

$$\Delta h_{i,2}^t(\mathbf{x}) \approx \langle \sigma(W_1 \mathbf{x}), \Delta W_i \rangle$$

$$\Delta W_2^i \propto w_i^3 \sigma(W_1 X^\top) f^\star(X) \odot \sigma'(\mathbf{h}_i^t(\mathbf{X}))$$



Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$


Gradient updates

\approx


Projections on the conjugate Kernel*

$$\Delta h_{i,2}^t(\mathbf{x}) \approx \langle \sigma(W_1 \mathbf{x}), \Delta W_i \rangle$$


$$\Delta W_2^i \propto w_i^3 \sigma(W_1 X^\top) f^\star(X) \odot \sigma'(\mathbf{h}_i^t(\mathbf{X}))$$



$$\Delta \mathbf{h}_i(\mathbf{x}) \propto w_i^3 \sigma(W_1 \mathbf{x})^\top \sigma(W_1 \mathbf{X}^\top) f^\star(\mathbf{X}) \odot \sigma'(\mathbf{h}_i^t(\mathbf{X}))$$



Conjugate Kernel



Perturbed target

Recovery by the second layer

$$\mathbf{h}(\mathbf{x}) = W_2 \sigma(W_1 \mathbf{x}) \in \mathbb{R}^{p_2}, X \in \mathbb{R}^{n \times d}$$


Gradient updates

\approx

Projections on the conjugate Kernel*

$$\Delta h_{i,2}^t(\mathbf{x}) \approx \langle \sigma(W_1 \mathbf{x}), \Delta W_i \rangle$$

$$\Delta W_2^i \propto w_i^3 \sigma(W_1 X^\top) f^\star(X) \odot \sigma'(\mathbf{h}_i^t(\mathbf{X}))$$



$$\Delta \mathbf{h}_i(\mathbf{x}) \propto w_i^3 \sigma(W_1 \mathbf{x})^\top \sigma(W_1 \mathbf{X}^\top) f^\star(\mathbf{X}) \odot \sigma'(\mathbf{h}_i^t(\mathbf{X}))$$

A large black curly brace spanning the width of the term $\sigma(W_1 \mathbf{x})^\top \sigma(W_1 \mathbf{X}^\top)$.

Conjugate Kernel

A large black curly brace spanning the width of the term $f^\star(\mathbf{X}) \odot \sigma'(\mathbf{h}_i^t(\mathbf{X}))$.

Perturbed target

$$\Delta \mathbf{h}_i(\mathbf{x}) \propto w_i^3 \sigma(W_1 \mathbf{x})^\top \sigma(W_1 \mathbf{X}^\top) f^\star(\mathbf{X})$$

Dimension reduction in action

Dimension reduction in action

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$

Dimension reduction in action

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \mathbf{u}_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}')$$

Dimension reduction in action

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \mathbf{u}_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}')$$

Dimension reduction in action


$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \mathbf{u}_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}')$$

Dimension reduction in action

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \mathbf{u}_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}')$$




$K_{CK}(\mathbf{x}, \mathbf{x}')$ before training W_1

Dimension reduction in action

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$


$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \mathbf{u}_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}')$$



$K_{CK}(\mathbf{x}, \mathbf{x}')$ before training W_1


Dimension reduction in action

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$



$K_{CK}(\mathbf{x}, \mathbf{x}')$ before training W_1

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \mathbf{u}_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}')$$




$K_{CK}(\mathbf{x}, \mathbf{x}')$ after training W_1

Dimension reduction in action

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sigma(\mathbf{W}_1 \mathbf{x})^\top \sigma(\mathbf{W}_1 \mathbf{x}')$$

$$K_{CK}(\mathbf{x}, \mathbf{x}') = \sum_{i=1}^{\infty} \mathbf{u}_i(\mathbf{x}) \mathbf{u}_i(\mathbf{x}')$$



$K_{CK}(\mathbf{x}, \mathbf{x}')$ before training W_1

- Conjugate Kernel ill-conditioned, $\lambda_k = \frac{1}{d^{\epsilon k}} \implies \mathcal{O}(d^{\epsilon k})$ steps for convergence
 $\implies \mathcal{O}(d^{2\epsilon k})$ sample complexity.

Caveats:

- Fix: Pre-conditioning: $W_2 = W_2 - \eta \left(\frac{1}{n} \sigma(W_1 X^\top) \sigma(W_1 X)^\top \right)^{-1} \nabla_{W_2} \mathcal{L}$

Spike+ bulk decomposition of Gram matrix

Spike+ bulk decomposition of Gram matrix

$$\mathbf{x}_\star = \mathbf{W}^\star \mathbf{x}$$

Spike+ bulk decomposition of Gram matrix

$$\mathbf{x}_\star = \mathbf{W}^\star \mathbf{x} \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

Spike+ bulk decomposition of Gram matrix

$$\mathbf{X}_\star = \mathbf{W}^\star \mathbf{x} \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$K \approx \kappa_0^2 \mathbf{1}\mathbf{1}^T + \kappa_1^2 \frac{X_\star X_\star^T}{d} + \kappa_2^2 \frac{H_2(X_\star) H_2(X_\star)^T}{d^2} + \dots + \kappa_k^2 \frac{H_k(X_\star) H_k(X_\star)^T}{d^k} + \kappa_{k+1}^2 \frac{H_{k+1}(X_\star) H_{k+1}(X_\star)^T}{d^{k+1}} + \dots$$

Spike+ bulk decomposition of Gram matrix

$$\mathbf{X}_\star = \mathbf{W}^\star \mathbf{x} \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$K \approx \kappa_0^2 \mathbf{1}\mathbf{1}^T + \kappa_1^2 \frac{X_\star X_\star^T}{d} + \kappa_2^2 \frac{H_2(X_\star) H_2(X_\star)^T}{d^2} + \dots + \kappa_k^2 \frac{H_k(X_\star) H_k(X_\star)^T}{d^k} + \kappa_{k+1}^2 \frac{H_{k+1}(X_\star) H_{k+1}(X_\star)^T}{d^{k+1}} + \dots$$

$$n = O(d^{k\epsilon + \delta})$$

Spike+ bulk decomposition of Gram matrix

$$\mathbf{x}_\star = \mathbf{W}^\star \mathbf{x} \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$K \approx \kappa_0^2 \mathbf{1}\mathbf{1}^T + \kappa_1^2 \frac{X_\star X_\star^T}{d} + \kappa_2^2 \frac{H_2(X_\star) H_2(X_\star)^T}{d^2} + \dots + \kappa_k^2 \frac{H_k(X_\star) H_k(X_\star)^T}{d^k} + \kappa_{k+1}^2 \frac{H_{k+1}(X_\star) H_{k+1}(X_\star)^T}{d^{k+1}} + \dots$$

$$n = O(d^{k\epsilon+\delta})$$


Generalization error of random features and kernel methods:
hypercontractivity and kernel matrix concentration

Song Mei*, Theodor Misiakiewicz†, Andrea Montanari†‡

Spike+ bulk decomposition of Gram matrix


$$\mathbf{x}_\star = \mathbf{W}^\star \mathbf{x} \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$K \approx \kappa_0^2 \mathbf{1}\mathbf{1}^T + \kappa_1^2 \frac{X_\star X_\star^T}{d} + \kappa_2^2 \frac{H_2(X_\star) H_2(X_\star)^T}{d^2} + \dots + \kappa_k^2 \frac{H_k(X_\star) H_k(X_\star)^T}{d^k} + \kappa_{k+1}^2 \frac{H_{k+1}(X_\star) H_{k+1}(X_\star)^T}{d^{k+1}} + \dots$$



Concentrates to informative spikes

$$n = O(d^{k\epsilon + \delta})$$



Noise/ Identity


Generalization error of random features and kernel methods:
hypercontractivity and kernel matrix concentration

Song Mei*, Theodor Misiakiewicz†, Andrea Montanari†‡

Spike+ bulk decomposition of Gram matrix


$$\mathbf{x}_\star = \mathbf{W}^\star \mathbf{x} \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$K \approx \kappa_0^2 \mathbf{1}\mathbf{1}^T + \kappa_1^2 \frac{X_\star X_\star^T}{d} + \kappa_2^2 \frac{H_2(X_\star) H_2(X_\star)^T}{d^2} + \dots + \kappa_k^2 \frac{H_k(X_\star) H_k(X_\star)^T}{d^k} + \kappa_{k+1}^2 \frac{H_{k+1}(X_\star) H_{k+1}(X_\star)^T}{d^{k+1}} + \dots$$



Concentrates to informative spikes

$$n = O(d^{k\epsilon + \delta})$$



Noise/ Identity

**dominant low-degree term along
 $h^\star(\mathbf{x})$**

Generalization error of random features and kernel methods:
hypercontractivity and kernel matrix concentration

Song Mei*, Theodor Misiakiewicz†, Andrea Montanari†‡

Spike+ bulk decomposition of Gram matrix

$$\mathbf{x}_\star = \mathbf{W}^\star \mathbf{x} \quad h^\star = \mathbf{a}^\star \cdot \frac{P_k(\mathbf{x}_\star)}{\sqrt{d^\varepsilon}} \in \mathbb{R}$$

$$K \approx \kappa_0^2 \mathbf{1}\mathbf{1}^T + \kappa_1^2 \frac{X_\star X_\star^T}{d} + \kappa_2^2 \frac{H_2(X_\star) H_2(X_\star)^T}{d^2} + \dots + \kappa_k^2 \frac{H_k(X_\star) H_k(X_\star)^T}{d^k} + \kappa_{k+1}^2 \frac{H_{k+1}(X_\star) H_{k+1}(X_\star)^T}{d^{k+1}} + \dots$$



$$n = O(d^{k\epsilon + \delta})$$

Concentrates to informative spikes

Noise/ Identity

**dominant low-degree term along
 $h^\star(\mathbf{x})$**

Generalization error of random features and kernel methods:
hypercontractivity and kernel matrix concentration

Song Mei*, Theodor Misiakiewicz†, Andrea Montanari†‡

Caveat:

With Gaussian inputs, some radial degree k polynomials require less than $O(d^k)$ samples.

Fitting the last layer

Fitting the last layer

$$\mathbf{h}(\mathbf{x}) = W_2\sigma(W_1\mathbf{x}) \approx c\mathbf{w}_3\mathbf{h}^\star(\mathbf{x})$$

Fitting the last layer

$$\mathbf{h}(\mathbf{x}) = W_2\sigma(\mathbf{W}_1\mathbf{x}) \approx c\mathbf{w}_3\mathbf{h}^{\star}(\mathbf{x})$$

- Reduction to Kernel on low-dimensional features.
- $K(\mathbf{x}_1, \mathbf{x}_2) = \mathbb{E}_{\mathbf{w} \sim \mathcal{N}(0,1)}[\sigma(c\mathbf{w}\mathbf{h}^{\star}(\mathbf{x}_1) + \mathbf{b})\sigma(c\mathbf{w}\mathbf{h}^{\star}(\mathbf{x}_2) + \mathbf{b})]$.
- $\hat{W}_3 \approx \text{KRR}(K(\cdot), \mathbf{X}, \mathbf{y})$.
- Sample complexity/width now dimension-independent.

Takeaways

Takeaways

**Hierarchical functions with robustness w.r.t intermediate features
allow exploitation of depth through dimension reduction**

Takeaways

Hierarchical functions with robustness w.r.t intermediate features
allow exploitation of depth through dimension reduction

Do we need narrowing of networks? No, consider:


$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}_1^{\star\top} P_2(W_1^\star \mathbf{x})}{\sqrt{d}}, \dots, \frac{\mathbf{a}_m^{\star\top} P_2(W_m^\star \mathbf{x})}{\sqrt{d}} \right), \mathbf{x} \in \mathbb{R}^d, m = \mathcal{O}(d)$$

Takeaways

Hierarchical functions with robustness w.r.t intermediate features
allow exploitation of depth through dimension reduction

Do we need narrowing of networks? No, consider:

$$f^\star(\mathbf{x}) = g^\star \left(\frac{\mathbf{a}_1^{\star\top} P_2(W_1^\star \mathbf{x})}{\sqrt{d}}, \dots, \frac{\mathbf{a}_m^{\star\top} P_2(W_m^\star \mathbf{x})}{\sqrt{d}} \right), \mathbf{x} \in \mathbb{R}^d, m = \mathcal{O}(d)$$



“approximately independent” $\mathcal{O}(d)$ features in $\mathcal{O}(d^2)$ space

Thanks to my collaborators!

